

Good morning !!!

SOLVING PROBLEMS BY SEARCHING

Chapter 3

Outline



- Problem-solving agents
- Problem formulation
- Example problems
- Uninformed search algorithms

Problem solving agent

- Particular type of **goal-based agent**
 - ▣ Goal-based agents act to achieve their goals
- **Problem-solving agents** use **atomic representation**
 - ▣ states of the world have no internal structure visible to the problem-solving algorithm
- **Solution:** fixed sequence of actions (in this chapter)

Problem solving agent



- **Intelligent agents** are supposed to maximize their **performance measure**

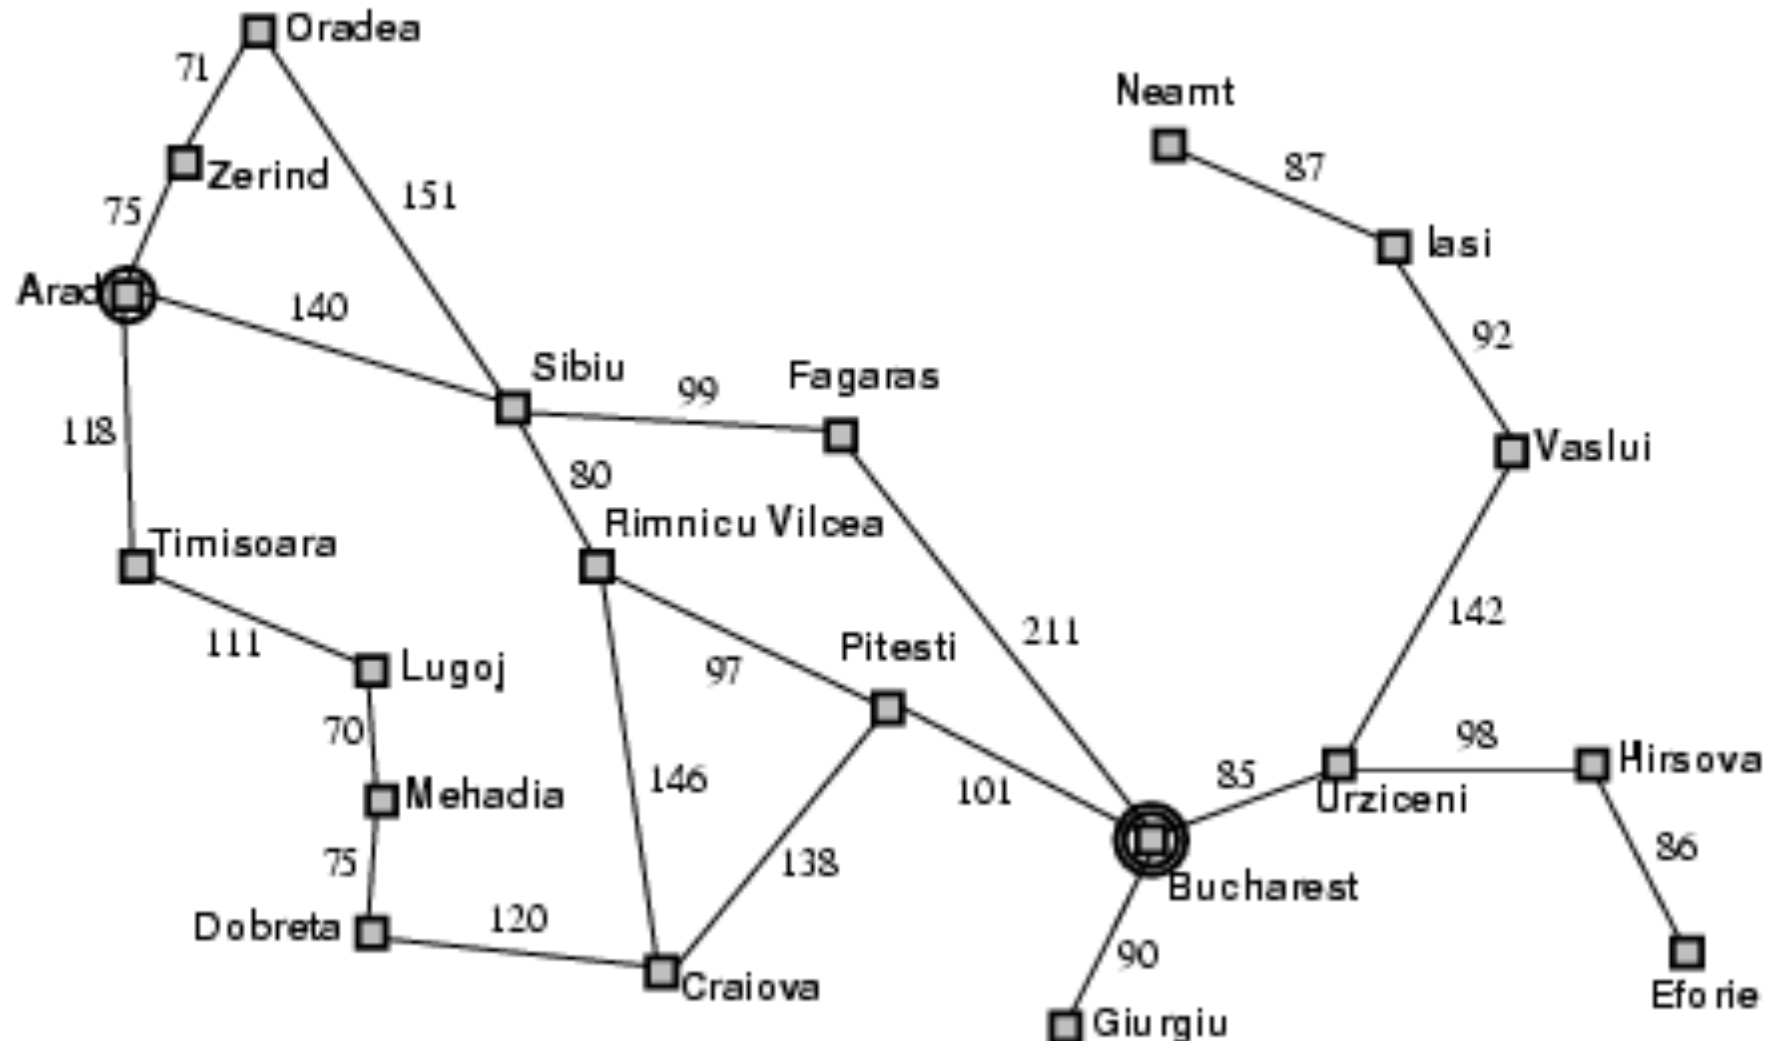
Achieving this is sometimes **simplified**

if the agent can **select a goal** and aim at **satisfying** it

Example: Romania

- On road **trip** in Romania. Currently in **Arad**
- **Flight back** leaves tomorrow **from Bucharest**
- **Formulate goal**
 - ▣ Be in Bucharest
- **Formulate problem**
 - ▣ **states**: various cities
 - ▣ **actions**: drive between cities
- **Find solution**
 - ▣ sequence of cities (e.g., Arad, Sibiu, Fagaras, Bucharest)

Example: Romania



Formulate, search and execute

- If the **environment** is:

- observable (the agent always knows the current state)
- known (the agent knows which states are reached by each action)
- deterministic (each action has exactly one outcome)

→ a **solution** to a problem is a fixed sequence of actions

- **Search:** process of looking for such a sequence

- **Search algorithm:**

- **input:** problem
- **output:** an action sequence

- The sequence in output can then be **executed**

Problem-solving agent

Goal formulation based on the current situation and the agent's performance measure

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
         state, some description of the current world state
         goal, a goal, initially null
         problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

Problem formulation: what **states** and **actions** to consider given the goal

Search: Decide **what to do next** by exploring consequences of actions in the future and if they lead to the goal

Problem-solving agents



This simple **problem-solving agent**

1. **Formulates** a **goal** and a **problem**
2. **Searches** for a *sequence of actions* that would solve the problem
3. **Executes the actions** one at a time

When this is complete,
it formulates **another goal** and **starts over**

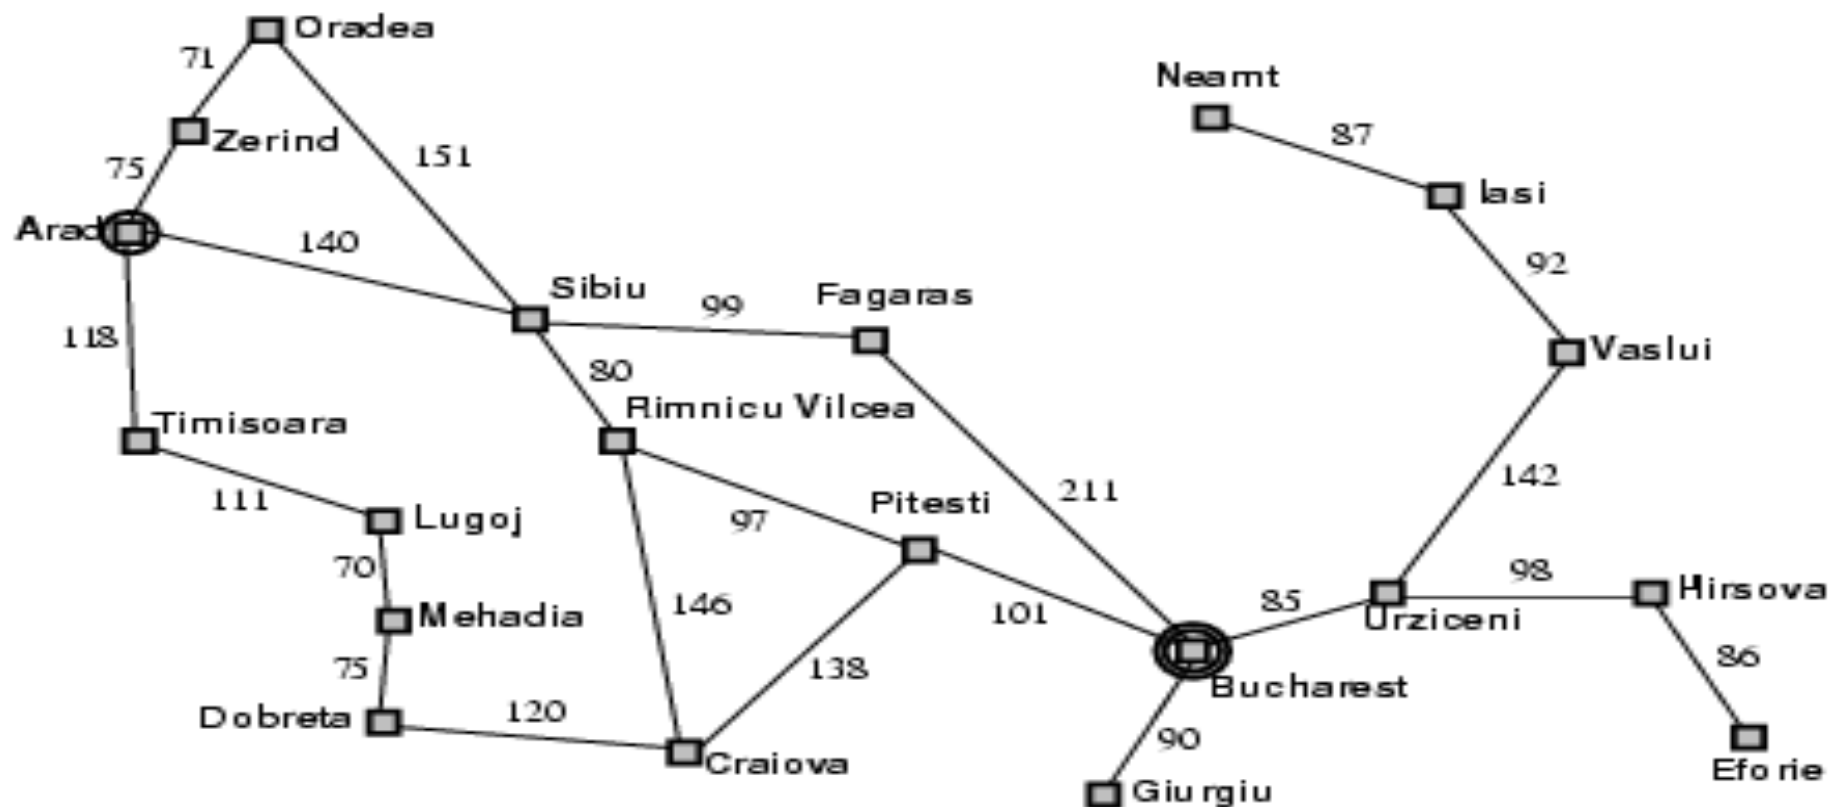
Outline



- Problem-solving agents
- Problem formulation
- Example problems
- Uninformed search algorithms

Example: Romania

- On road **trip** in Romania.
- Currently in **Arad**
- **Flight back** leaves tomorrow **from Bucharest**



Example: Romania

- Currently in **Arad**
- **Flight back** leaves tomorrow **from Bucharest**

- **Formulate goal**

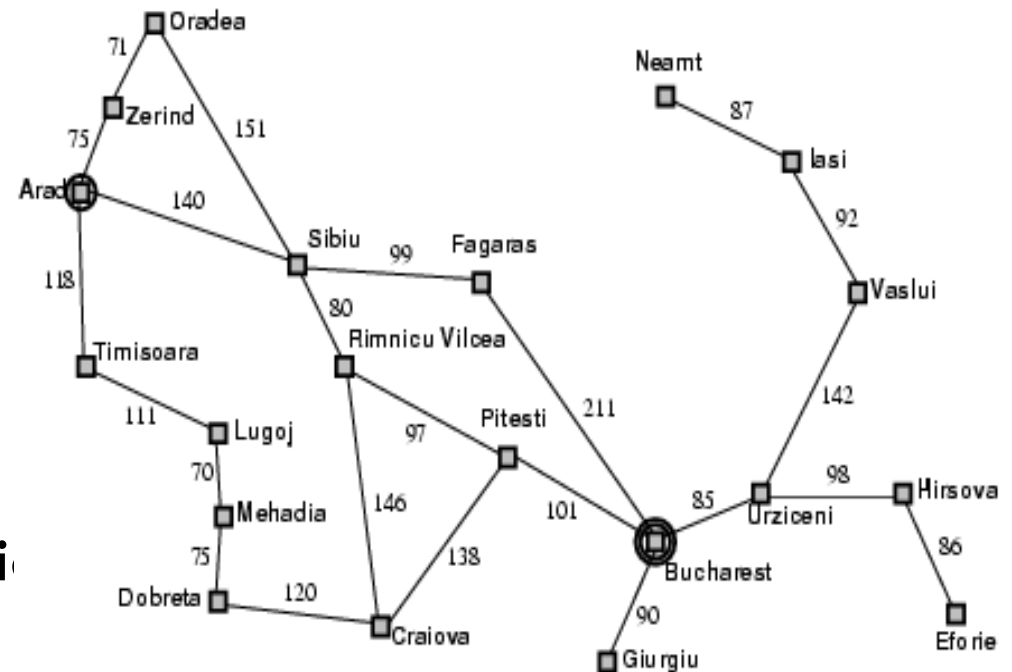
- Be in Bucharest

- **Formulate problem**

- **states**: various cities
 - **actions**: drive between cities

- **Find solution**

- sequence of cities (e.g., Arad, Sibiu, Fagaras, Bucharest)



Problem formulation



A **problem** can be defined formally by:

- **Initial state**
- **Actions**
- **Transition model**
- **Goal test**
- **Path cost**

Solution: sequence of actions (path) leading from the initial state to a goal state

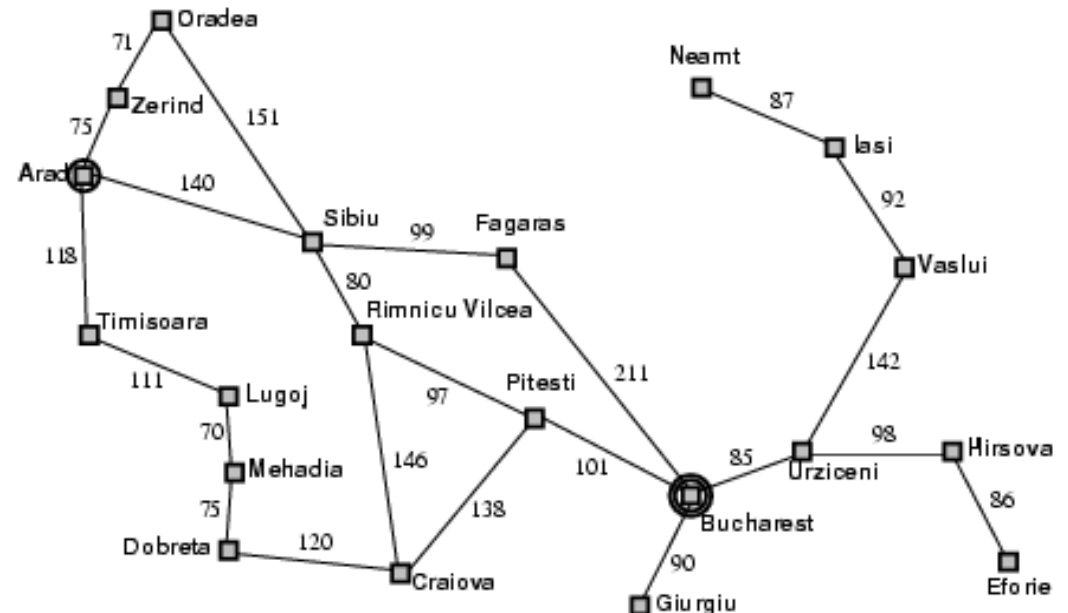
Optimal solution: a solution with minimal path-cost

Problem formulation

A **problem** can be defined formally by:


- **Initial state** that the agent starts in

- **Example:** $\text{In}(\text{Arad})$



Problem formulation

- **Actions** available to the agent
 - ▣ Given a state s , **ACTIONS**(s) returns the set of actions that can be executed in s
 - ▣ We say that each of these actions is **applicable** in s
- **Example:** from the state $\text{In}(\text{Arad})$, the applicable actions are $\{\text{Go}(\text{Sibiu}), \text{Go}(\text{Timisoara}), \text{Go}(\text{Zerind})\}$.

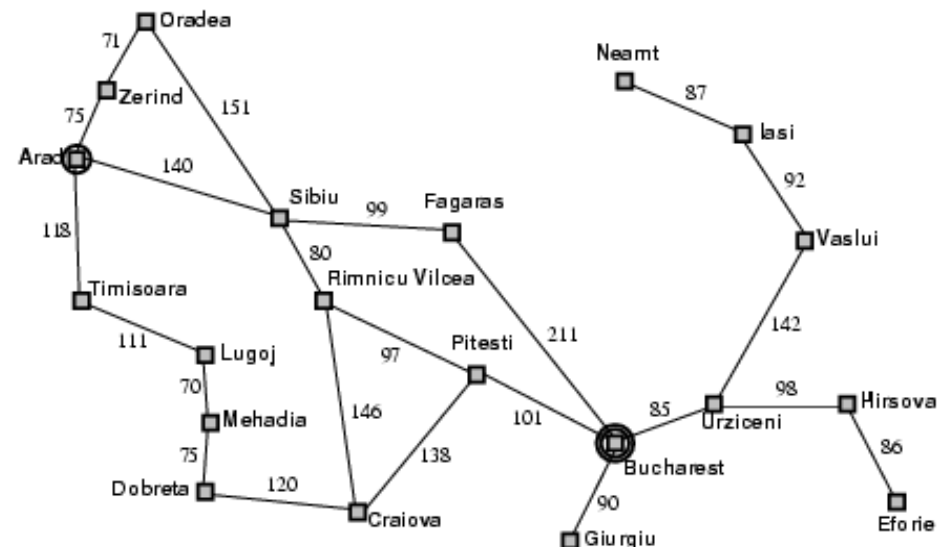


Problem formulation

□ Transition model:

- **RESULT**(s, a) returns the state obtained from doing action a in state s
- **Successor**: any state reachable from a given state by a single action

■ **Example:** **RESULT** (In(Arad), Go(Zerind)) = In(Zerind)



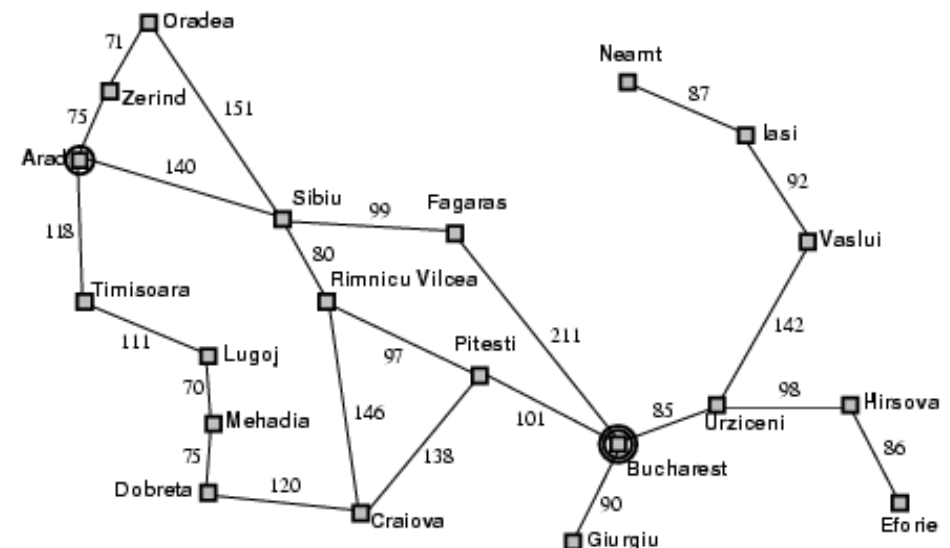
Problem formulation

- Problem **state space** = (initial state, actions, transition model)
- **State space** = the set of all states reachable from the initial state by any sequence of actions
- The **state space** can be depicted as a **directed graph**:
 - nodes** → states
 - edges** → actions
 - path** → sequence of states *connected* by a sequence of actions

Problem formulation

□ **Goal test:** allows to check if a state is a **goal**

■ **Example:** The agent's goal in Romania is the singleton set $\{ \text{In(Bucharest)} \}$



Problem formulation

The **state space** can be depicted as a **directed graph**:

nodes \rightarrow states

edges \rightarrow actions

path \rightarrow sequence of states connected by a sequence of actions


Path cost: **numeric value** associated to each path reflecting the desired performance measure

- e.g., sum of distances, number of actions executed, etc.
- $c(x,a,y)$ is the **step cost**, for going from **state x** to **state y** by performing **action a**
- assumed to be ≥ 0
- We assume path costs to be additive: **sum of step costs**

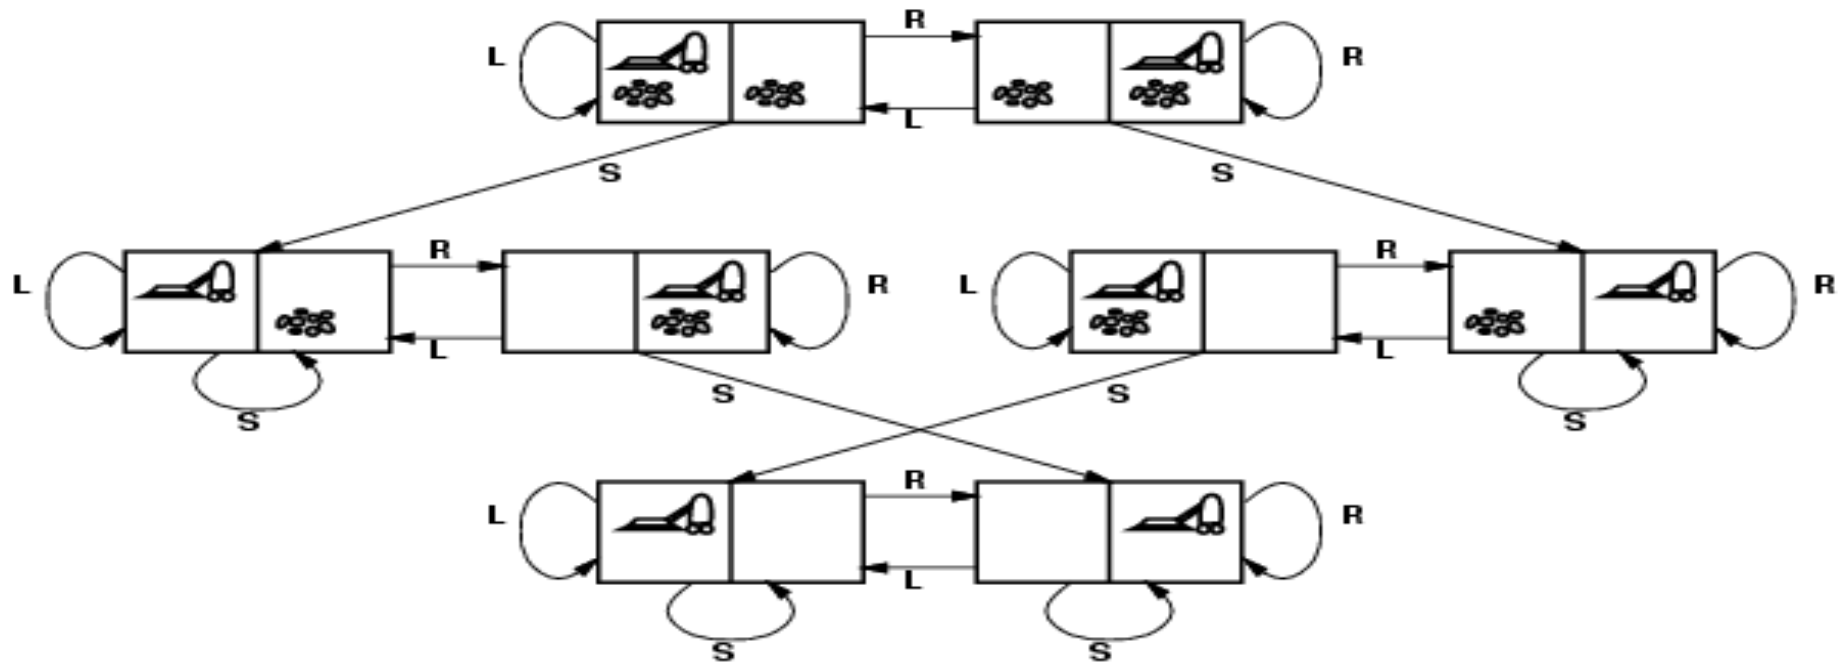
Problem formulation

- **Solution:** a sequence of actions (path) leading from the initial state to a goal state
- **Optimal solution:** a solution with minimal path-cost

Example: Vacuum-cleaner world




Example: Vacuum world



- states?
- initial state?
- actions?
- transition model?
- goal test?
- path cost?

Example: Vacuum world



- states? dirt locations and robot location
- initial state? any state
- actions? Left, Right, Suck
- Transition model? The actions have the expected effects, except that moving Left in the leftmost square, moving Right in the rightmost square, and Sucking in a clean square have no effect
- goal test? no dirt at all locations
- path cost? Each step costs 1 → path cost is the number of steps in the path

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

The 8-puzzle

- A 3×3 board with **eight numbered tiles** and a **blank space**
- A **tile adjacent** to the blank space **can slide** into that space
- **Object:** to reach a specified goal state such as the one shown on the right of the figure

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- initial state?
- actions?
- transition model?
- goal test?
- path cost?

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states? locations of eight tiles and the blank in one of the nine squares
- initial state? any
- actions? move blank left, right, up, down
- transition model? Given a state and an action, it returns the resulting state (that is, the puzzle configuration after moving the blank)
- goal test? goal state (given)
- path cost? Each step cost 1 per move, path cost is the number of steps in the path

Outline



- Problem-solving agents
- Problem formulation
- Example problems
- Uninformed **search algorithms**

Review: Problem-solving agent



This simple **problem-solving agent**

1. **Formulates** a **goal** and a **problem**
2. **Searches** for a **sequence of actions** that would solve the problem
3. **Executes** the **actions** one at a time

When this is complete,
it formulates **another goal** and **starts over**

Review: Problem formulation

A **problem** can be defined formally by:

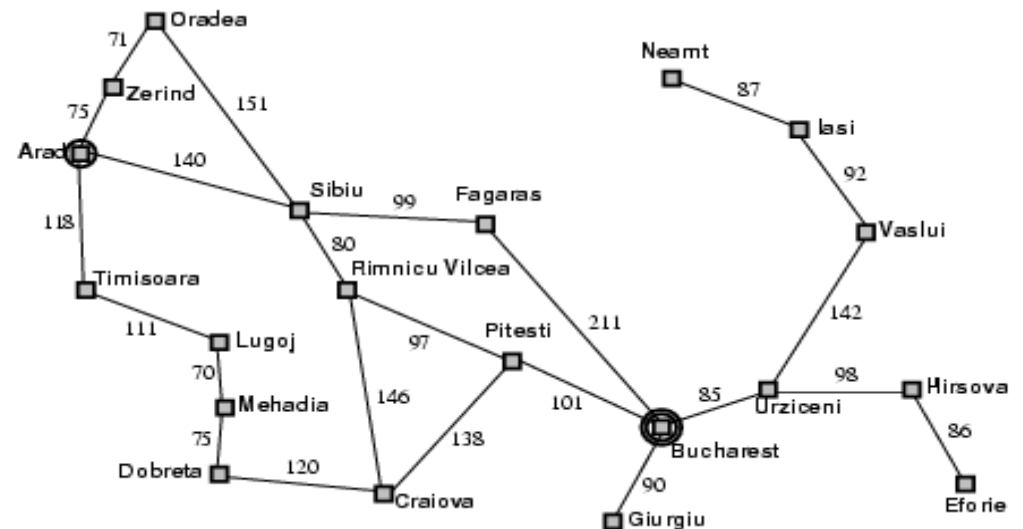
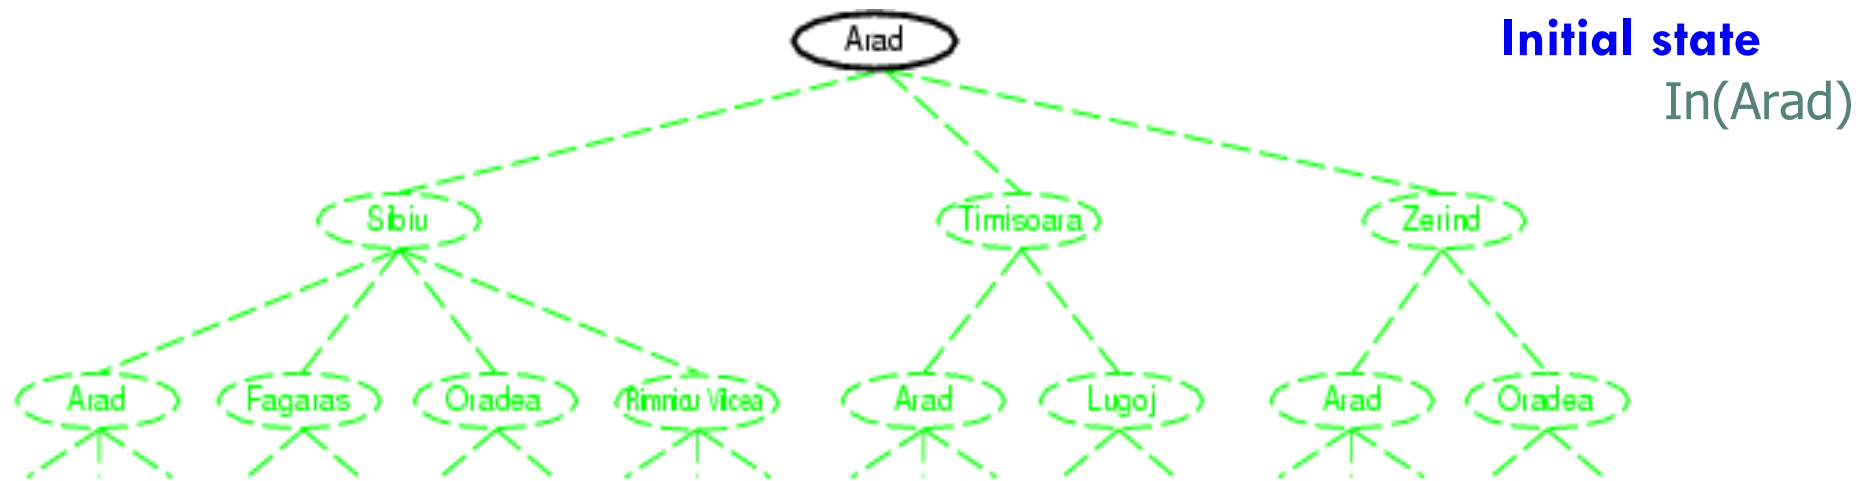
- **Initial state** that the agent starts in
- **Actions** available to the agent
- **Transition model:** A description of what each action does
- **Goal test:** Determines whether a given state is a **goal state**
- **Path cost: Numeric value** associated to each path reflecting the desired performance measure

The Search Tree

Solution: a **sequence of actions** (path) leading from the initial state to a goal state


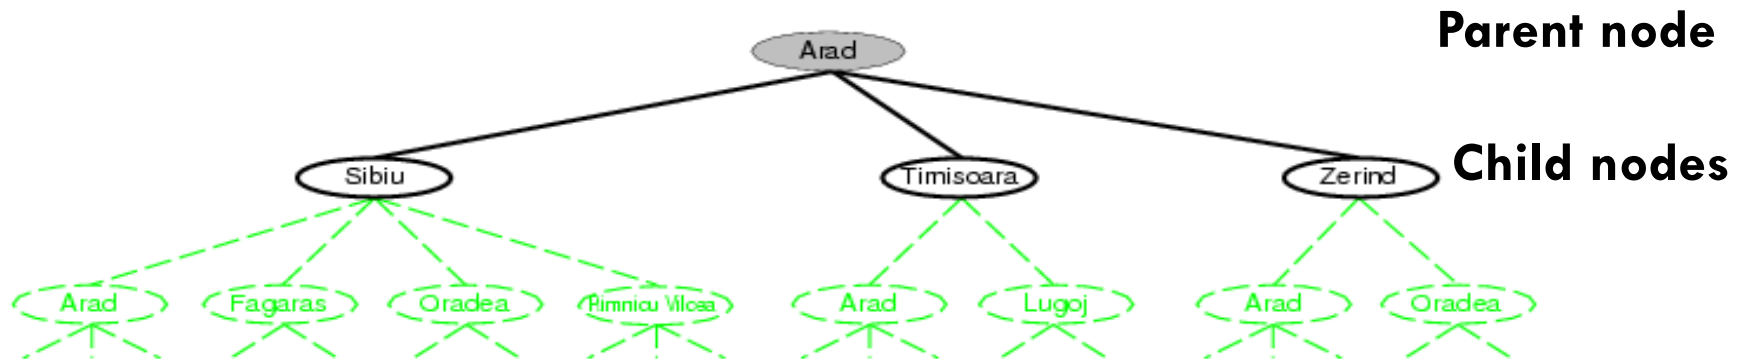
- **Search algorithms** consider possible sequences of actions
- Possible sequences of actions from initial state form a **search tree**
 - **Root** → initial state
 - **Nodes** → states
 - **Branches** → actions
 - The **same state** can appear **multiple time**
 - **Outgoing edges** from a node →
all possible actions available in the state represented by the node

Tree search example



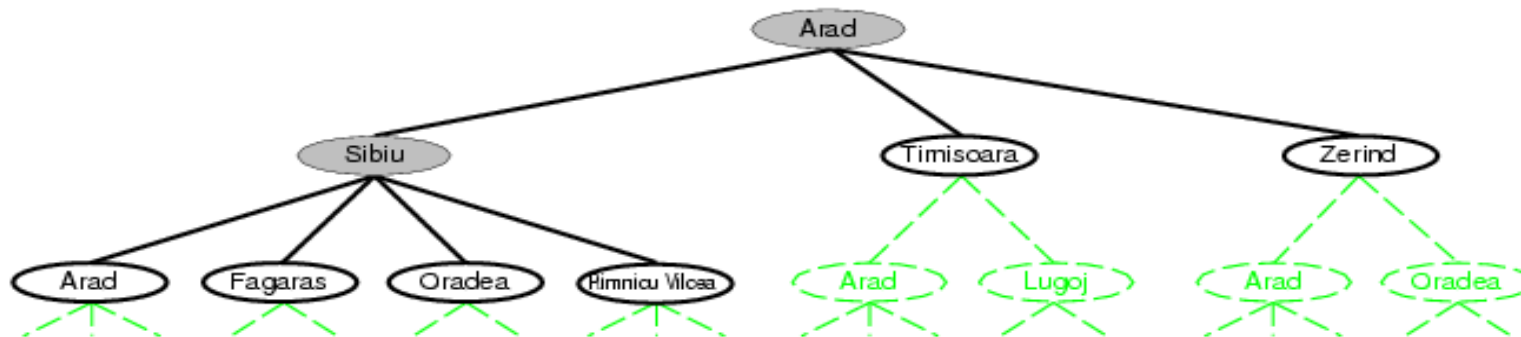
Tree search example

After expanding Arad



Tree search example

After expanding Sibiu

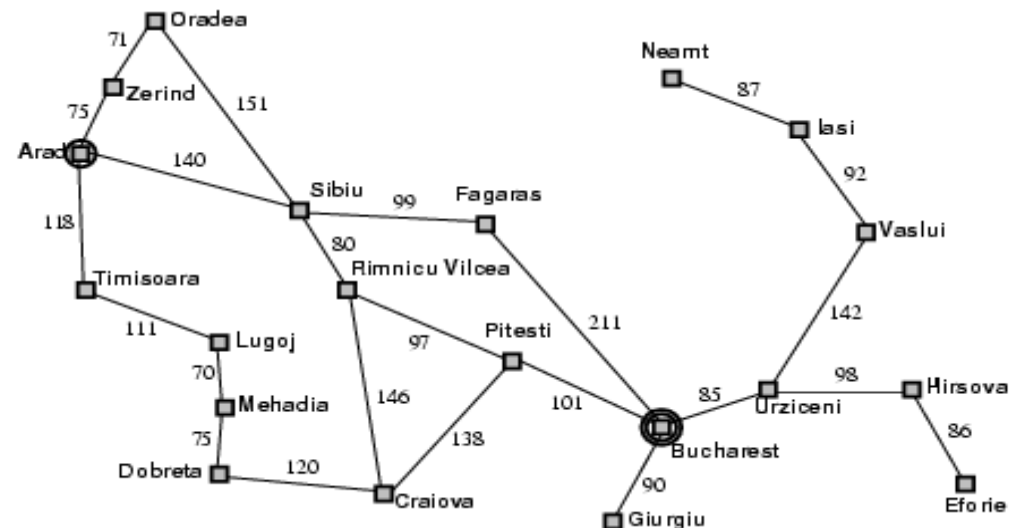


Leaf node:

a node with no children in the tree

Frontier:

The set of **all leaf nodes** available for expansion at any given point



Tree search algorithm

function **TREE-SEARCH**(problem) **returns** a **solution**, or **failure**

initialize the **frontier** using the **initial state** of problem

loop do

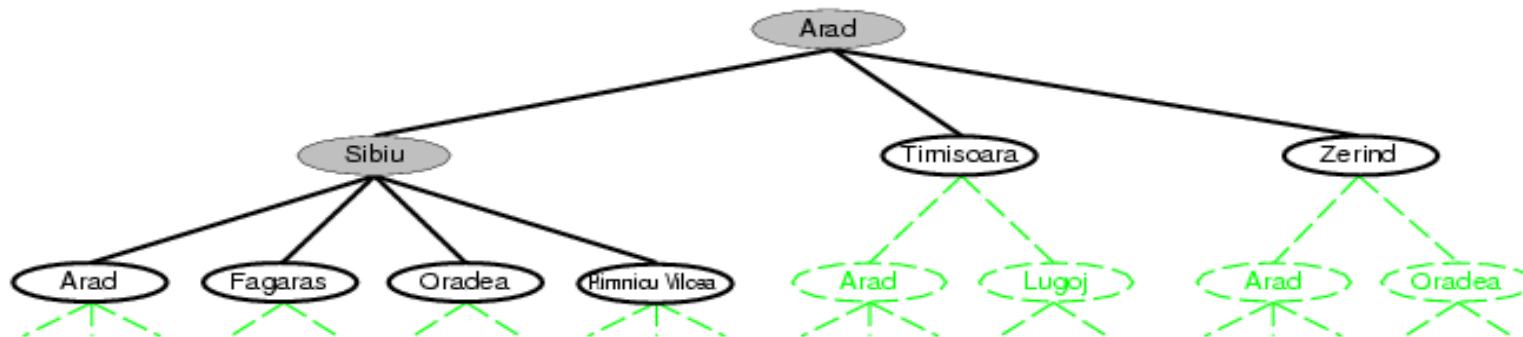
if the **frontier** is empty **then return failure**

choose a **leaf node** and remove it **from the frontier**

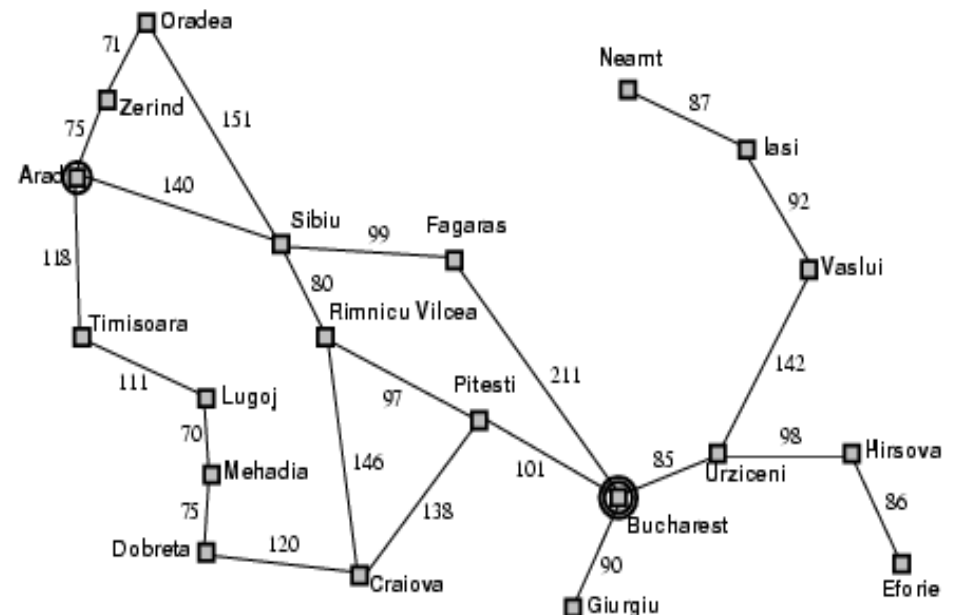
if the node contains a goal state **then return** the corresponding **solution**

expand the **chosen node**, adding the resulting nodes to the **frontier**


Tree search example




- The search tree includes the path **from Arad to Sibiu** and **back to Arad again!**
- The state **In(Arad)** is a **repeated state** generated by a **loopy path**




Tree search example



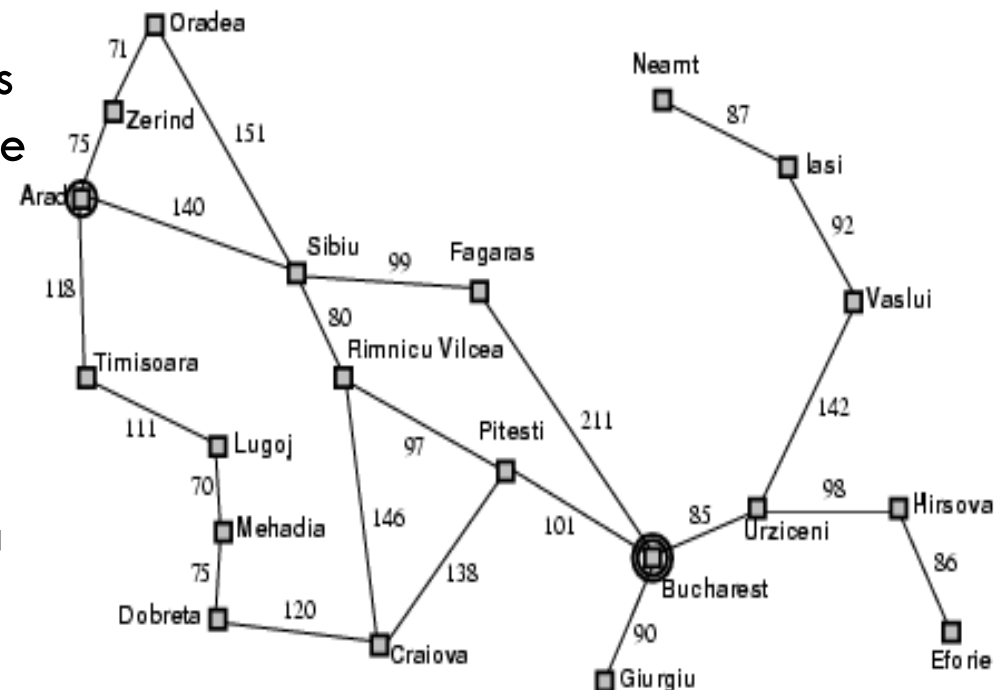
- **Loopy paths:** special case of **redundant paths**
- **Redundant paths:** exist whenever there is **more than one way** to get from one state to another




Tree search example



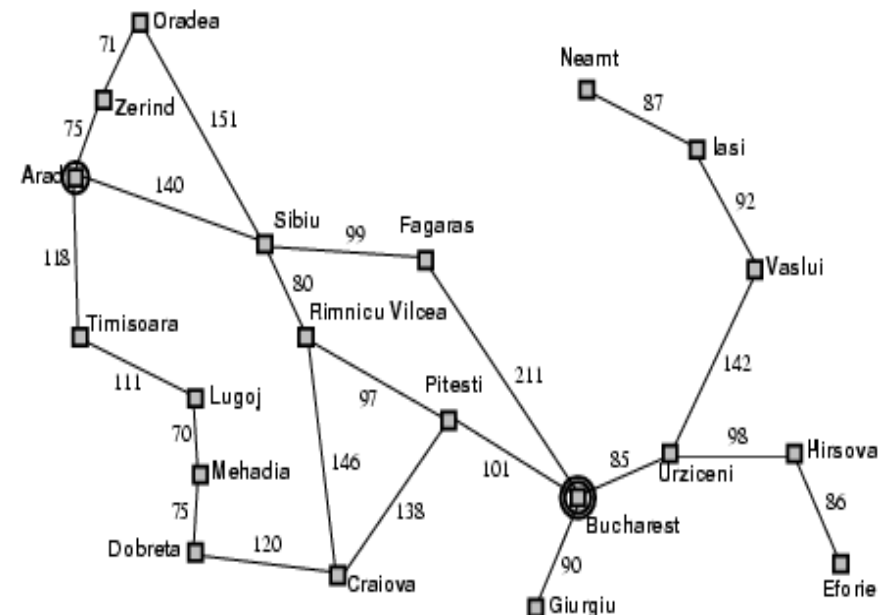
- **Redundant paths:** exist whenever there is **more than one way** to get from one state to another
- **Example:** Consider the path
 - Arad–Sibiu (140)
 - Arad–Zerind–Oradea–Sibiu (297)
 - the **second** path is **redundant**—it's just a worse way to get to the same state



Tree search example



To avoid exploring **redundant paths**
TREE-SEARCH algorithm is
augmented with **explored set** that
remembers **every expanded node**



Graph search algorithm (le parti in marrone sono nuove)

function **GRAPH-SEARCH**(problem) **returns** a **solution**, or **failure**

initialize the **frontier** using the **initial state** of problem

initialize the explored set to be empty

loop do

if the **frontier** is empty **then return failure**

choose a **leaf node** and remove it from the **frontier**

if the node contains a goal state **then return** the corresponding **solution**


add the node to the explored set

expand the **chosen node**, adding the resulting nodes to the **frontier**


only if not in the frontier or explored set

Infrastructure for search algorithms

- **Search algorithms** require a **data structure** to keep track of the **search tree**
- **For each node n** of the tree, we have a structure with:
 - **n .STATE**: the **state** in the state space to which the **node corresponds**
 - **n .PARENT**: the **node** in the search tree that **generated this node**
 - **n .ACTION**: the **action** that was **applied** to the parent **to generate the node**
 - **n .PATH-COST**: the **cost**, traditionally denoted by **$g(n)$** , of the **path from the initial state to the node**, as indicated by the parent pointers



Infrastructure for search algorithms



- A **state** corresponds to a configuration of the world
- A **node** is a **data structure** used to represent the search tree

Frontier: The set of **all leaf nodes** available **for expansion** at any given point

Infrastructure for search algorithms

- **Frontier** needs to be stored in such a way that the **search algorithm** can easily choose the **next node** to expand
 - The appropriate **data structure** for this is a **queue**

Operations on a queue:

- **EMPTY?(queue)** returns true only if there are no more elements in the queue
- **POP(queue)** removes the first element of the queue and returns it
- **INSERT(element, queue)** inserts an element and returns the resulting queue

Infrastructure for search algorithms

- **Queues** are characterized by the *order* in which they store the inserted nodes

Three common variants:

- **FIFO queue**

- which pops the *oldest* element of the queue

- **LIFO queue**

- which pops the *newest* element of the queue

- **Priority queue**

- which pops the element of the queue with the *highest priority* according to some ordering function

Search strategies

- A **search strategy** is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it **always** find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it **always** find a least-cost solution?

Search strategies



- **Time** and **space complexity** are measured in terms of
 - **b : branching factor** of the search tree
(i.e., maximum number of successors of any node)
 - **d : depth** of the **least-cost solution**
 - **m : the maximum depth** of the state space

Uninformed search strategies



- **Uninformed** strategies use **only the information** available in the **problem definition**
 - generate successors
 - distinguish goal

- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search