

INFORMED SEARCH ALGORITHMS

Chapter 3

Best-first search

- Informed search algorithms

use **problem-specific knowledge** to speed up the search process

Outline



- Best-first search
 - **Greedy** best-first search
 - **A*** search
- Heuristics

Review: Tree search algorithms

function **TREE-SEARCH**(problem) **returns** a solution, or failure

initialize the frontier using the initial state of problem

loop do

if the frontier is empty **then return** failure

choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

expand the chosen node, adding the resulting nodes to the frontier

A search strategy is defined by picking the **order of node expansion**

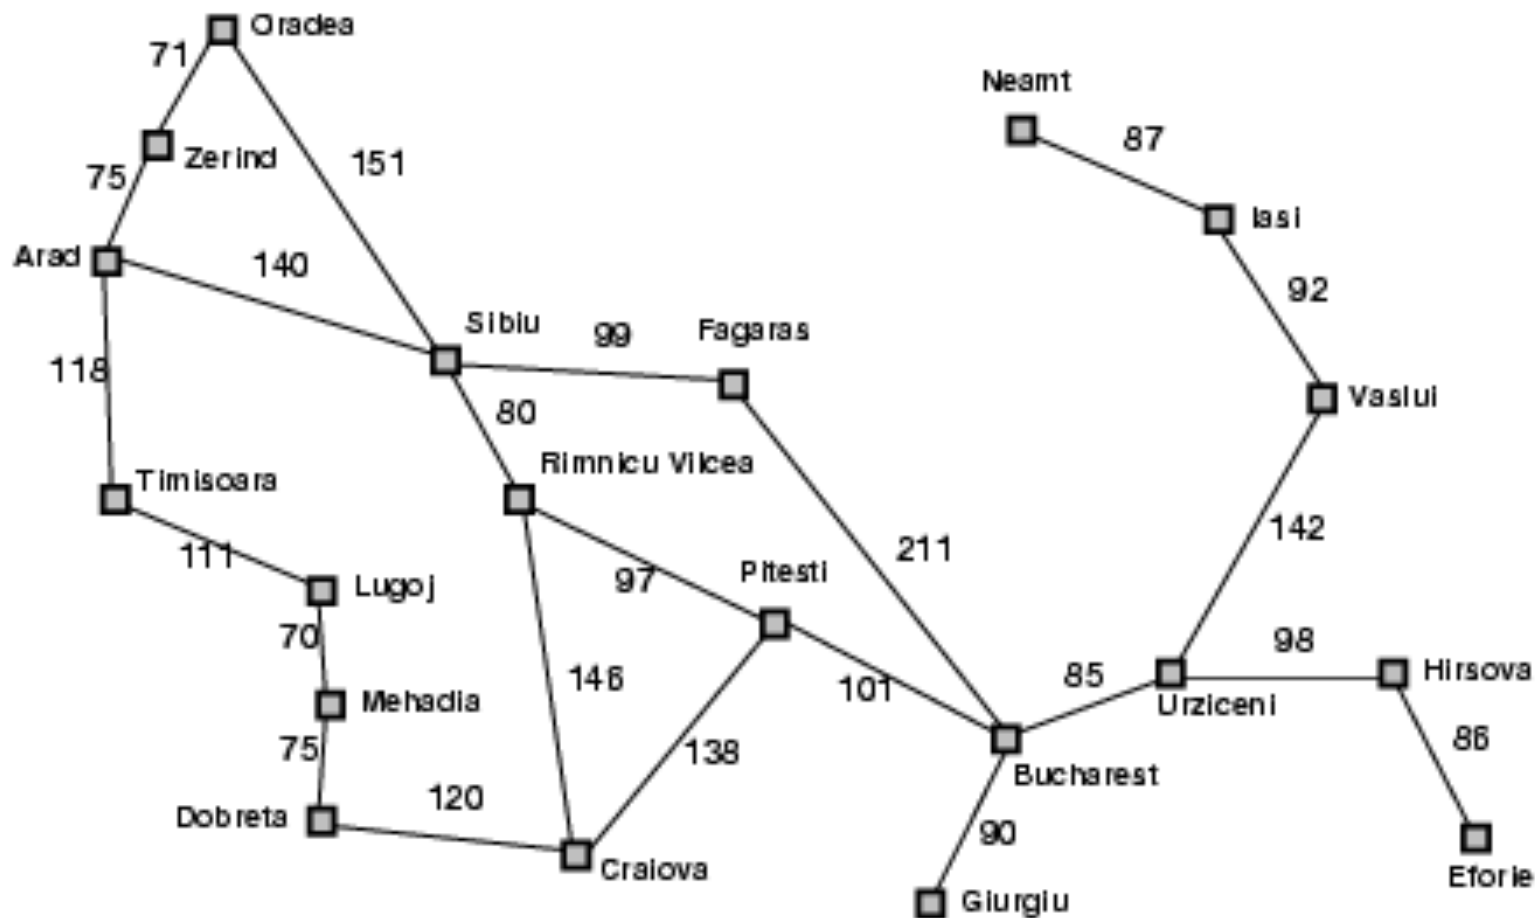
Best-first search

- Informed search algorithms
 - use **problem-specific knowledge** to speed up the search process
- Idea: use an **evaluation function $f(n)$** for each node
 - ▣ Estimate of “desiderability”
 - ▣ to choose the node selected for expansion
 - ▣ is construed as a cost estimate
 - ▣ Expand node with the **lowest $f(n)$** first
- Special cases
 - ▣ **Greedy** best-first search
 - ▣ **A*** search

Best-first search

- Most algorithms use a **heuristic function** $h(n)$ as a **component of f**
 - **$h(n)$** = estimated cost of the cheapest path
from the state of the **node n** to a **goal state**
 - depends only on the state associated with n
 - Assumptions:
 - $h(n)$ is non negative
 - $h(n)=0$ at every **goal state**
 - **Example:** in Romania, we may estimate
 - the **cost of the cheapest path** from Arad to Bucharest
 - **via** the straight-line distance from Arad to Bucharest
 - $h_{SLD}(n)$ = **straight-line distance** from n to **Bucarest**

Romania with step costs in km



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search

- Evaluation function $f(n) = h(n)$ (heuristic)
 $h(n)$: estimated cost of the cheapest path
from n to **goal**
- **Example**: $h_{SLD}(n) = \text{straight-line distance from } n \text{ to Bucharest}$
- **Greedy best-first** search expands the node that **appears** to be **closest** to goal

We expand first the node with the **lowest value** of $h(n)$

Greedy best-first search example

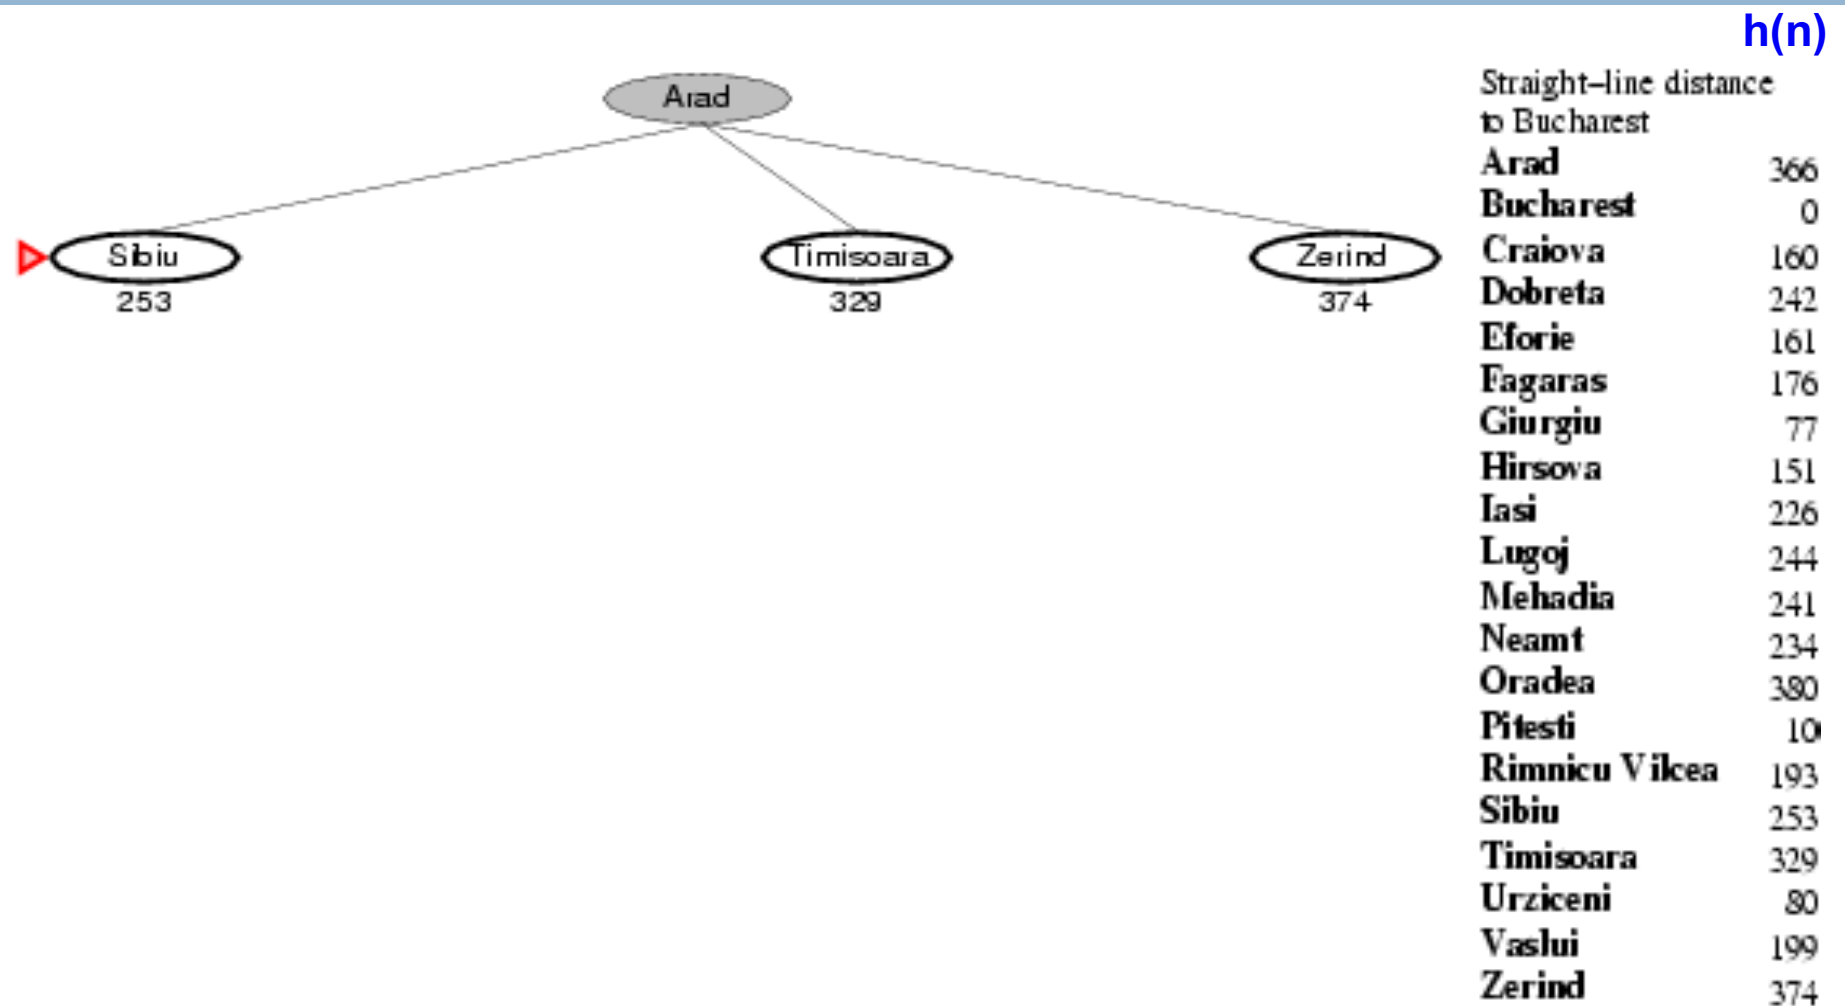


$h(n)$

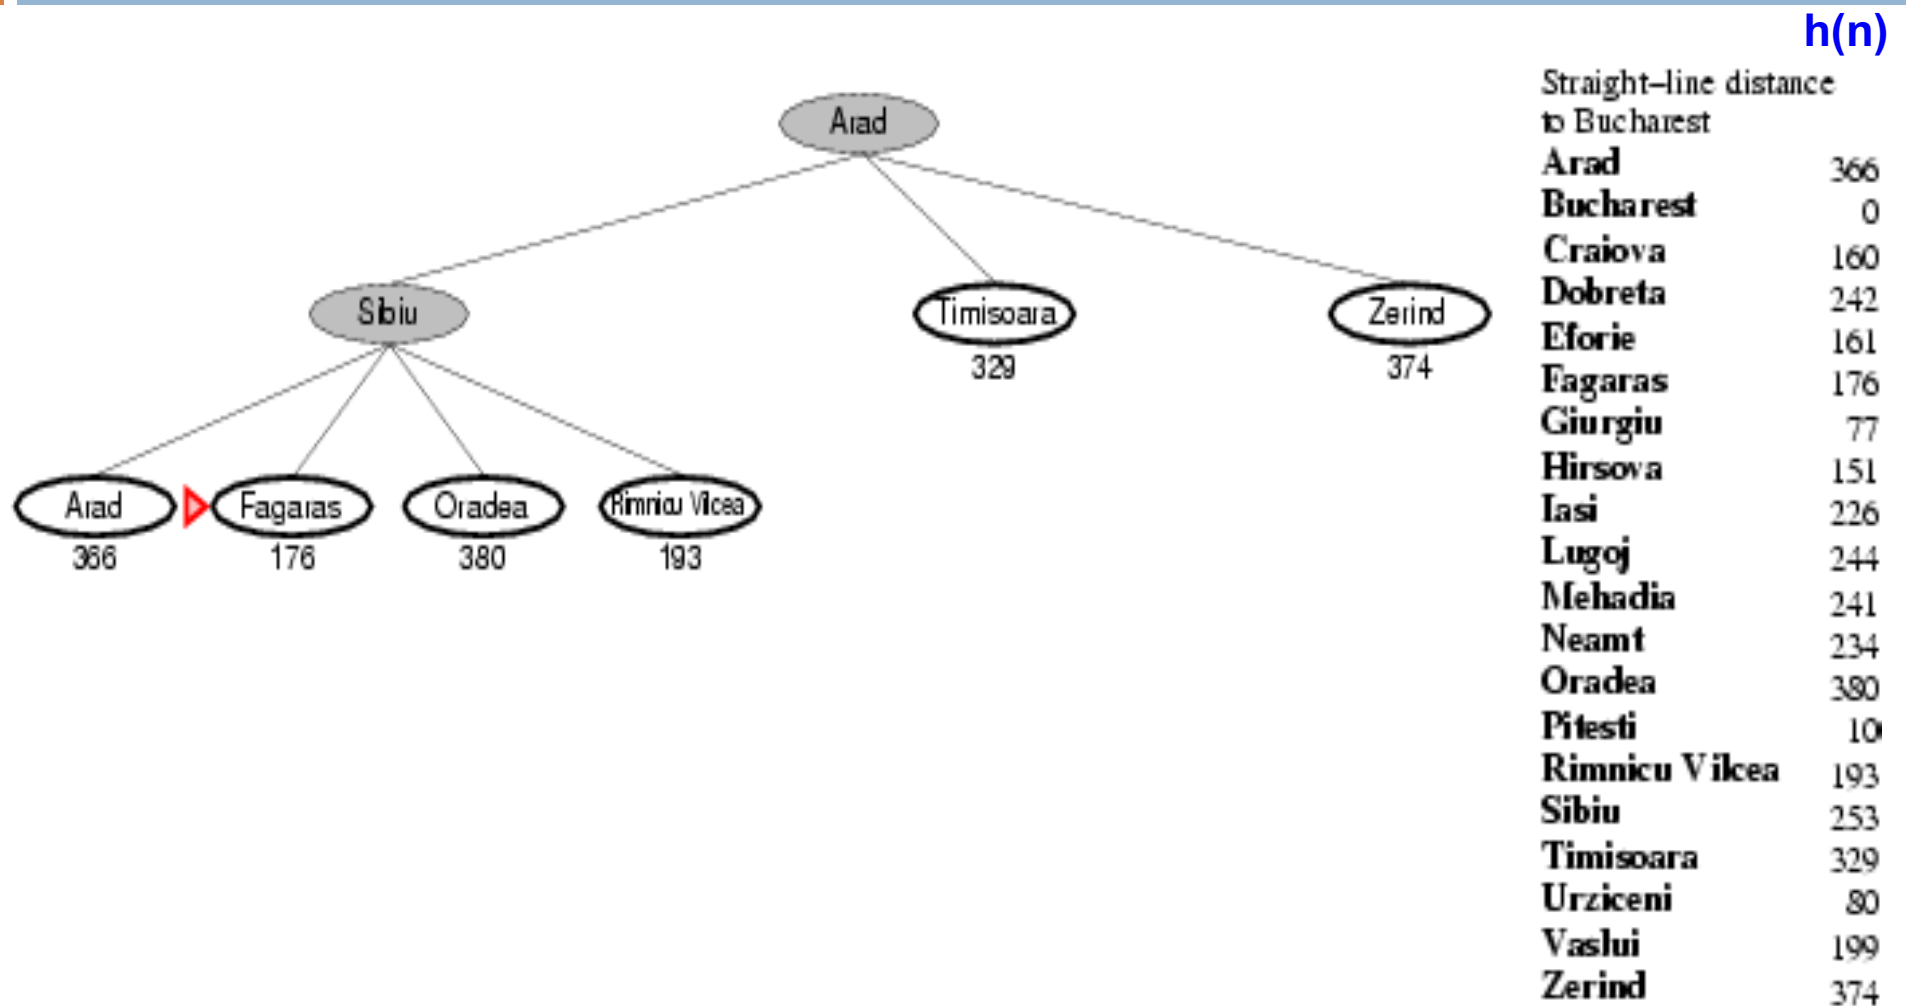
Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

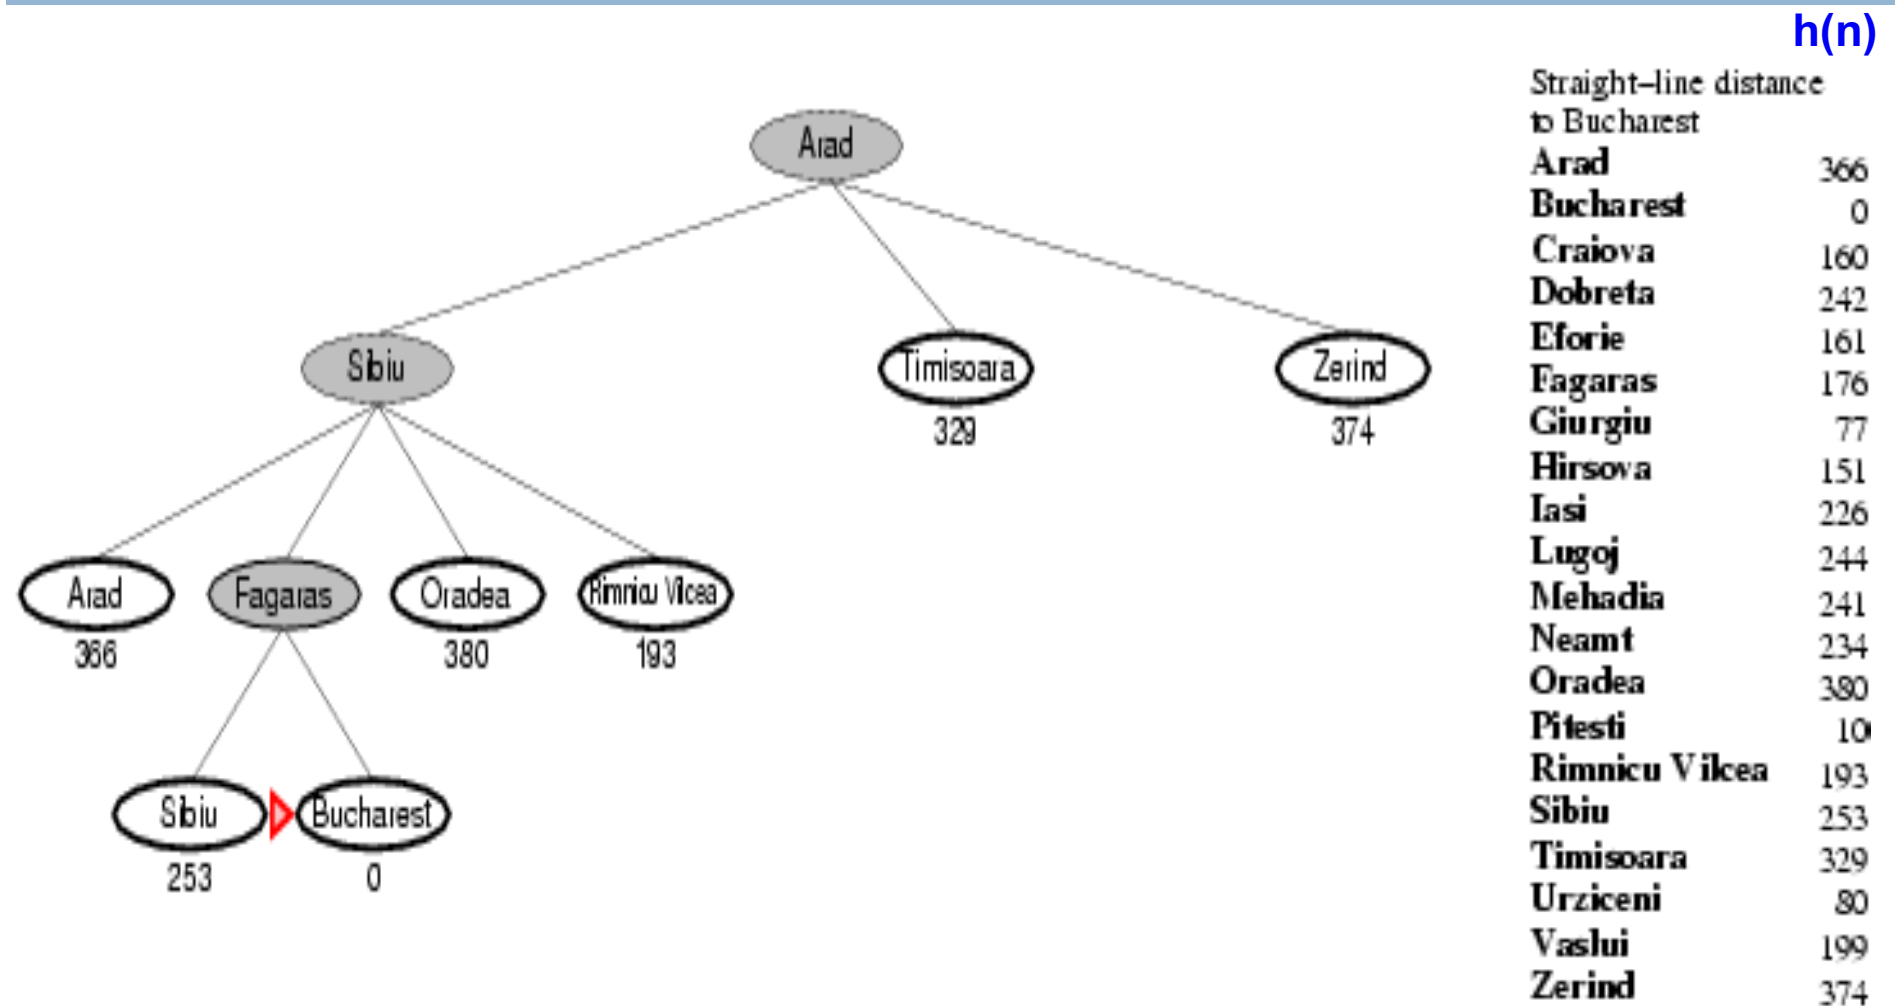
Greedy best-first search example



Greedy best-first search example




Greedy best-first search example



— returned by **greedy BFS**
— optimal

$h(n)$



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Properties of greedy best-first search

b : branching factor

m : the maximum depth of the search space

- Complete? No (tree-search)

- NO, consider the problem of going from Iasi to Fagaras: Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt


The heuristic suggests Neamt be expanded first because it is closest to Fagaras, but it is a dead end

- Complete in finite state with repeated-state checking

- Time? $O(b^m)$, but a good heuristic can give dramatic improvement

- Space? $O(b^m)$ -- keeps all nodes in memory

- Optimal? No



Outline



- Best-first search
 - **Greedy** best-first search
 - **A*** search
- Heuristics

A* search

- **Idea:** avoid expanding **paths** that are **already expensive**
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = **path cost** from the start node to **node n**
 - $h(n)$ = **estimated cost** of the cheapest path from n to **goal**
 - $f(n)$ = **estimated cost** of the cheapest solution through n
- Like **uniform-cost search** with **$f=g+h$** instead of **g**

A* search example

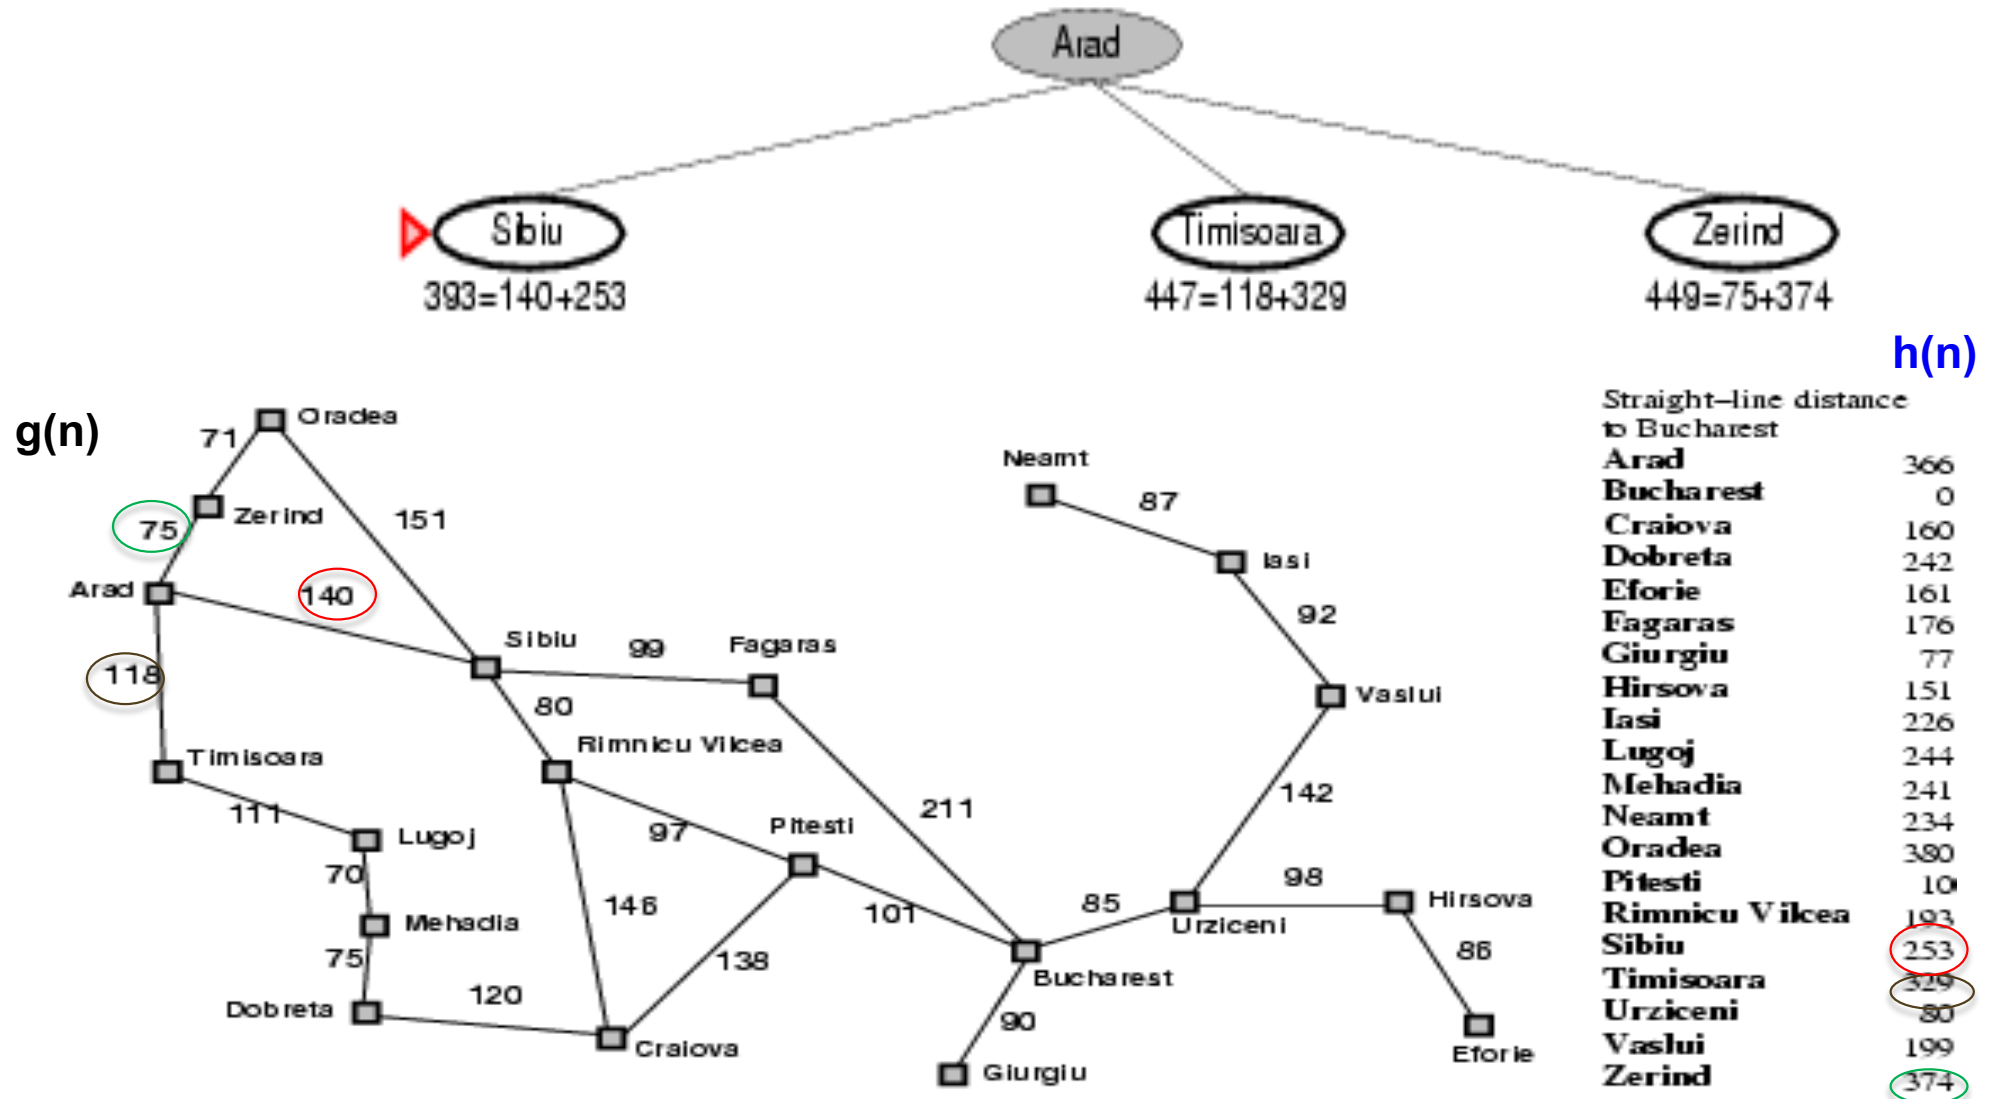
We expand first the node with the **lowest value** of $g(n) + h(n)$



Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

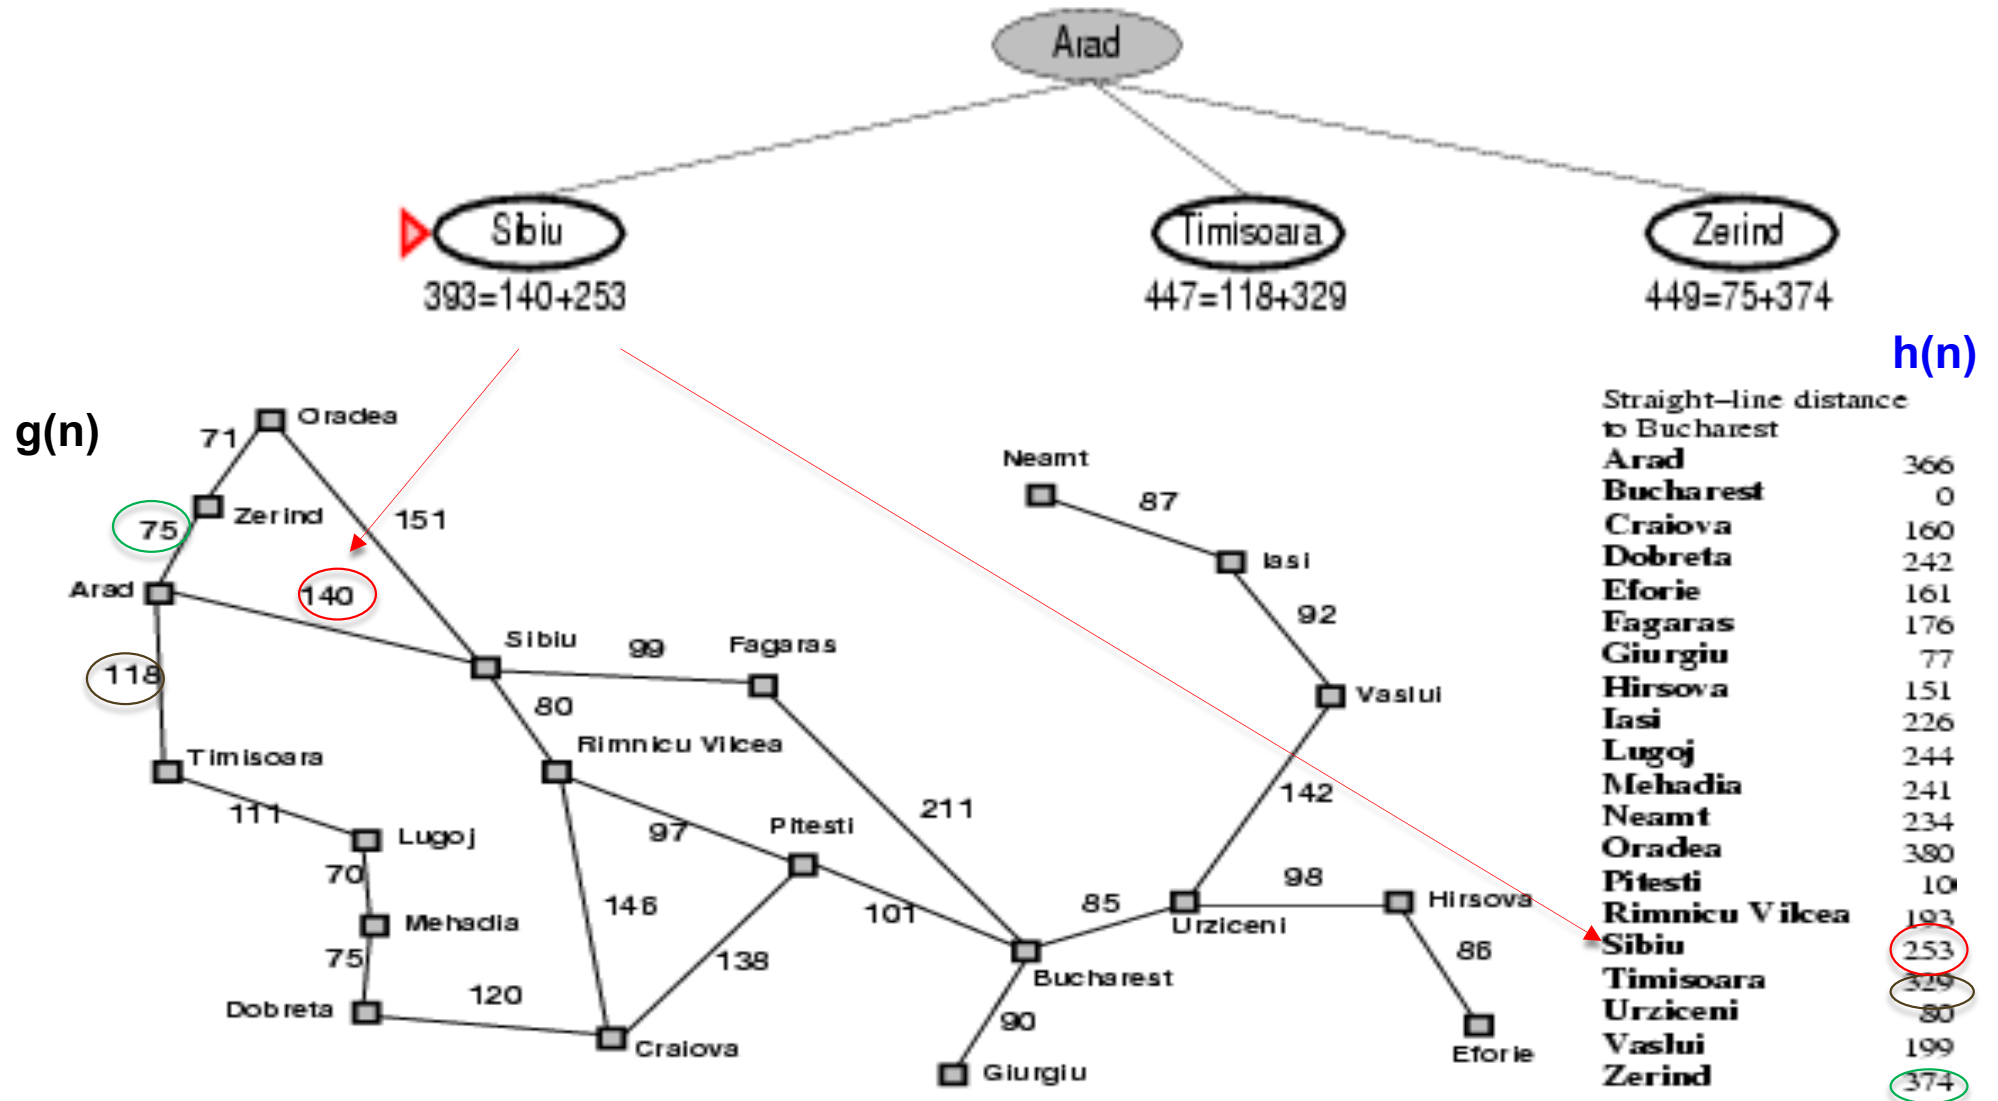
A* search example

We expand first the node with the **lowest value** of $g(n) + h(n)$



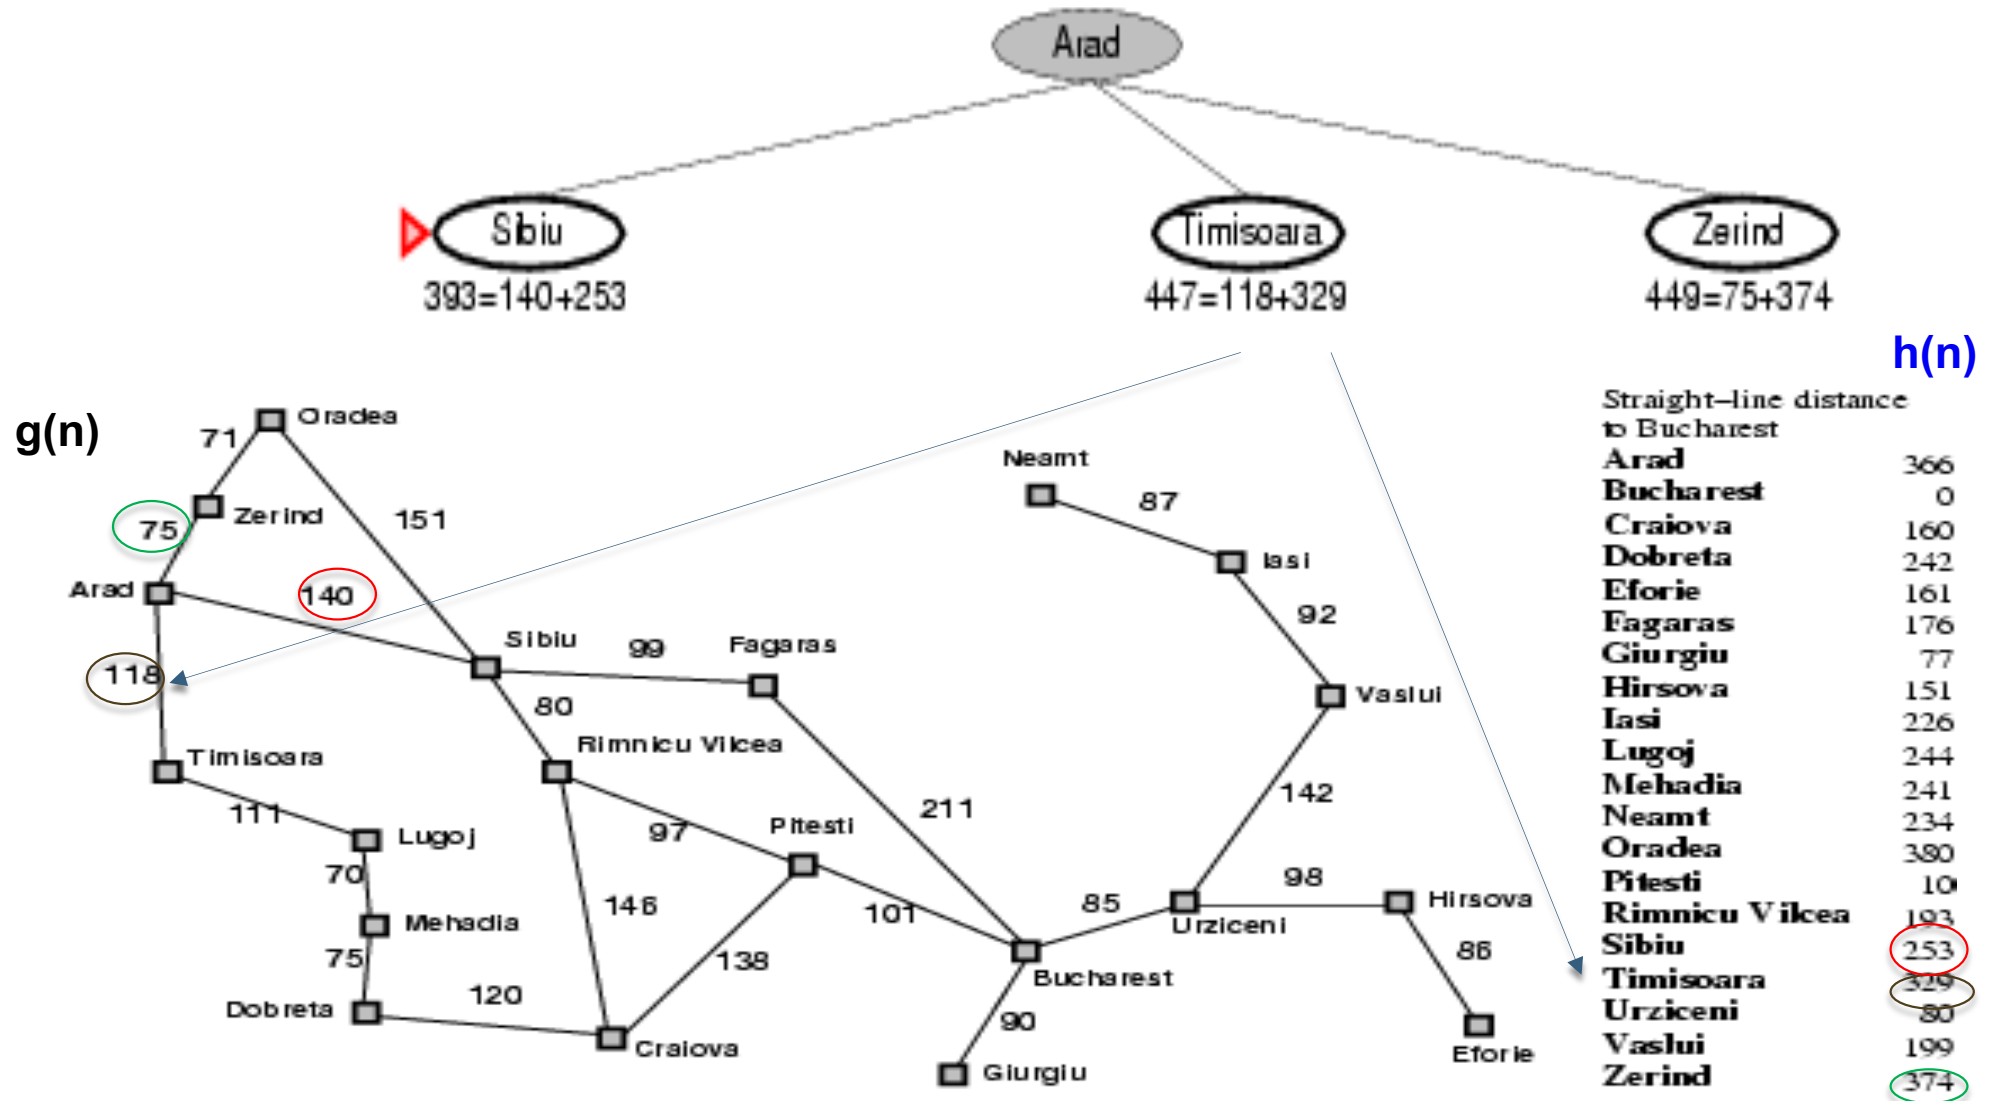
A* search example

We expand first the node with the **lowest value** of $g(n) + h(n)$



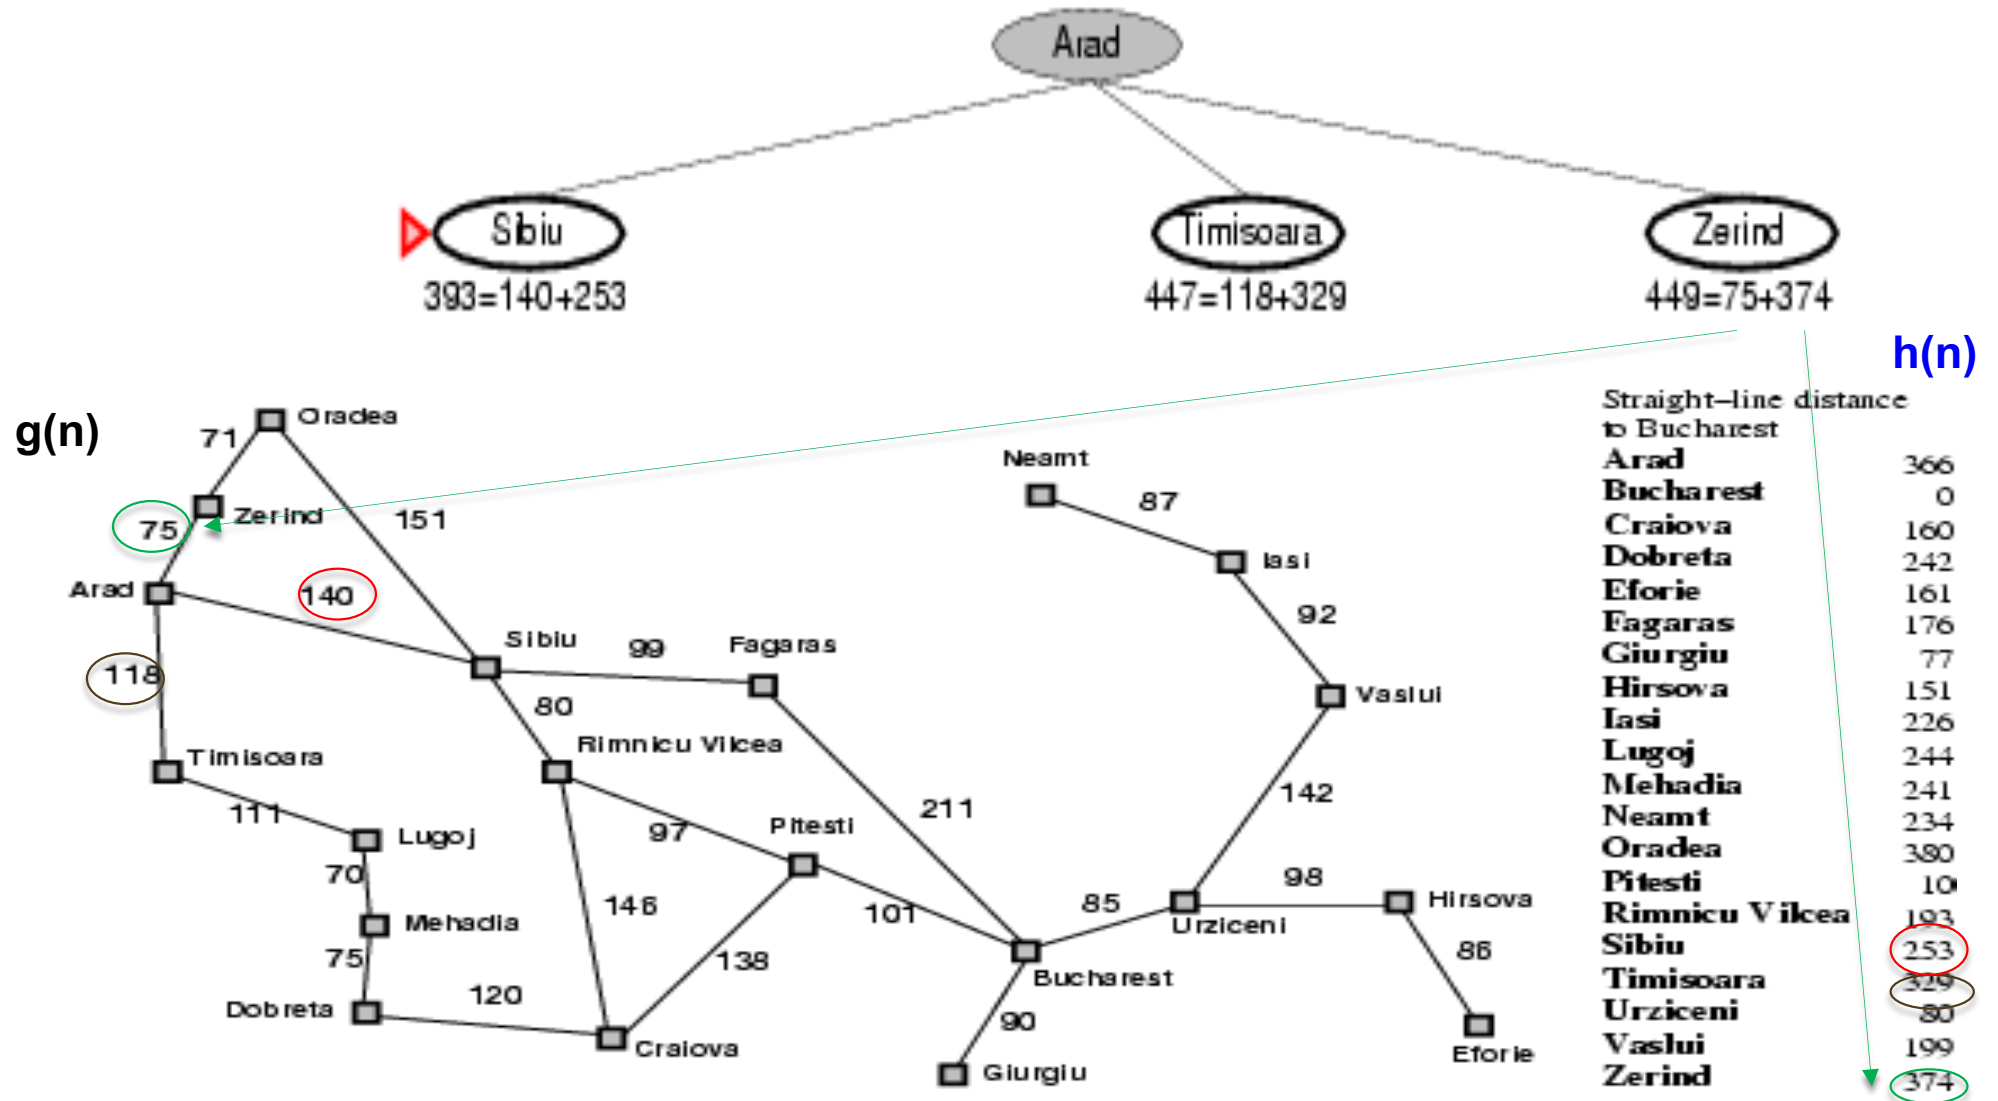
A* search example

We expand first the node with the **lowest value** of $g(n) + h(n)$




A* search example

We expand first the node with the **lowest value** of $g(n) + h(n)$

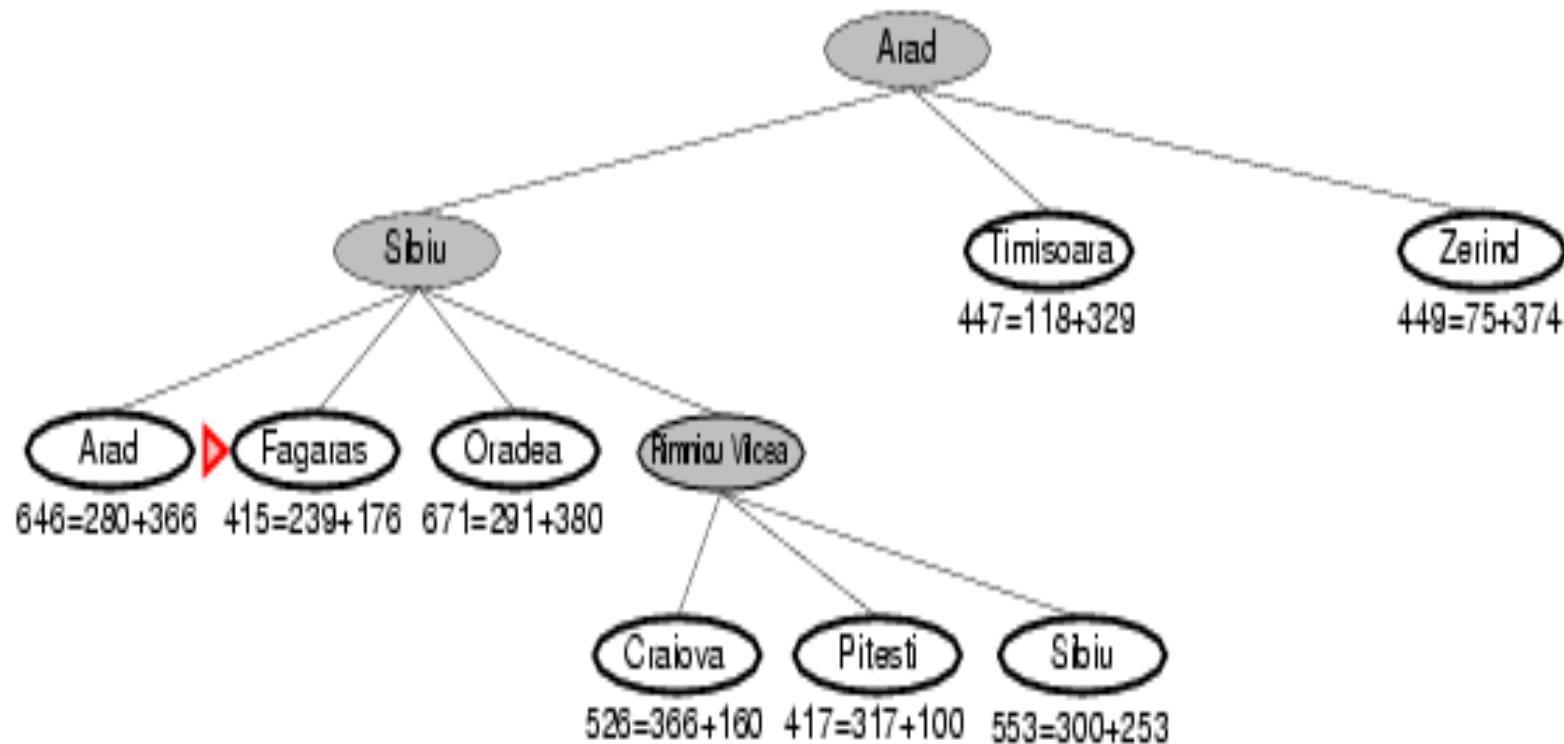


A* search example



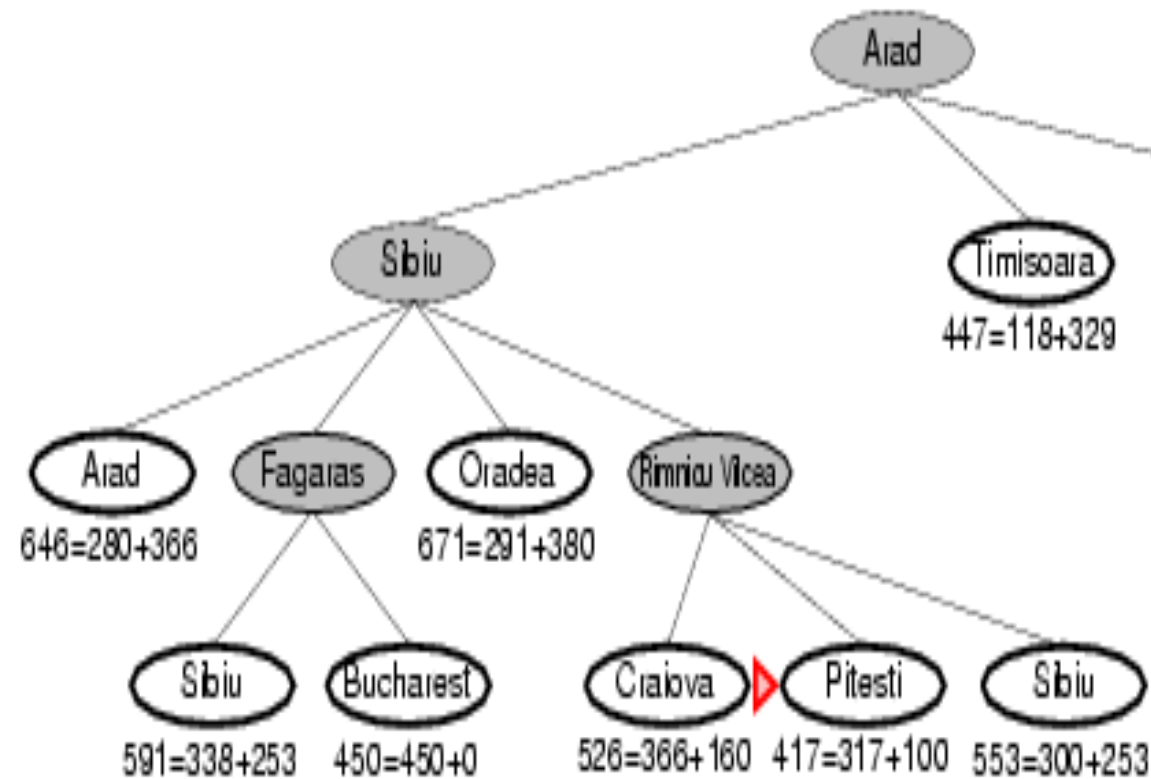
Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

A* search example



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

A* search example



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374


Admissible heuristics

- A heuristic h is **admissible** if
for every node n , $h(n) \leq h^*(n)$
where $h^*(n)$ is the true cost from n to goal
- An admissible heuristic **never overestimates** the cost to reach the goal
- **Example:** $h_{SLD}(n)$ (never overestimates the actual road distance)
- **Theorem:** If h is **admissible**,
 A^* using TREE-SEARCH is **optimal**

Consistent heuristics

- A heuristic h is **consistent**
if **for every** node n ,
for every successor n' of n generated **by an action** a ,
$$h(n) \leq c(n, a, n') + h(n')$$

- Instance of triangular inequality
- Consistency \rightarrow admissibility



h consistent $\rightarrow A^*$ graph-search optimal



- **Theorem:** If h is consistent
 A^* using GRAPH-SEARCH is optimal

Properties of A*

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? Exponential
- Space? Keeps all nodes in memory
- Optimal? Yes (with TREE-SEARCH, admissible heuristic;
GRAPH-SEARCH, consistent heuristic)

Admissible heuristics

Review: **8-puzzle**

goal: to slide the tiles

horizontally or vertically into the empty space **until** the configuration matches the goal configuration

E.g., for the 8-puzzle:

- $h_1(n)$ = number of **misplaced** tiles
- $h_2(n)$ = total Manhattan distance
(i.e, the **sum** of the horizontal and vertical **distances** of the tiles from their goal positions)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $\underline{h_1(S)} = ?$
- $\underline{h_2(S)} = ?$

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S)?$
- $h_1(S) = 8$
- **Admissible:** any tile out of place must be moved at least once

Admissible heuristics

E.g., for the 8-puzzle:

□ $h_2(n)$ = Manhattan distance

(i.e., the sum of the horizontal and vertical distances of the tiles from their goal positions)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

□ $h_2(S) = ?$

□ $h_2(S) = 3+1+2+2+2+3+3+2 = \mathbf{18}$

□ **Admissible:** any move can move at most one tile one step closer to the goal

Dominance

8-puzzle:

h_1 number of misplaced tiles

h_2 Manhattan distance

- If $h_2(n) \geq h_1(n)$ for all node n (both admissible)
then h_2 dominates h_1
- h_2 is better than h_1 for search
 h_1 will expand all the nodes h_2 expands and possibly more
- **Comparison** of the **search costs** (average number of generated nodes)
for **A*** with h_1 , **A*** with h_2 and **IDS** (Iterative Deepening Search)
 - $d=12$ **IDS** = 3,644,035 nodes
 $A^*(h_1) = 227$ nodes
 $A^*(h_2) = 73$ nodes
 - $d=24$ **IDS** = too many nodes
 $A^*(h_1) = 39,135$ nodes
 $A^*(h_2) = 1,641$ nodes

d = solution length

Generating admissible heuristics from relaxed problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- A **superset of actions** is available from **each state**
- The **graph** of the relaxed problem is a **supergraph** of the original
- Any optimal solution of the **original problem** is also a solution of the **relaxed problem**
- The **cost** of an optimal solution of the **relaxed problem** may be the **same or lower than** the **cost** of an optimal solution of the **original problem**



The **cost** of an optimal solution to a **relaxed problem** is an **admissible heuristic** for the **original problem**

Generating admissible heuristics from relaxed problems

Admissible heuristics for the 8-puzzle:

h_1 number of misplaced tiles

h_2 Manhattan distance

- **Admissible heuristics** can be derived from the **exact solution cost** of a **relaxed** version of the problem
- If the **rules** of the **8-puzzle** are **relaxed** so that a tile **can move anywhere** (instead of just to the adjacent **empty square**) then $h_1(n)$ gives the **exact** number of steps of the optimal solution
- If the **rules** of the **8-puzzle** are **relaxed** so that a tile **can move one square in any direction**, even onto an occupied square, then $h_2(n)$ gives the **exact** number of steps of the optimal solution