

INTRODUCTION

Chapter 1


Outline



- What is AI?
- A brief history
- The state of the art

What is AI?

Views of AI fall into four categories:




Thinking humanly	Thinking rationally
Acting humanly	Acting rationally

The textbook advocates "acting rationally"

- Ideally, an **intelligent agent** takes the best possible action in a situation
- We study the problem of building intelligent agents


Acting humanly: Turing Test

- Turing (1950) "Computing machinery and intelligence":
- "Can machines behave intelligently?"
- Operational test for intelligent behavior



Acting humanly: Turing Test

- Turing (1950) "Computing machinery and intelligence":
- "Can machines behave intelligently?"
- Operational test for intelligent behavior



- The test suggested major components of AI:
 - Natural language understanding
 - Knowledge representation
 - Automatic reasoning
 - Learning
 - Computer vision and robotics, if we consider the "total Turing test"

Thinking humanly: cognitive modeling

- 1960s "cognitive revolution": information-processing psychology
- Requires scientific theories (e.g. computational models) of internal activities of the brain
- How to validate? Requires
For ex, predicting and testing behavior of human subjects
- This approach (roughly, **Cognitive Science**) is now distinct from AI

Thinking rationally: "laws of thought"

- **Aristotele**: what are correct reasoning processes?
- **Syllogisms**: patterns that lead from correct premises to correct conclusions
 - Example:**
 - Premises:
 - Socrates is a man;
 - All men are mortal
 - Conclusion:
 - Socrates is mortal
- The study of these laws of thought initiated the field called **logic**
- Some problem:
 - Casting informal knowledge into the formal structure required by the logic (e.g. uncertainty)



Acting rationally: rational agent

- An **agent** is an entity that **perceives** and **acts**
- A **rational agent**: **agent** that **acts** so as **to achieve**
 - ▣ the best outcome or
 - ▣ the best expected outcome when there is uncertainty
- Doesn't necessarily involve thinking – e.g., reflex action – but thinking should be in the service of rational action
- This course is about **designing rational agents**

Rational agents

- For any given class of environments and tasks, **we seek the agent** (or class of agents) **with the best performance**
- This course is about designing **rational agents**
- **Computational limitations** make **perfect rationality unachievable**
 - design **best program** for given machine resources

INTRODUCTION – PART II


Chapter 1


Outline



- What is AI?
- A brief history
- The state of the art

AI prehistory

- 
- Philosophy Logic, methods of reasoning, mind as physical system, foundations of learning, language, rationality
 - Mathematics Formal representation and proof algorithms, computation, (un)decidability, (in)tractability, probability
 - Economics Utility, decision theory
 - Neuroscience How does the brain process information, neural models
 - Psychology Phenomena of perception, action, experimental techniques
 - Computer engineering Building fast computers
 - Control theory Design systems that maximize an objective function over time
 - Linguistics Knowledge representation, grammar



A brief history of AI

- 1943 McCulloch & Pitts: Boolean circuit model of brain
- 1950 Turing's "Computing Machinery and Intelligence"
- 1956 Dartmouth meeting: "Artificial Intelligence" adopted
- 1952—69 Early enthusiasm, great expectations
- 1950s Early AI programs, Newell & Simon's Logic Theorist, Gelernter's Geometry Theorem Prover
- 1965 Robinson's complete algorithm for logical reasoning
- 1966—73 AI discovers computational complexity
Neural network research almost disappears
- 1969—79 Early development of knowledge-based systems
- 1980-- AI becomes an industry
- 1986-- Neural networks return to popularity
- 1987-- AI adopts the scientific method
- 1995-- The emergence of intelligent agents
- 2001-- Big data



State of the art

- **Logistic planning**: During the 1991 Gulf War, US forces deployed an **AI logistics planning and scheduling program** that

- involved up to **50,000** vehicles, cargo, and people
- had to account for starting points, destinations, routes, and conflict resolution among all parameters

The **AI planning techniques** generated in hours a plan that would have taken weeks with older methods

DARPA (The Defense Advanced Research Project Agency) stated that this single application more than **paid back DARPA's 30 year** investment in AI

State of the art



- **Autonomus planning and scheduling:** NASA's on-board **autonomous planning program** controlled the **scheduling** of operations for a spacecraft
- **Robotics:** iRobot Corporation has sold over two million Roomba robotic vacuum cleaners for home use
- **Speech recognition:** A traveler calling United Airlines to book a flight can have the entire conversation guided by an automated speech recognition and dialog management system

State of the art

- **Games:** An **AI** beats the world best human player at GO

<https://www.youtube.com/watch?v=g-dKXOIsf98>

- Self Driving Cars

<https://www.youtube.com/watch?v=7oCe0aLye-U>

- Wigdog entertains conversations with patients with Alzheimer

<http://www.youtube.com/watch?v=-Xx5hgjD-Mw>

- ...