


# Agent Architectures & Hierarchical Control



Gloria Beraldo




Book:

<https://artint.info/3e/html/ArtInt3e.html>

## Chapter 2

- Agent systems and architectures
- Agent controllers
- Hierarchical controllers

# Agents in the environments



- **Abilities**: the set of primitive actions it is capable of carrying out.
- **Goals/Preferences**: what it wants, its desires, its values,...
- **Prior Knowledge**: what it comes into being knowing, what it doesn't get from experience,...
- **History of stimuli**
  1. (current) stimuli: what it receives from environment now (observations, percepts)
  2. past experiences: what it has received in the past

# Examples

## ROBOT


**abilities:** movement, grippers, speech, facial expressions,. . .

**goals:** deliver food, rescue people, score goals, explore,. . .

**prior knowledge:** what is important feature, categories of objects, what a sensor tell us,. . .

**stimuli:** vision, sonar, sound, speech recognition, gesture recognition,. . .

**past experiences:** effect of steering, slipperiness, how people move,. . .



## THERMOSTAT FOR HEATER

**abilities:** turn heater on or off

**goals:** conformable temperature, save fuel, save money

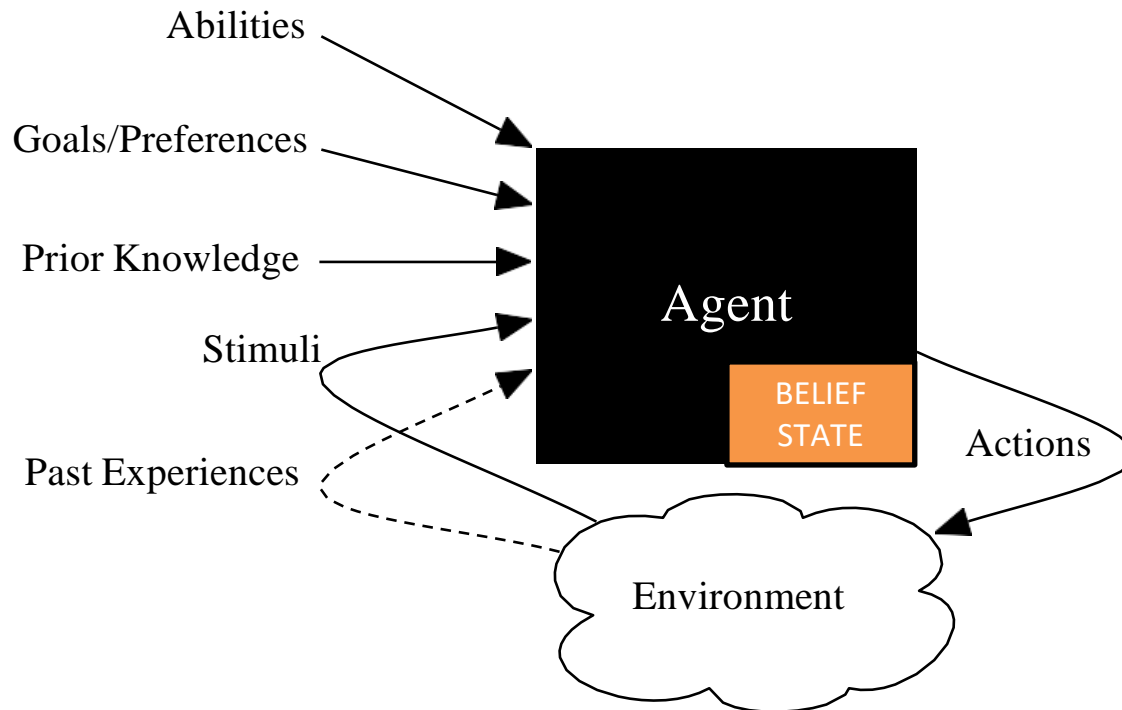
**prior knowledge:** 24 hour cycle, weekends

**stimuli:** temperature, set temperature, who is home, outside temperature

**past experiences:** when people come and go, who likes what temperature



# Agents in the environments




- **Abilities**: the set of primitive actions it is capable of carrying out.
- **Goals/Preferences**: what it wants, its desires, its values,...
- **Prior Knowledge**: what it comes into being knowing, what it doesn't get from experience,...

- **History of stimuli**
  1. (current) stimuli: what it receives from environment now (observations, percepts)
  2. past experiences: what it has received in the past

**BELIEF STATE** that encodes beliefs about its environment and it is used to **choose actions**

# The Agent & its components

An agent interacts with the environment through its **body**.



An agent receives **stimuli** from the environment

An agent carries out **actions** in the environment.

An **agent** is made up of a **body** and a **controller**.

The body is made up of:


- **sensors** that interpret stimuli
- **actuators** that convert commands into actions

The controller receives **percepts** from the body.

The controller sends **commands** to the body.

The body can also have reactions that are not controlled (e.g., after emergency button, reflexes) .

# Implementing a controller



Agents are situated in time, they receive sensory data in time, and do actions in time.

Controllers have (limited) memory and (limited) computational capabilities.

The controller specifies the command at every time.

The command at any time can depend on the current and previous percepts.

# The Agent functions

A **percept trace** is a sequence of all past, present, and future percepts received by the controller.

A **command trace** is a sequence of all past, present, and future commands output by the controller.

A **transduction** is a function from percept traces into command traces.

A transduction is **causal** if the command trace up to time  $t$  depends only on percepts up to  $t$ .

An **agent's history** at time  $t$  is sequence of past and present percepts and past commands.

A **controller** is an implementation of a causal transduction that specifies a function from an agent's history at time  $t$  into its action at time  $t$ .



An agent doesn't have access to its entire history. It only has access to what it has remembered.

The **memory** or **belief state** of an agent at time  $t$  encodes all of the agent's history that it has access to.

The belief state of an agent encapsulates the information about its past that it can use for current and future actions.

At every time a controller has to decide on:


What should it do?

What should it remember?

(How should it update its memory?)

as a function of its percepts and its memory.

# Functions in a controller



For discrete time, a controller implements:

- **belief state function**  
*remember(belief state, percept)*  
returns the next belief state.
- **command function**  
*command(memory, percept)*  
returns the command for the agent.

You don't need to implement an intelligent agent as:



as three independent modules, each feeding into the the next.

- It's too slow.
- High-level strategic reasoning takes **more time than the reaction time** needed (e.g. to avoid obstacles).
- The output of the perception depends on what you will do with it.

**Complex agents are built modular in terms of  
interacting hierarchical layers**

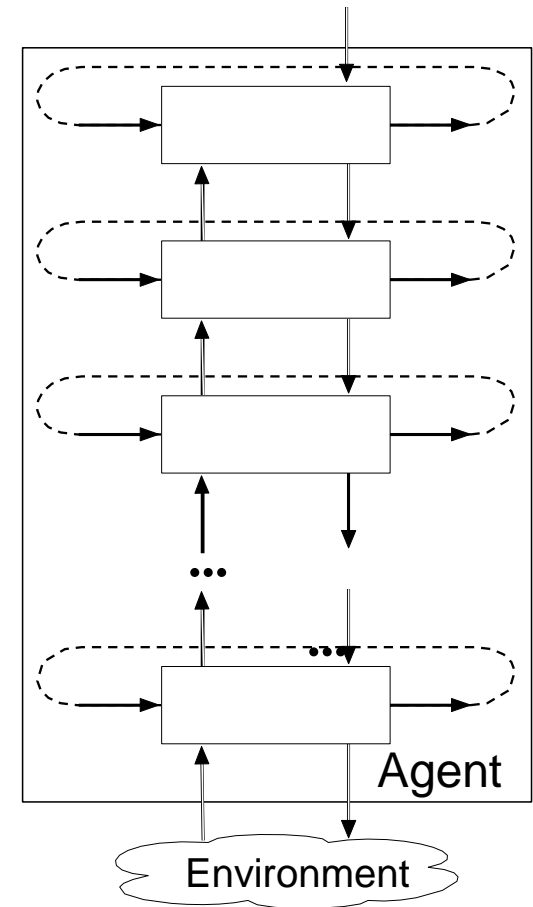

# Hierarchical Control

A better architecture is a **hierarchy** of controllers.


Each controller sees the controllers below it as a virtual body from which it gets percepts and sends commands.

The lower-level controllers can

- **run much faster**, and react to the world more quickly
- deliver a **simpler view of the world** to the higher-level controllers.



# Functions per layer



memory function

*remember(memory, percept, command)*

command function

*do(memory, percept, command)*

percept function

*higher percept(memory, percept, command)*

Two levels of abstraction:

- The **knowledge level** is in terms of what an agent knows and what its goals are.
- The **symbol level** is a level of description of an agent in terms of what reasoning it is doing.

The knowledge level is about the external world to the agent.

The symbol level is about what symbols an agent uses to implement the knowledge level.

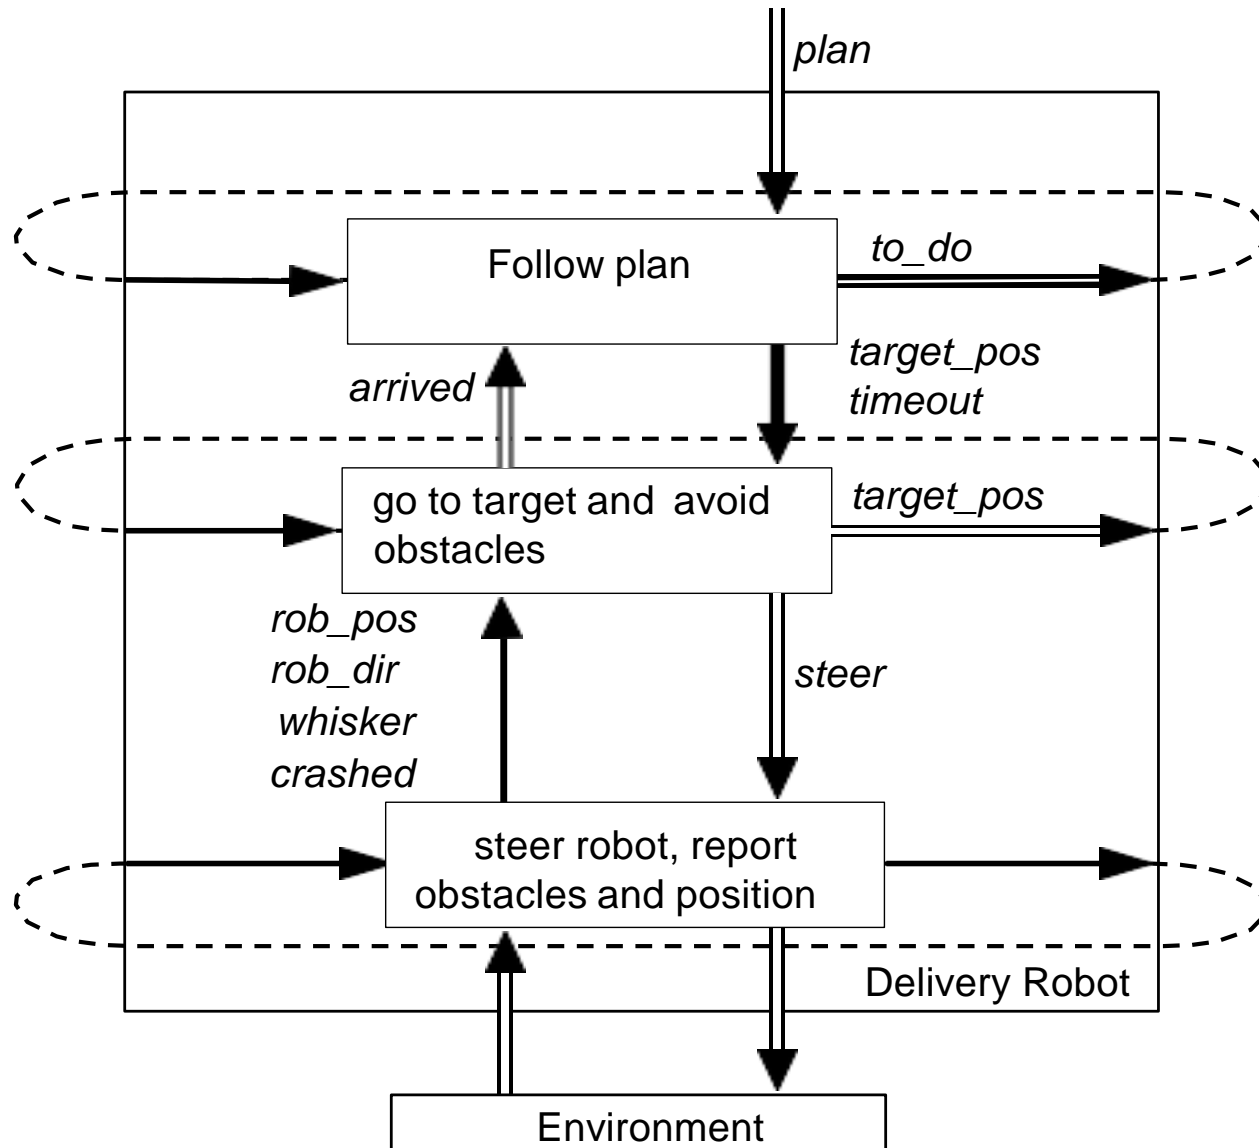
- A high-level description is **easier for a human** to specify and understand.
- A low-level description can be **more accurate and more predictive**. High-level descriptions abstract away details that may be important for actually solving the problem.
- The lower the level, the **more difficult it is to reason** with. You may not know the information needed for a low-level description.
- It is sometime possible to use multiple levels of abstraction.

# Example: delivery robot

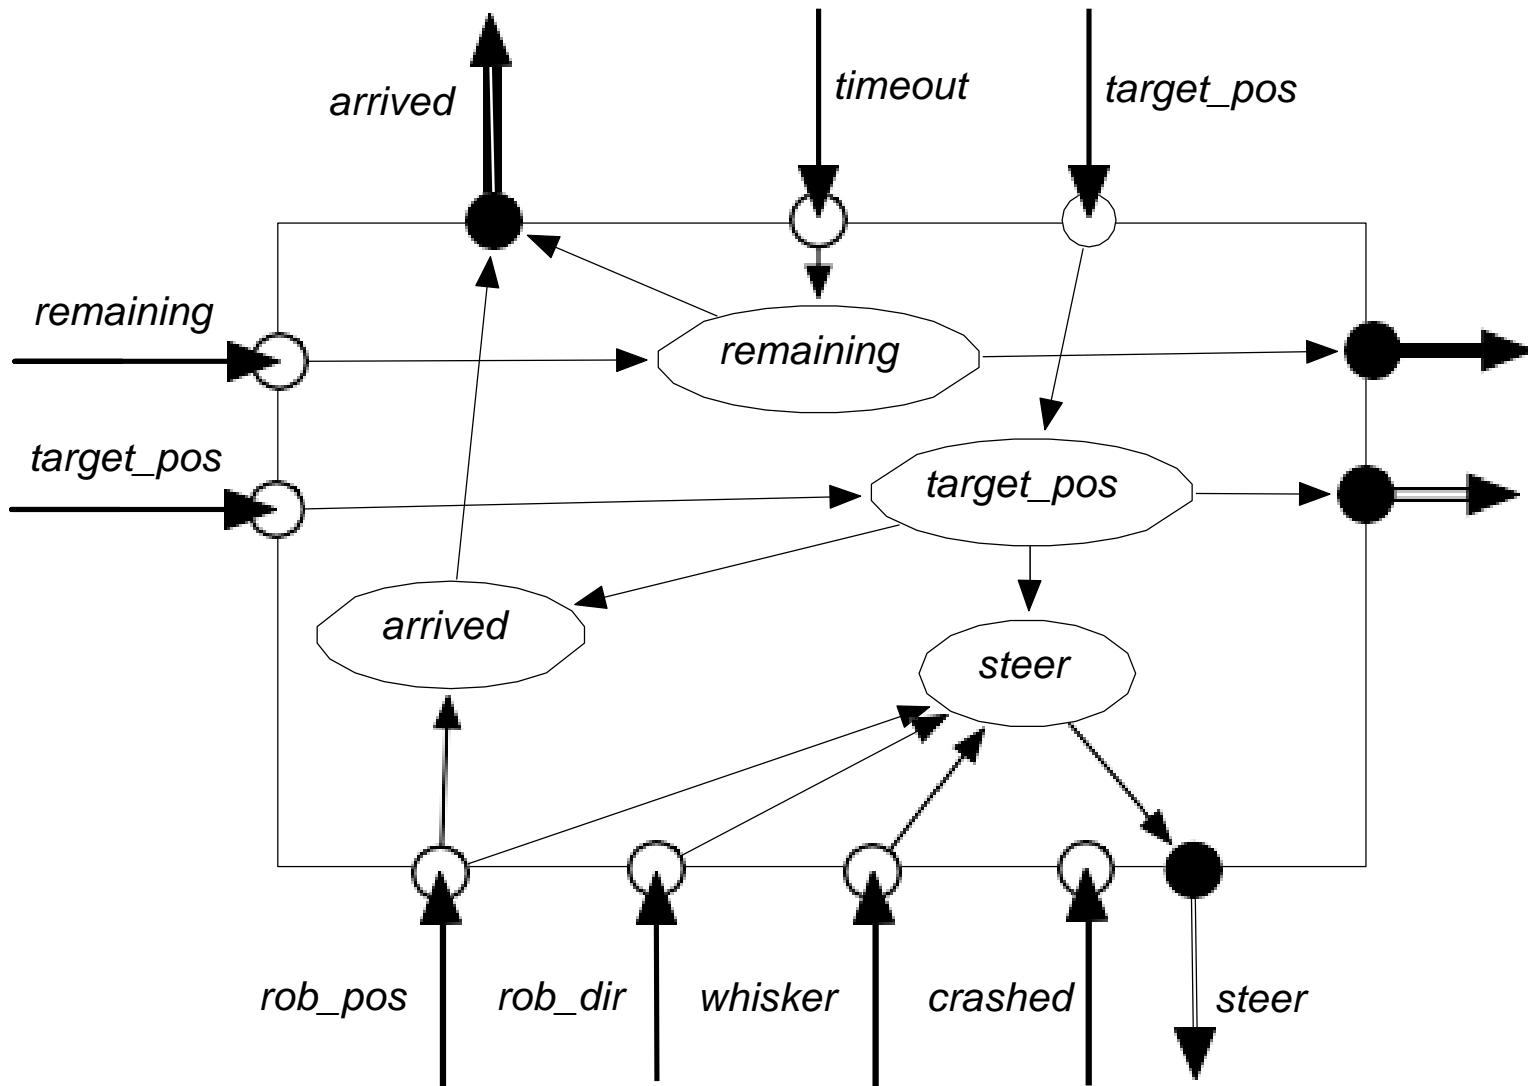
- The robot has **three actions**: go straight, go right, go left. (Its velocity doesn't change).
- It can be given a **plan** consisting of sequence of named locations for the robot to go to in turn.
- The robot **must avoid obstacles**. It has a single whisker sensor pointing forward and to the right. The robot can detect if the whisker hits an object. The robot knows where it is.
- The obstacles and locations **can be moved dynamically**. Obstacles and new locations can be created dynamically.



# Delivery robot architecture



# Middle layer of the delivery robot



given *timeout* and *target pos*:

*remaining* := *timeout*

while not *arrived()* and *remaining*  $\neq$  0

if *whisker\_sensor* = on

then *steer* := left

else if *straight\_ahead*(*rob\_pos*, *robot\_dir*, *target\_pos*)

then *steer* := straight

else if *left\_of*(*rob\_pos*, *robot\_dir*, *target\_pos*)

then *steer* := left

else *steer* := right

*lower.do(steer)*

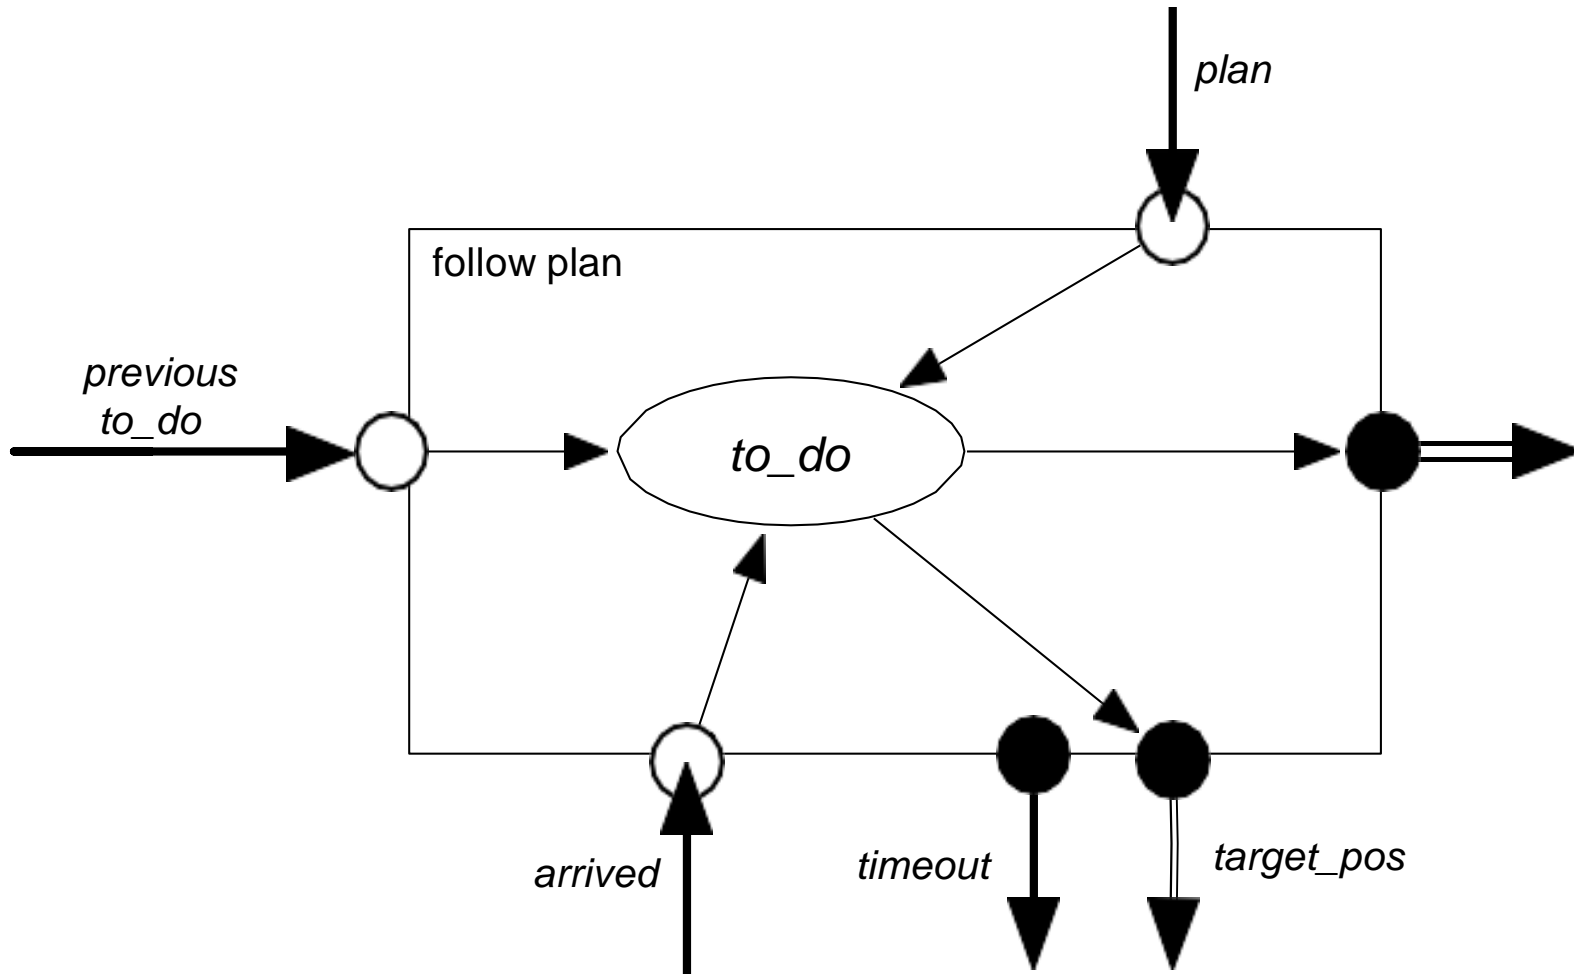
*remaining* := *remaining* - 1

tell upper layer *arrived()*

# Top layer of the delivery robot

- The top layer is given a **plan** which is a sequence of named locations.
- The top layer tells the middle layer **the goal position** of the current location.
- It has to **remember the current goal** position and the **locations still to visit**.
- When the middle layer reports the robot has arrived, **the top layer takes the next location from the list of positions** to visit, and there is a new goal position.

# Top layer of the delivery robot



# Top layer of the delivery robot

given *plan*:

*to\_do* := *plan*

*timeout* := 200

while not *empty*(*to\_do*)

*target\_pos* := *coordinates*(*first*(*to\_do*))

*middle.do*(*timeout*, *target\_pos*)

*to\_do* := *rest*(*to\_do*)

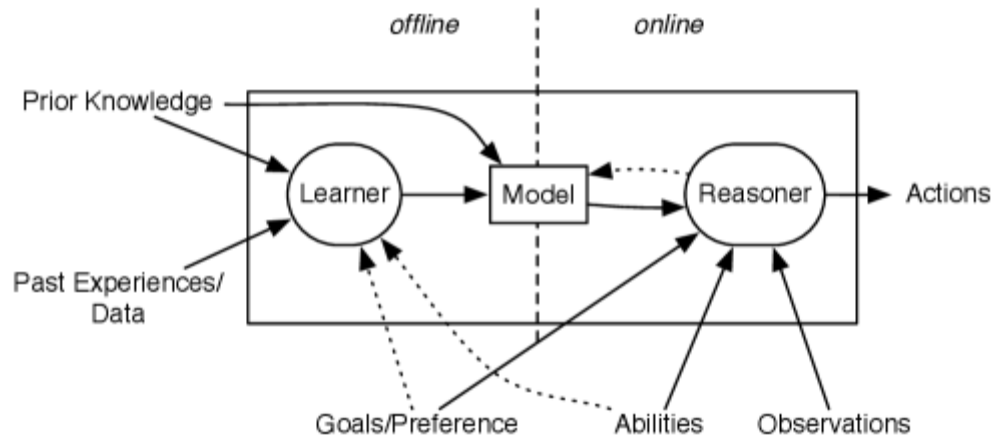
*to\_do* list is the first location

*rest*(*to\_do*) the rest of the *to\_do* list

# What should be in an agent's belief state?

- An agent decides what to do **based on its belief state** and **what it observes**.
- A **purely reactive** agent doesn't have a belief state (i.e. action based on its perception).
- A **dead reckoning** agent doesn't perceive the world (i.e., it exploits forward prediction for the events in the words).  
— neither work very well in complicated domains (e.g., affected by noise, not deterministic situations).
- It is often useful for the agent's belief state to be a model of the world (itself and the environment).
- The next belief state can be estimated using Bayes's rule (we will see in the next lectures)

# Offline and Online Computation



An intelligent agent requires knowledge that is acquired at the design time, offline or online.

**Offline:** before the agent has to act, the agent uses prior knowledge and past experiences (either its own past experiences or data it has been given) to learn (parts of) a model that is useful for acting online.

**Online:** when the agent is acting, the agent uses its model, its observations of the world, and its goals and abilities to choose what to do and update its belief state. Online, the information about the particular situation becomes available, and the agent has to act.

An agent typically has **much more time for offline computation** than for online computation. During online computation it can **take advantage of particular goals and particular observations**.