

Machine Learning

Regularization and Feature Selection

Fabio Vandin

November 27th, 2023

Ridge Regression: Matrix Form

Linear regression: pick

$$\arg \min_{\mathbf{w}} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

Ridge regression: pick

$$\arg \min_{\mathbf{w}} \left(\lambda \|\mathbf{w}\|^2 + (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \right)$$

Want to find \mathbf{w} which minimizes

$$f(\mathbf{w}) = \lambda \|\mathbf{w}\|^2 + (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}).$$

How? $2\lambda \vec{w}$ $-2 \mathbf{X}^T (\vec{y} - \mathbf{X} \vec{w})$

gradient

Compute gradient $\frac{\partial f(\mathbf{w})}{\partial \mathbf{w}}$ of objective function w.r.t \mathbf{w} and compare it to 0.

Want to find \mathbf{w} which minimizes

$$f(\mathbf{w}) = \lambda \|\mathbf{w}\|^2 + (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}).$$

How?

Compute gradient $\frac{\partial f(\mathbf{w})}{\partial \mathbf{w}}$ of objective function w.r.t \mathbf{w} and compare it to 0.

$$\frac{\partial f(\mathbf{w})}{\partial \mathbf{w}} = 2\lambda \mathbf{w} - 2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

Then we need to find \mathbf{w} such that

$$2\lambda \mathbf{w} - 2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

$$\lambda \vec{w} - \mathbf{X}^T(\vec{y} - \mathbf{X}\vec{w}) = 0$$

$$\lambda \vec{w} + \mathbf{X}^T \mathbf{X} \vec{w} = \mathbf{X}^T \vec{y}$$

$$(\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X}) \vec{w} = \mathbf{X}^T \vec{y} \Rightarrow \vec{w} = (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \vec{y}$$

$$2\lambda \mathbf{w} - 2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

is equivalent to

$$(\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

Note:

- $\mathbf{X}^T \mathbf{X}$ is positive semidefinite
- $\lambda \mathbf{I}$ is positive definite

$\Rightarrow \lambda \mathbf{I} + \mathbf{X}^T \mathbf{X}$ is positive definite

$\Rightarrow \lambda \mathbf{I} + \mathbf{X}^T \mathbf{X}$ is invertible

Ridge regression solution:


$$\mathbf{w} = (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Exercise 5

Consider the ridge regression problem

$\arg \min_{\mathbf{w}} \lambda \|\mathbf{w}\|^2 + \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$. Let: h_S be the hypothesis obtained by ridge regression on with training set S ; h^* be the hypothesis of minimum generalization error among all linear models.

- (A) Draw, in the plot below, a *typical* behaviour of (i) *the training error* and (ii) *the test/generalization error* of h_S as a function of λ .
- (B) Draw, in the plot below, a *typical* behaviour of (i) $L_{\mathcal{D}}(h_S) - L_{\mathcal{D}}(h^*)$ and (ii) $L_{\mathcal{D}}(h_S) - L_S(h_S)$ as a function of λ .



Feature Selection

In general, in machine learning one has to decide what to use as features (= input) for learning.

Even if somebody gives us a representation as a feature vector, maybe there is a “better” representation?

What is “better”?

Example

- features x_1, x_2 , output y
- $x_1 \sim \text{Uniform}(-1, 1)$
- $y = x_1^2$
- $x_2 \sim y + \text{Uniform}(-0.01, 0.01)$

If we want to predict y , which feature is better: x_1 or x_2 ?

- x_1 : because with $(x_1)^2$ we perfectly predict y .
But what if you use x_2 as feature of a linear model?
Prediction is not great!
- x_2 , because with linear models we predict y exactly up to noise ($\text{Uniform}(-0.01, 0.01)$). But we could do better by looking at x_1^2 .

Example

- features x_1, x_2 , output y
- $x_1 \sim \text{Uniform}(-1, 1)$
- $y = x_1^2$
- $x_2 \sim y + \text{Uniform}(-0.01, 0.01)$

If we want to predict y , which feature is better: x_1 or x_2 ?

No-free lunch...

Feature Selection: Scenario

We have a large pool of features

Goal: select a small number of features that will be used by our (final) predictor

Assume $\mathcal{X} = \mathbb{R}^d$.

Goal: learn (final) predictor using $k \ll d$ predictors ^(features)

Feature Selection: Scenario

We have a large pool of features

Goal: select a small number of features that will be used by our (final) predictor

Assume $\mathcal{X} = \mathbb{R}^d$.

Goal: learn (final) predictor using $k \ll d$ predictors

Motivation?

- prevent overfitting: less predictors \Rightarrow hypotheses of lower complexity!
- predictions can be done faster
- useful in many applications!

Feature Selection: Computational Problem

Assume that we use the Empirical Risk Minimization (ERM) procedure.

Assumption: an hypothesis $h \in \mathcal{H}$ corresponds to a vector of weights $\mathbf{w} \in \mathbb{R}^d$.
The problem of selecting k features that minimize the empirical risk can be written as:

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \text{ subject to } \|\mathbf{w}\|_0 \leq k$$

where $\|\mathbf{w}\|_0 = |\{i : w_i \neq 0\}|$

Feature Selection: Computational Problem

Assume that we use the Empirical Risk Minimization (ERM) procedure.

The problem of selecting k features that minimize the empirical risk can be written as:

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \quad \text{subject to} \quad \|\mathbf{w}\|_0 \leq k$$

where $\|\mathbf{w}\|_0 = |\{i : w_i \neq 0\}|$

How can we solve it?

Subset Selection

How do we find the solution to the problem below?

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \quad \text{subject to} \quad \|\mathbf{w}\|_0 \leq k$$

Note: the solution will always include k features

Let:

- $\mathcal{I} = \{1, \dots, d\}$;
- given $p = \{i_1, \dots, i_k\} \subseteq \mathcal{I}$: \mathcal{H}_p = hypotheses/models where only features $w_{i_1}, w_{i_2}, \dots, w_{i_k}$ are used

$$P^{(k)} \leftarrow \{J \subseteq \mathcal{I} : |J| = k\};$$

foreach $p \in P^{(k)}$ **do**

$$\quad \left[\quad h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h); \right]$$

return $p_k \leftarrow \arg \min_{p \in P^{(k)}} L_S(h_p);$

Complexity? Learn $\Theta\left(\binom{d}{k}\right) \in \Theta(d^k)$ models \Rightarrow exponential algorithm!

Can we do better?

*≈ it is unlikely that there is
a poly-time alg. to solve
it.*

Proposition

The optimization problem of feature selection NP-hard.

Can we do better?

Proposition

The optimization problem of feature selection NP-hard.

What can we do?

Heuristic solution \Rightarrow greedy algorithms

Greedy Algorithms for Feature Selection

Forward Selection: start from the empty solution, add one feature at the time, until solution has cardinality k

```
 $sol \leftarrow \emptyset;$   
while  $|sol| < k$  do  
  foreach  $i \in \mathcal{I} \setminus sol$  do  
     $p \leftarrow sol \cup \{i\};$   
     $h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h);$   
   $sol \leftarrow sol \cup \arg \min_{i \in \mathcal{I} \setminus sol} L_S(h_{sol \cup \{i\}});$   
return  $sol;$ 
```

Complexity? Learns $\Theta(kd)$ models

Backward Selection: start from the solution which includes all features, remove one feature at the time, until solution has cardinality k

Pseudocode: analogous to forward selection [Exercise!]

Complexity? Learns $\Theta((d - k)d)$ models

Notes

We have used only training set to select the best hypothesis...

⇒ we may overfit!

Solution? Use validation! (or cross-validation)

Split data into training data and validation data, learn models on training, evaluate (= pick among different hypothesis models) on validation data. Algorithms are similar.

Note: now the best model (in terms of validation error) may include less than k features!

Subset Selection with Validation Data

S = training data (from data split)

V = validation data (from data split)

Using training and validation:

```
for  $\ell \leftarrow 0$  to  $k$  do  
   $P^{(\ell)} \leftarrow \{J \subseteq \mathcal{I} : |J| = \ell\};$   
  foreach  $p \in P^{(\ell)}$  do  
     $h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h);$   
   $p_\ell \leftarrow \arg \min_{p \in P^{(\ell)}} L_V(h_p);$   
return  $\arg \min_{p \in \{p_0, p_1, \dots, p_k\}} L_V(h_p)$ 
```

Forward Selection with Validation Data

Using training and validation:

```
sol  $\leftarrow \emptyset$ ;  
while  $|sol| < k$  do  
  foreach  $i \in \mathcal{I} \setminus sol$  do  
     $p \leftarrow sol \cup \{i\}$ ;  
     $h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h)$ ;  
   $sol \leftarrow sol \cup \arg \min_{i \in \mathcal{I} \setminus sol} L_V(h_{sol \cup \{i\}})$ ;  
return sol;
```

Backward Selection with validation: similar [Exercise]

Similar approach for all algorithms with cross-validation [Exercise]

Bibliography [UML]

Regularization and Ridge Regression: Chapter 12

- no Section 13.3;
- Section 13.4 only up to Corollary 13.8 (excluded)

Feature Selection and LASSO: Chapter 25

- only Section 25.1.2 (introduction and “Backward Elimination”) and 25.1.3