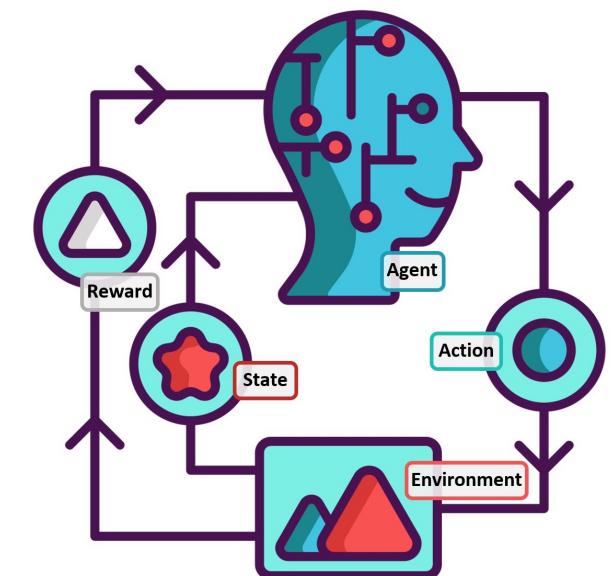


Lecture #12

n-step Bootstrapping

& TD(λ)

Gian Antonio Susto



MC vs. TD(0)

MC: no bias

TD(0): low variance

MC uses real return

TD(0) just rewards and
bootstrap

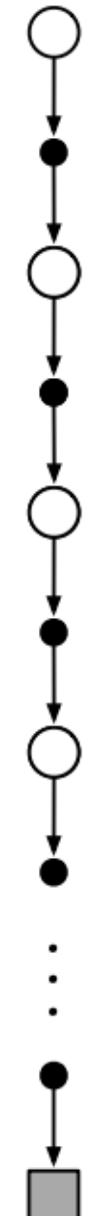
MC updates just at the
end of the episode

TD(0) updates along
the way

1-step TD
and TD(0)

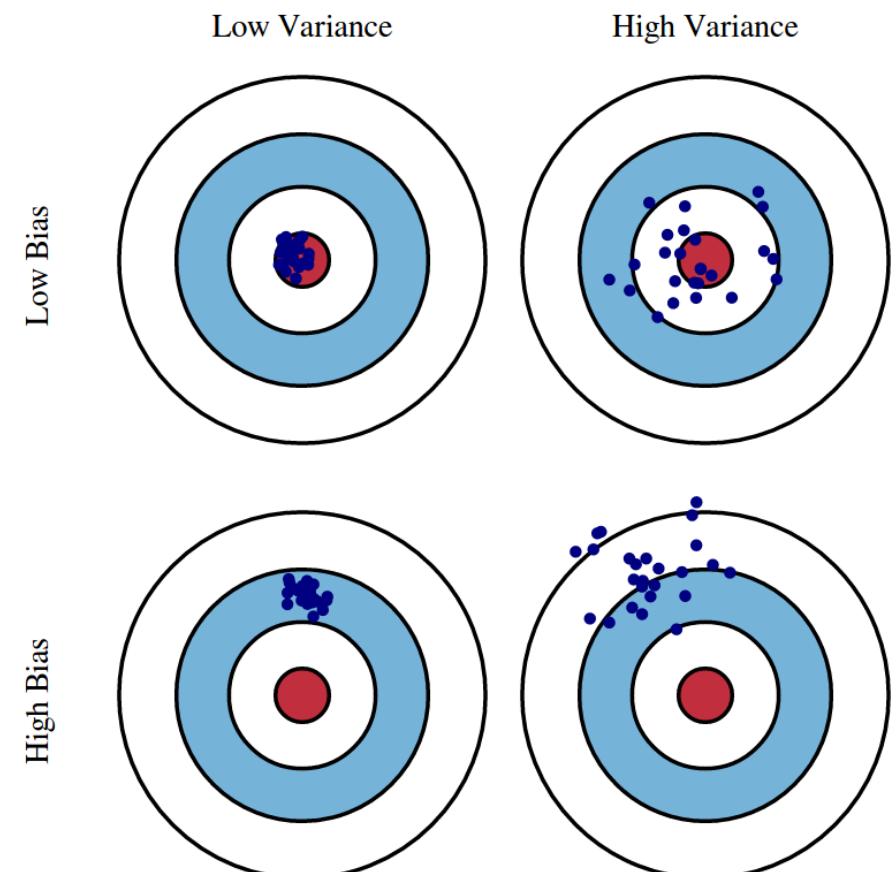
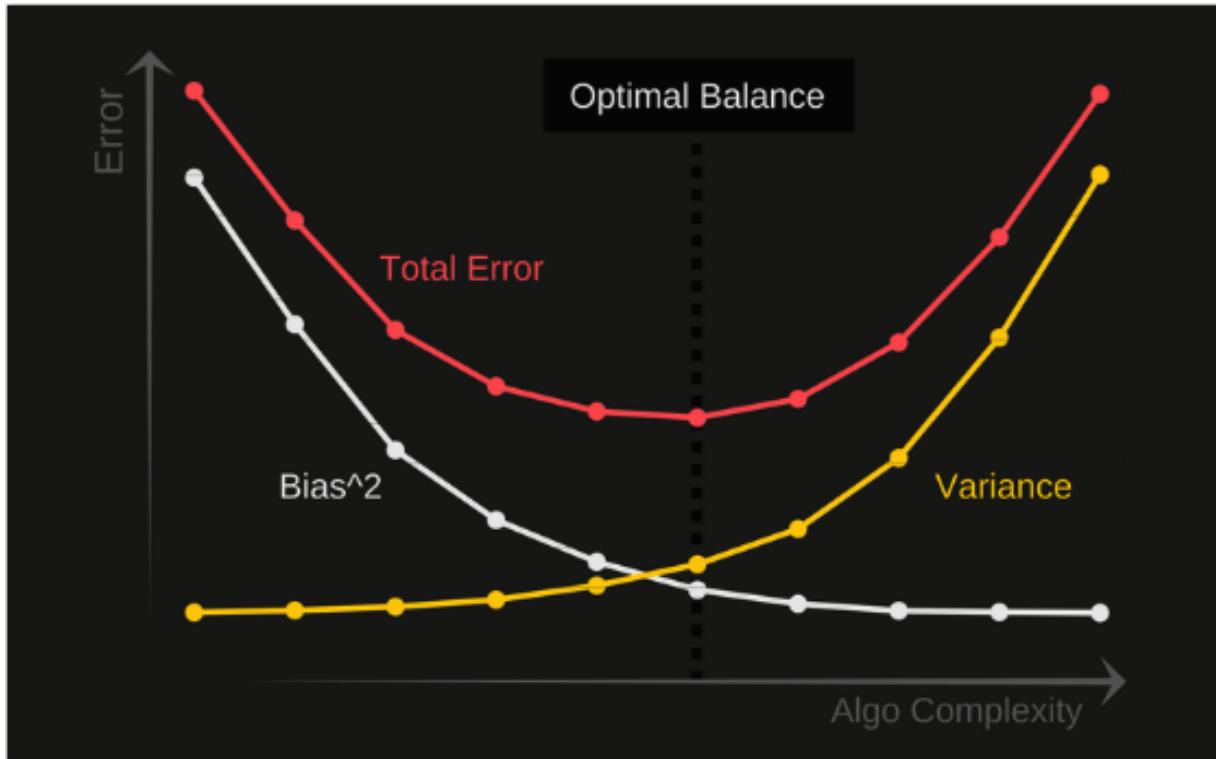


Monte Carlo



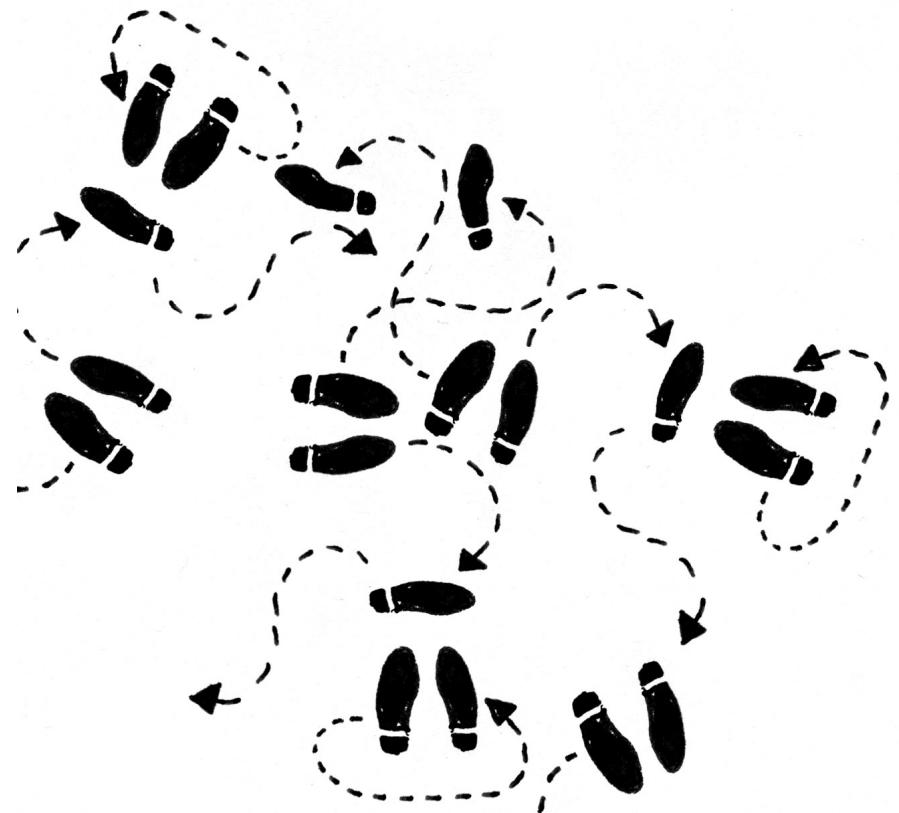
Bias/Variance Trade-off

- A trade-off present in many Machine Learning settings

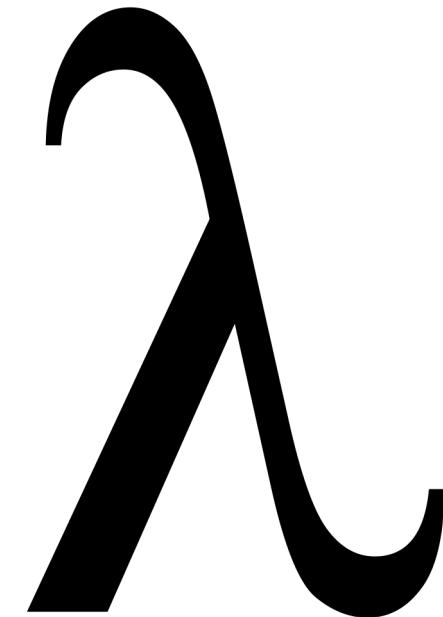


Approaches to balance Bias/Variance

n-step bootstrapping -
chapter 7

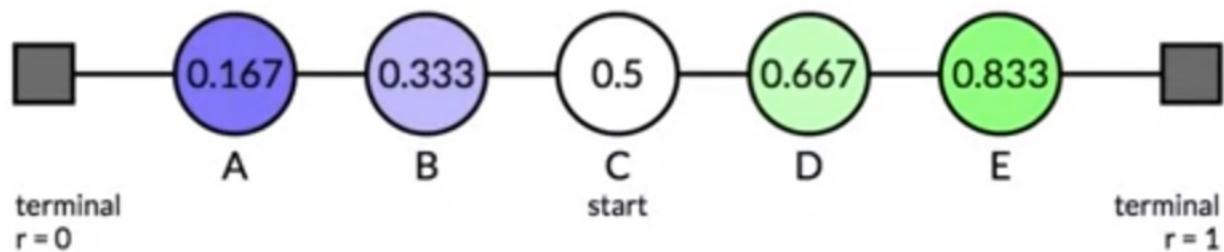


TD(λ) - chapter 12



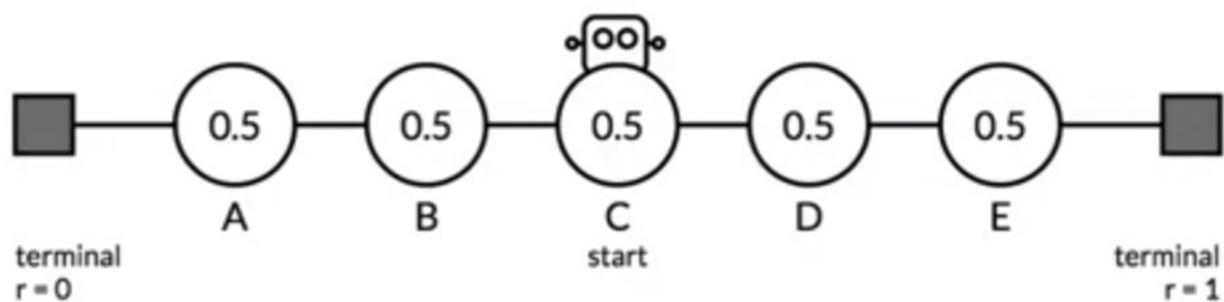
n -step Bootstrapping (Chapter 7)





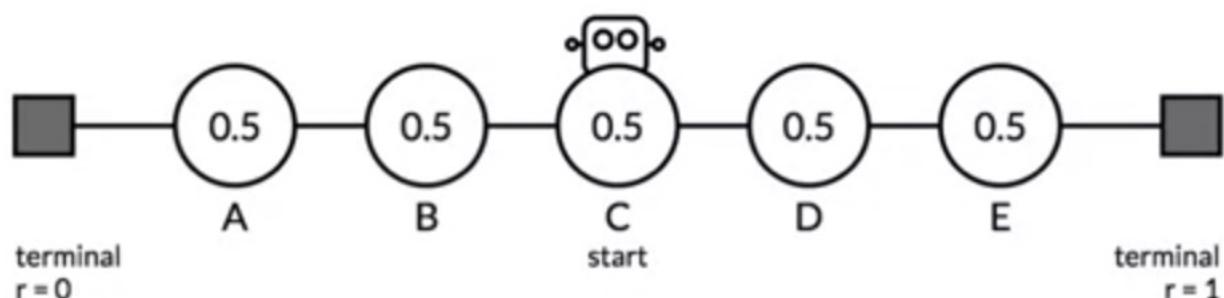
Updates using TD Learning

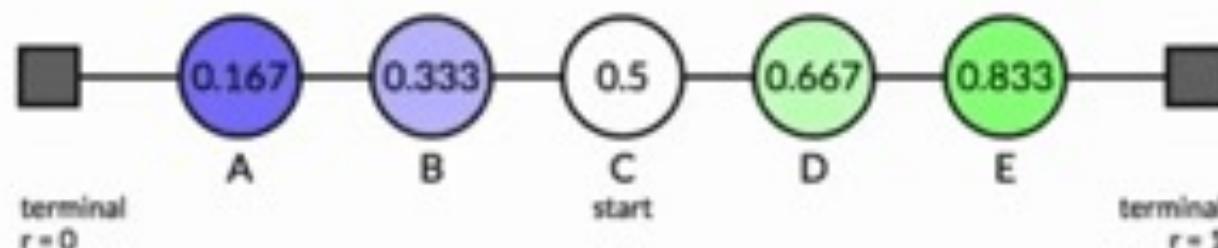
$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$



In the following:
 $\alpha = 0.5$

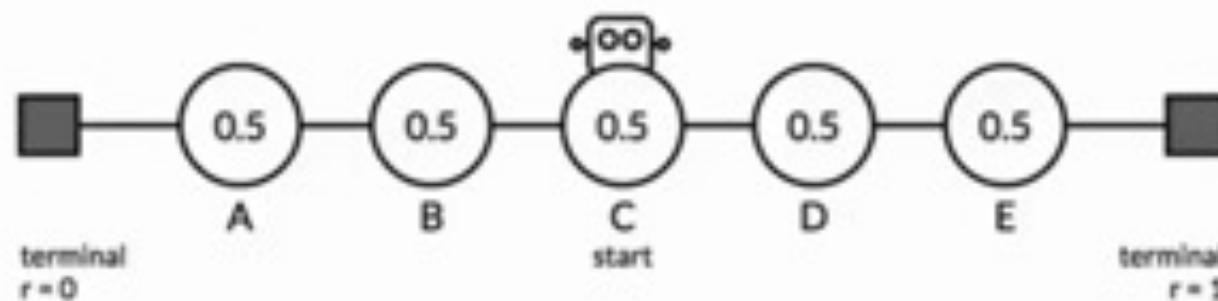
Updates using Monte Carlo





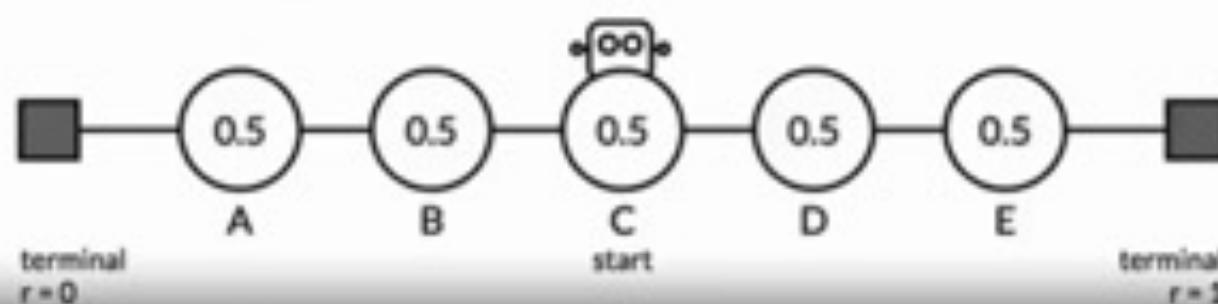
Updates using TD Learning

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$



Updates using Monte Carlo

↳



Prediction: MC vs. TD(0)

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

MC: $V(S_t) \leftarrow V(S_t) + \alpha (\textcolor{red}{G_t} - V(S_t))$

TD(0): $V(S_t) \leftarrow V(S_t) + \alpha (\textcolor{red}{R_{t+1}} + \gamma V(S_{t+1}) - V(S_t))$

The TD target can be seen as the ‘best’ estimation of G_t at time t+1:

$$G_t = R_{t+1} + \gamma G(S_{t+1}) \sim R_{t+1} + \gamma V(S_{t+1})$$

Prediction: why not considering a different estimation of G , for example after 2 steps?

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

$$n = 1$$

$$G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2$$

$$G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

Prediction: why not considering a different estimation of G , for example after 2 steps?

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

This is the best estimation
of G at time $t+2$!

$$n = 1$$

$$G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2$$

$$G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

Prediction: why not considering a different estimation of G , for example after 2 steps?

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

This is the best estimation
of G at time $t+2$!

$$n = 1$$

$$G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2$$

$$G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

I can use this quantity as a
different *TD target*

Prediction: why not considering a different estimation of G , for example after 2 steps?

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots$$

This is the best estimation
of G at time $t+2$!

$n = 1$

$$G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$n = 2$

$$G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

Pay attention
to the notation!

I can use this quantity as a
different *TD target*

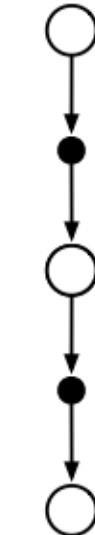
MC vs. TD(0)

MC: low bias

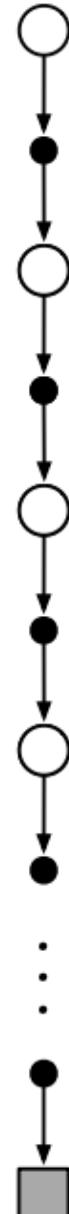
TD(0): low variance

Why don't using a
compromise between
MC and TD(0)?

1-step TD and TD(0)	2-step TD
------------------------	-----------



Monte Carlo



MC vs. TD(0)

MC: low bias

TD(0): low variance

Why don't using a compromise between MC and TD(0)?

We can generalize by considering n-step TD!

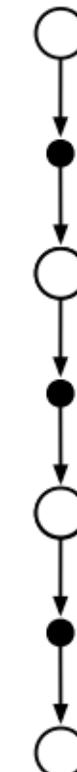
1-step TD and TD(0)



2-step TD



3-step TD



n-step TD



Monte Carlo



MC vs. TD(0)

MC: low bias

TD(0): low variance

Why don't using a
compromise between
MC and TD(0)?

We can generalize by
considering n-step TD!

1-step TD
and TD(0)



2-step TD



3-step TD



n-step TD

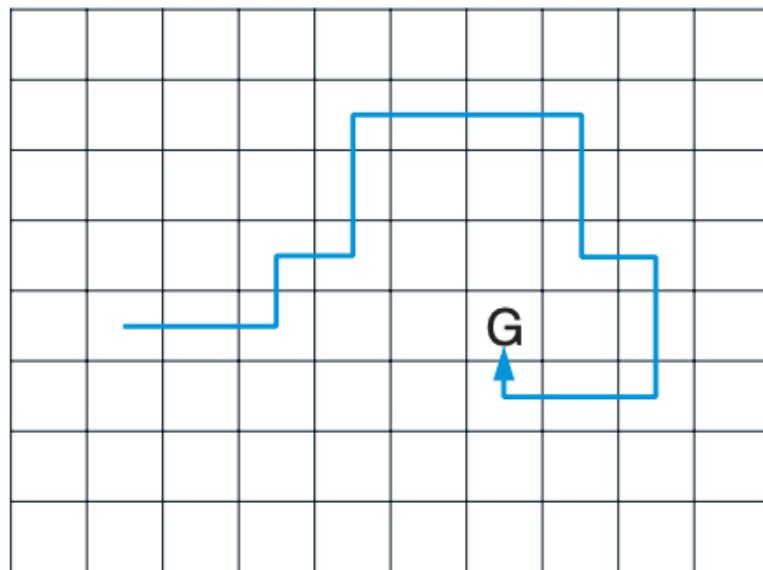


∞ -step TD
and Monte Carlo

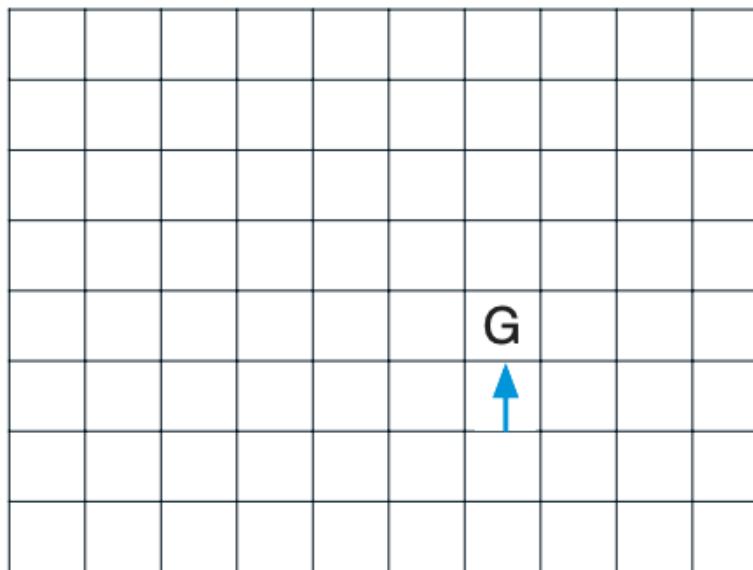


n-step TD learning (Prediction or Control)

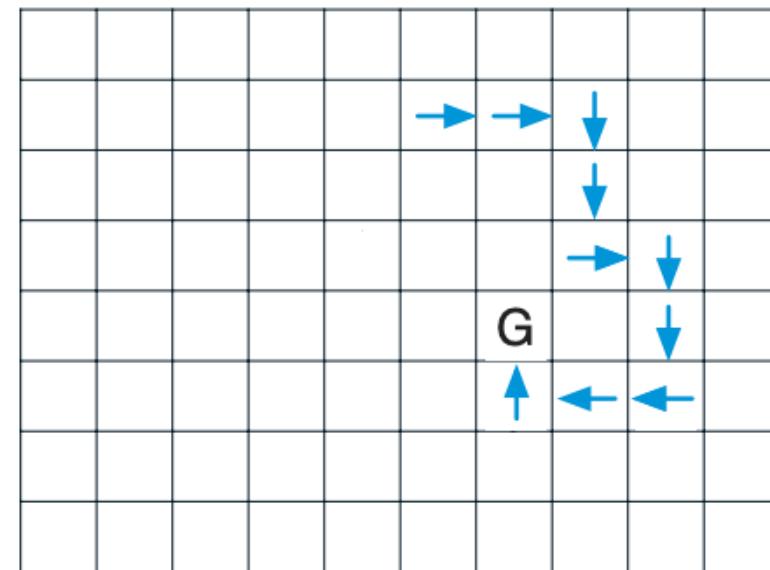
Path taken



Action values increased by one-step Sarsa



Action values increased by 10-step Sarsa



- The one-step method strengthens only the last action of the sequence of actions that led to the high reward
 - The n-step method strengthens the last n actions of the sequence, so much more is learned from the one episode.

Prediction: n -step return

Let's define the n -step returns:

$$n = 1 \quad (\text{TD}) \quad G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2 \quad G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

⋮

⋮

$$n = \infty \quad (\text{MC}) \quad G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$$

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$

A terminal state

Prediction: n -step return

Let's define the n -step returns:

$$n = 1 \quad (\text{TD}) \quad G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2 \quad G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

⋮

⋮

$$n = \infty \quad (\text{MC}) \quad G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$$

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$

n -step TD learning will consider the following updates:

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)]$$

Prediction: n -step return

Let's define the n -step returns:

$$n = 1 \quad (\text{TD}) \quad G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2 \quad G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

⋮

⋮

$$n = \infty \quad (\text{MC}) \quad G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$$

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$

n -step TD learning will consider the following updates:

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)]$$

Where this is not defined?

Prediction: n -step return

Let's define the n -step returns:

$$n = 1 \quad (\text{TD}) \quad G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$$

$$n = 2 \quad G_{t:t+2} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

:

:

$$n = \infty \quad (\text{MC}) \quad G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$$

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$

n -step TD learning will consider the following updates:

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)]$$

Where this is not defined?
When there are less than
 n step before the terminal
state!

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can take their index mod $n + 1$

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$

$$T \uparrow \infty$$

Loop for $t = 0, 1, 2, \dots$:

If $t < T$, then:

Take an action according to $\pi(\cdot|S_t)$

Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n,T)} \gamma^{i-\tau-1} R_i$$

If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$

$$(G_{\tau:\tau+n})$$

$$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$$

Until $\tau = T - 1$

A fixed value for n

n-step TD for estimating $V \approx v_\pi$

twill be our step counter inside the episode

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positiv

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$

$$T \leftarrow \infty$$

Loop for $t = 0, 1, 2, \dots$:

If $t < T$, then:

Take an action according to $\pi(\cdot|S_t)$

Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n,T)} \gamma^{i-\tau-1} R_i$$

If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$

$$(G_{\tau:\tau+n})$$

$$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$$

Until $\tau = T - 1$

n -step TD for estimating $V \approx v_\pi$

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can

Loop for each episode:

Initialize and store $S_0 \neq$ terminal

$T \leftarrow \infty$

Loop for $t = 0, 1, 2, \dots :$

If $t < T$, then:

Take an action according to $\pi(\cdot | S_t)$

Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$

$$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$$

t will be our step counter inside the episode

We will use:

- τ for starting the updates only when n steps are actually seen in the trajectory
- T for checking if we have arrived to close to a terminal state (and therefore a different update must be given)

We need to have at least n steps to do our updates: if we are too close to a terminal state we can't update anymore

Until $\tau = T - 1$

n -step TD for estimating $V \approx v_\pi$

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can

Loop for each episode:

 Initialize and store $S_0 \neq$ terminal

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots :$

 If $t < T$, then:

 Take an action according to $\pi(\cdot | S_t)$

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

 If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$

$$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$$

 Until $\tau = T - 1$

t will be our step counter inside the episode
We will use:

- τ for starting the updates only when n steps are actually seen in the trajectory
- T for checking if we have arrived to close to a terminal state (and therefore a different update must be given)

We need to have at least n steps to do our updates: if we are too close to a terminal state we can't update anymore

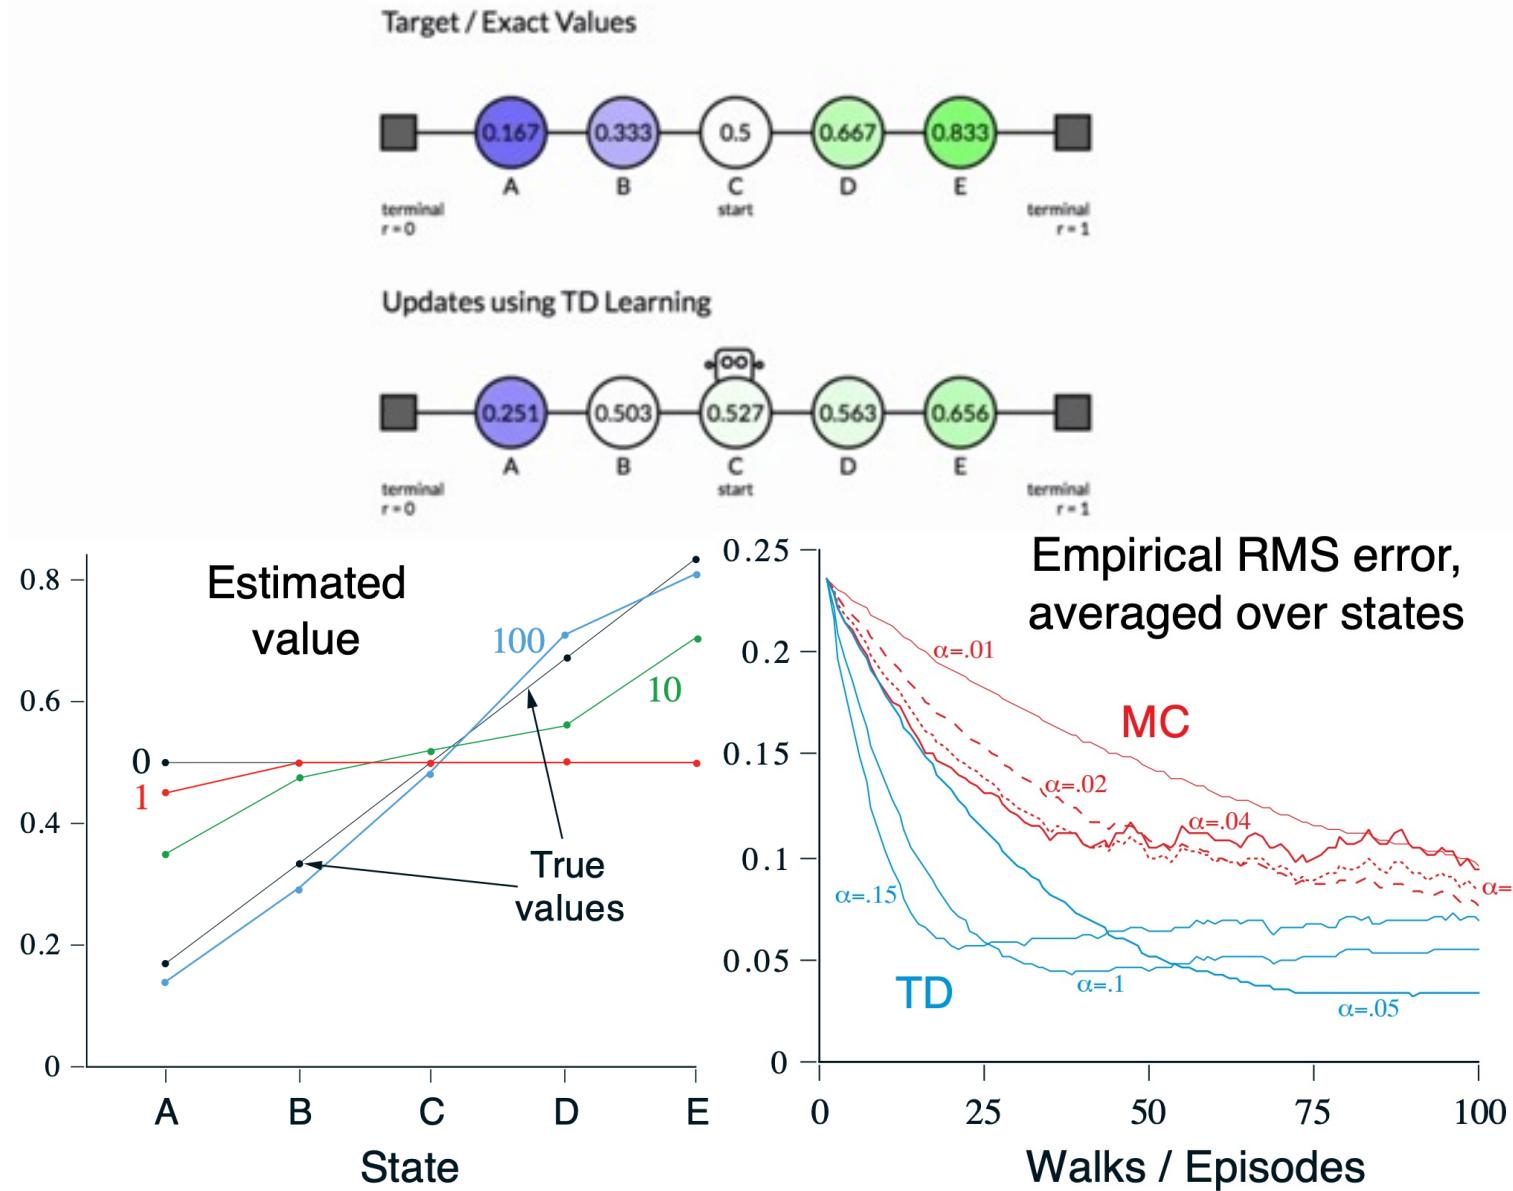
$$(G_{\tau:\tau+n})$$

Policy-evaluation

Prediction: n -step return – example

What is the ‘best’ n ? It depends!

Let’s consider the random walk process example (we are evaluating a random uniform policy): instead of 5 states, we consider 19 states!

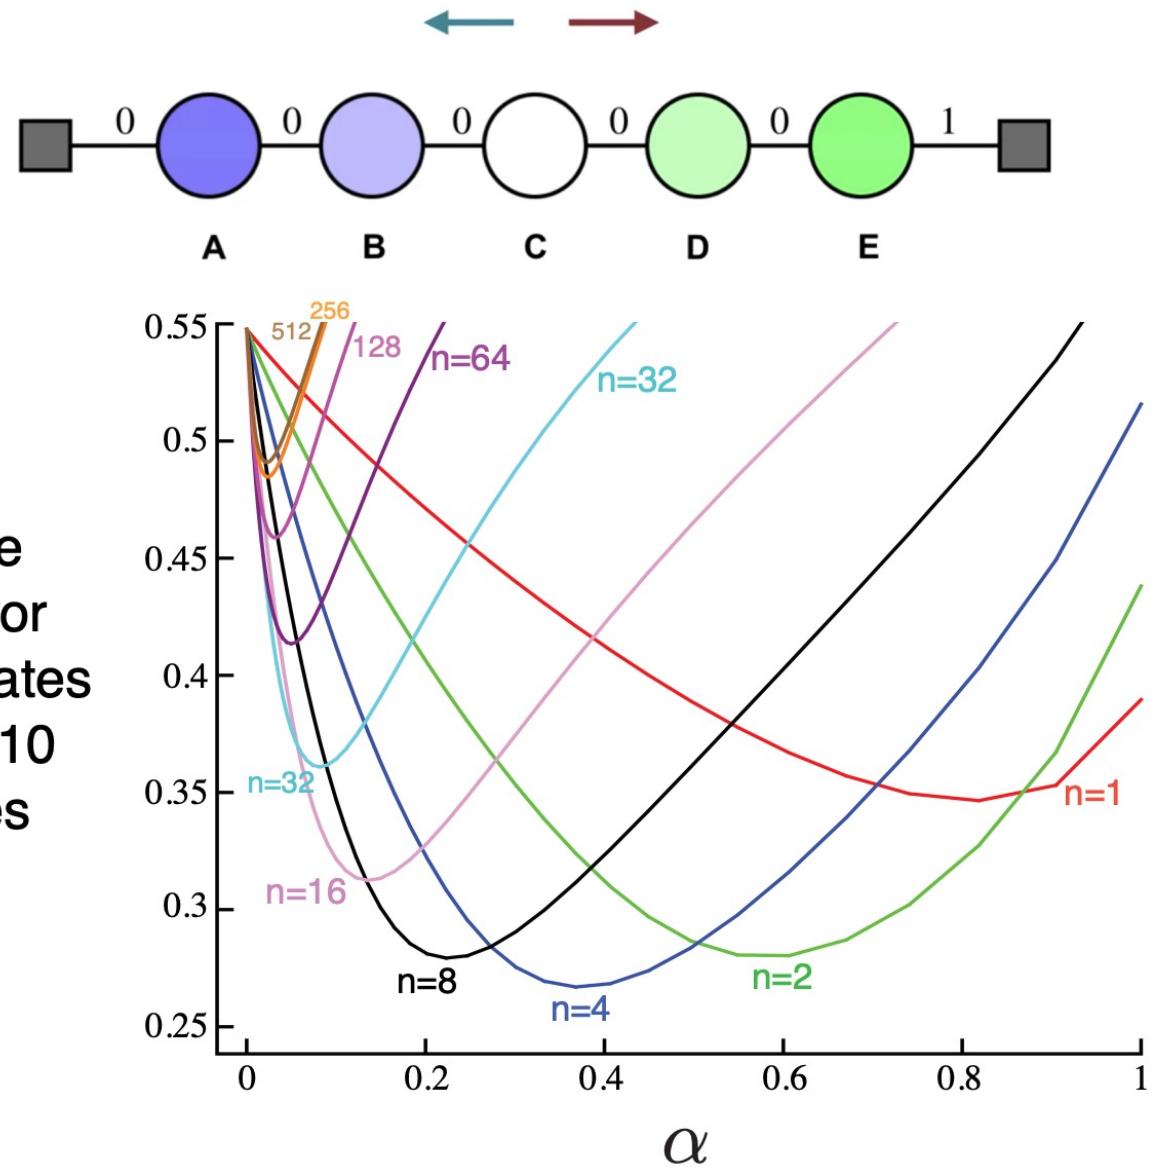


Prediction: n -step return – example

What is the ‘best’ n ? It depends!

Let’s consider the random walk process example (we are evaluating a random uniform policy): instead of 5 states, we consider 19 states!

Average RMS error over 19 states and first 10 episodes



Control: n -step SARSA

Let's define the n -step returns:

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n})$$

n -step SARSA updates:

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be ε -greedy with respect to Q , or to a fixed given policy

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$, a positive integer n

All store and access operations (for S_t , A_t , and R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq$ terminal

 Select and store an action $A_0 \sim \pi(\cdot | S_0)$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim \pi(\cdot | S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$ $(G_{\tau:\tau+n})$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$

 If π is being learned, then ensure that $\pi(\cdot | S_\tau)$ is ε -greedy wrt Q

Until $\tau = T - 1$

A fixed value for n

Loop for each episode:

Initialize and store $S_0 \neq$ terminal

Select and store an action $A_0 \sim \pi(\cdot | S_0)$

$T \leftarrow \infty$

Loop for $t = 0, 1, 2, \dots :$

If $t < T$, then:

Take action A_t

Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

else:

Select and store an action $A_{t+1} \sim \pi(\cdot | S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

$$\text{If } \tau + n < T, \text{ then } G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n}) \quad (G_{\tau:\tau+n})$$

$$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$$

If π is being learned, then ensure that $\pi(\cdot | S_\tau)$ is ε -greedy wrt Q

Until $\tau = T - 1$

t will be our step counter inside the episode

We will use:

- τ for starting the updates only when n steps are actually seen in the trajectory
- T for checking if we have arrived to close to a terminal state (and therefore a different update must be given)

Loop for each episode:

 Initialize and store $S_0 \neq$ terminal

 Select and store an action $A_0 \sim \pi(\cdot|S_0)$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots :$

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim \pi(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

 If $\tau + n < T$, then $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$ $(G_{\tau:\tau+n})$

$$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$$

 Policy-evaluation

 If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is ε -greedy wrt Q

 Until $\tau = T - 1$

Policy-improvement

Control: off-policy n -step SARSA

We can also have off-policy n -step, SARSA. We need to consider:

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha \rho_{t:t+n-1} [G_{t:t+n} - V_{t+n-1}(S_t)]$$

where an importance sampling ratio is in place

$$\rho_{t:h} \doteq \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k | S_k)}{b(A_k | S_k)}$$

Off-policy n -step Sarsa for estimating $Q \approx q_*$ or q_π

Input: an arbitrary behavior policy b such that $b(a|s) > 0$, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be greedy with respect to Q , or as a fixed given policy

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

All store and access operations (for S_t , A_t , and R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq$ terminal

 Select and store an action $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots :$

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n, T-1)} \frac{\pi(A_i|S_i)}{b(A_i|S_i)} \quad (\rho_{\tau+1:\tau+n})$$

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

$$\text{If } \tau + n < T, \text{ then: } G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n}) \quad (G_{\tau:\tau+n})$$

$$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$$

 If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is greedy wrt Q

Until $\tau = T - 1$

Loop for each episode:

Initialize and store $S_0 \neq$ terminal

Select and store an action $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

Loop for $t = 0, 1, 2, \dots :$

If $t < T$, then:

Take action A_t

Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

else:

Select and store an action $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

If $\tau \geq 0$:

$$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n, T-1)} \frac{\pi(A_i|S_i)}{b(A_i|S_i)} \quad (\rho_{\tau+1:\tau+n})$$

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i \quad \text{Target policy}$$

If $\tau + n < T$, then: $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$

$$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$$

If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is greedy wrt Q

Until $\tau = T - 1$

Actions are taken according
to a behaviour policy

$(\rho_{\tau+1:\tau+n})$

$(G_{\tau:\tau+n})$

Loop for each episode:

 Initialize and store $S_0 \neq$ terminal

 Select and store an action $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots :$

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n, T-1)} \frac{\pi(A_i|S_i)}{b(A_i|S_i)}$$

Importance Sampling must be considered

$(\rho_{\tau+1:\tau+n})$

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

Target policy

$$\text{If } \tau + n < T, \text{ then: } G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$$

$(G_{\tau:\tau+n})$

$$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$$

If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is greedy wrt Q

Until $\tau = T - 1$

n-step bootstrap: Exam

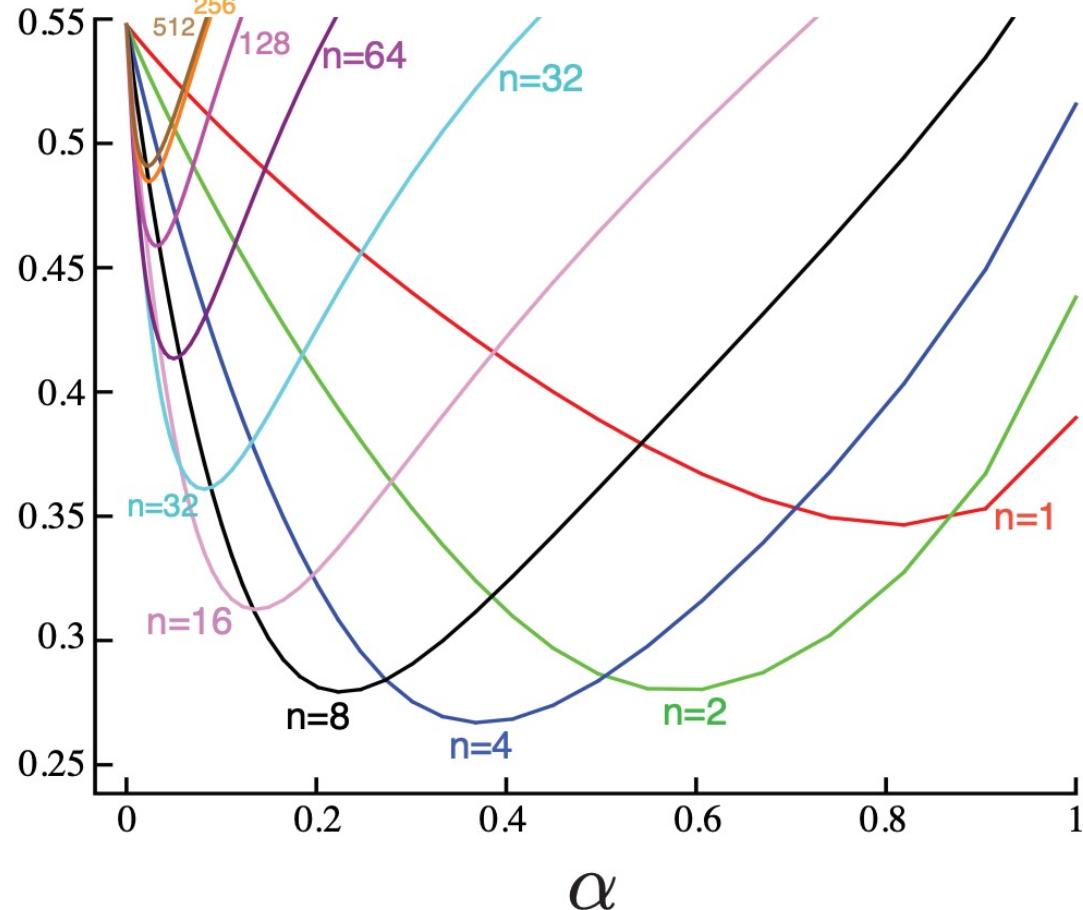
- Section 7.1 and 7.2 are exam material
- Section 7.3 (off-policy SARSA) is optional
- (Again!) Pay particular attention to the algorithms!

TD(λ)

(Chapter 12)

λ

n -step: which n ?



1-step TD
and TD(0)



2-step TD



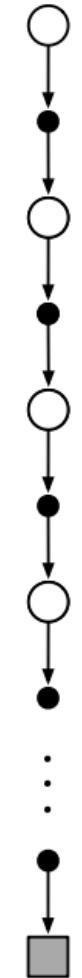
3-step TD



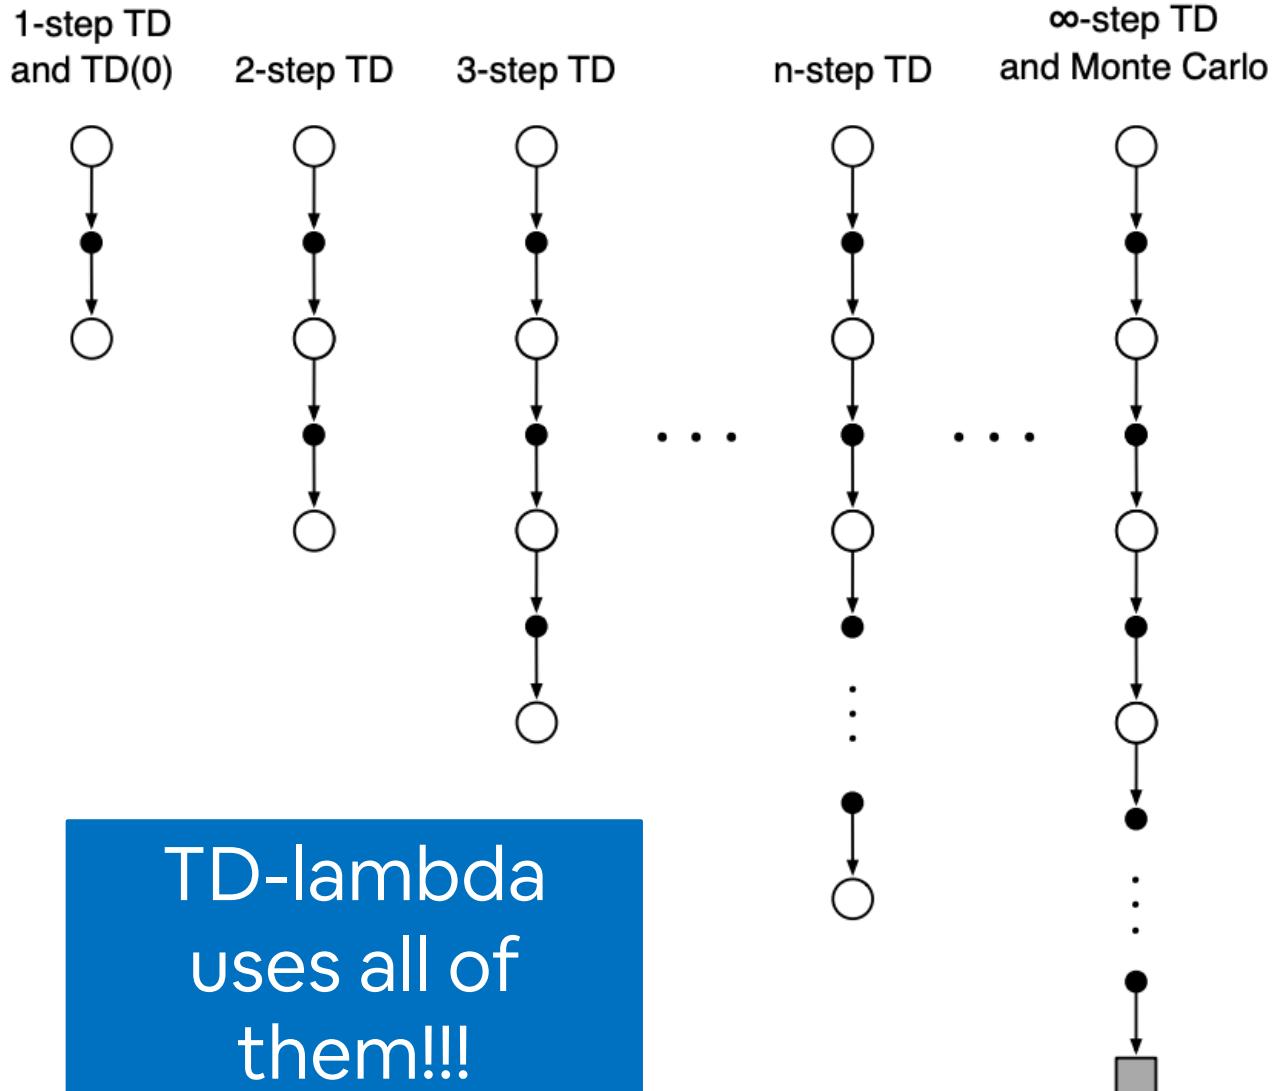
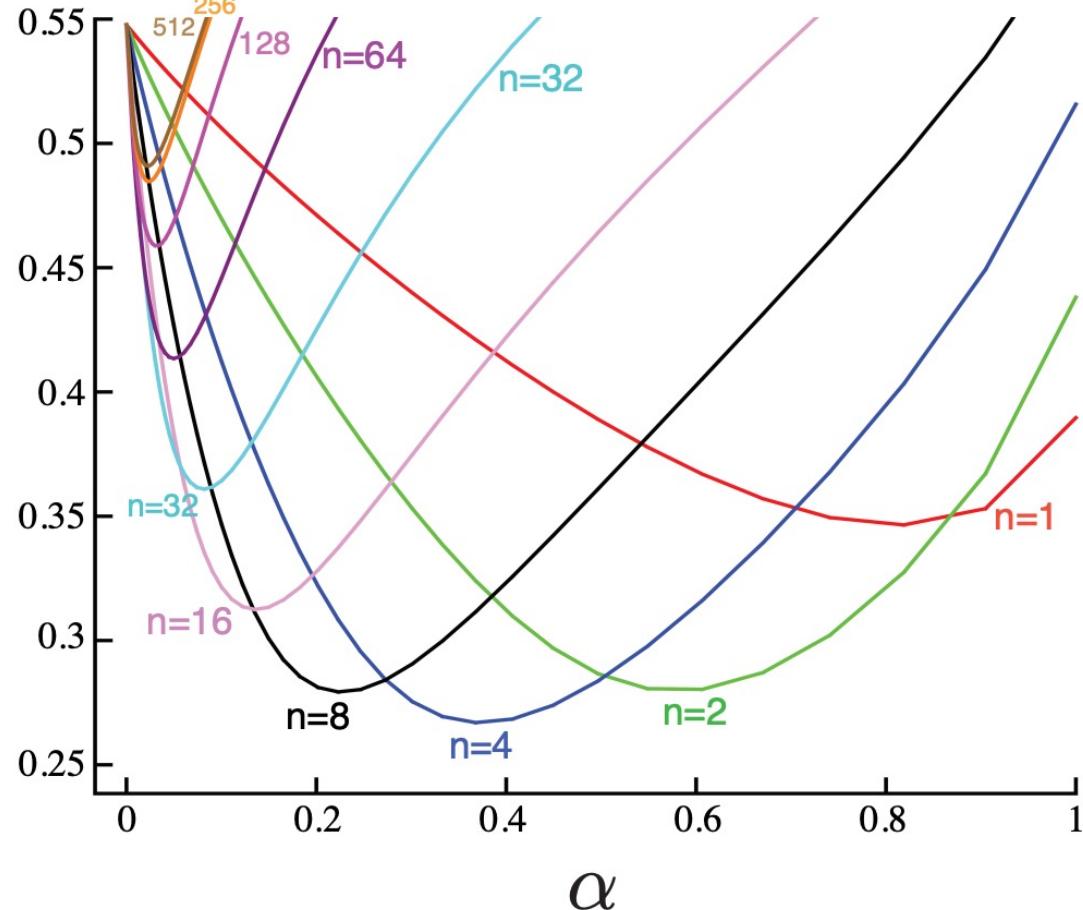
n -step TD



∞ -step TD
and Monte Carlo

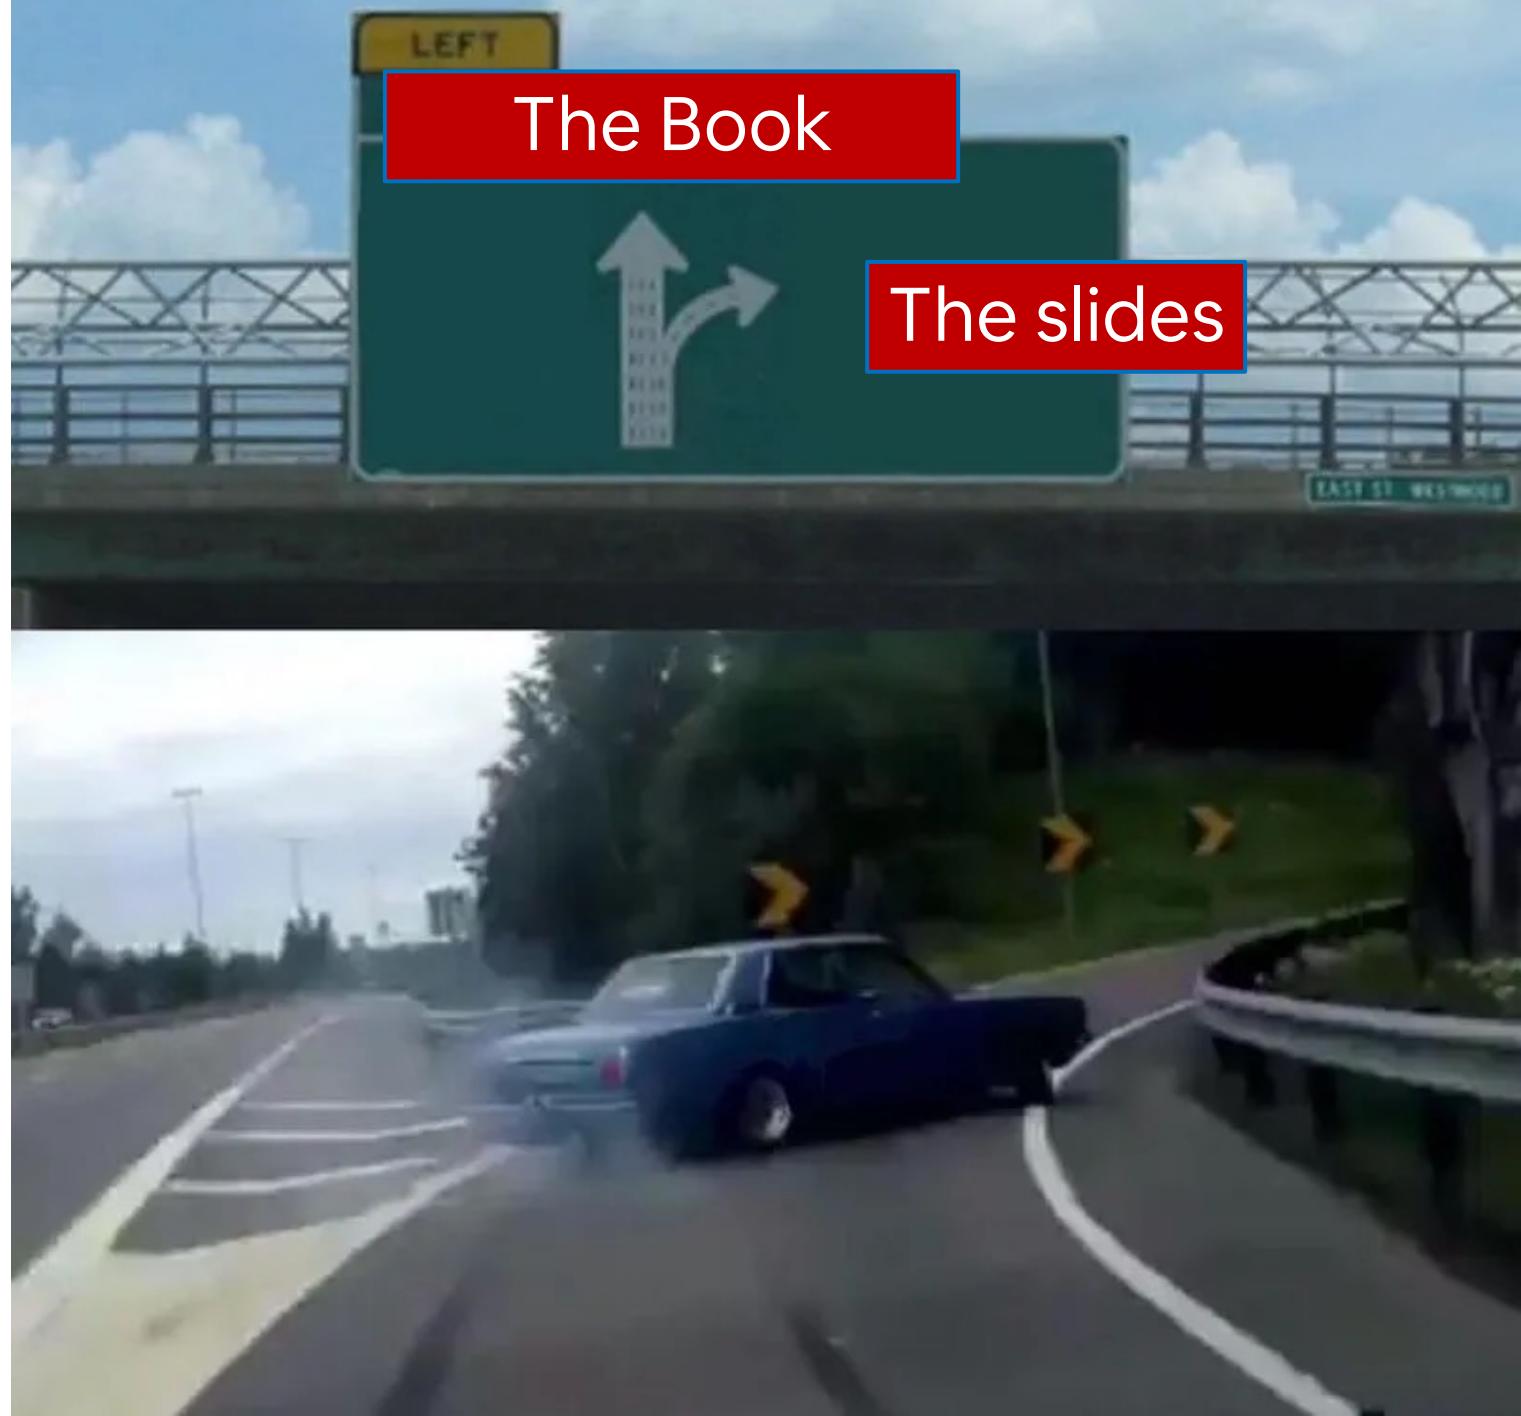


n -step: which n ?



Chapter 12

- The last version of the book described TD-lambda and *eligibility traces* (one of the main ‘tool’ in TD-lambda) in Chapter 12...
- ... but it is mixed with value function approximation VFA algorithms (in my opinion it is a bit confusing)
- For the moment, slides are the reference, after we do VFA (starting next week) the TD-lambda algorithms reported in the book will be clear

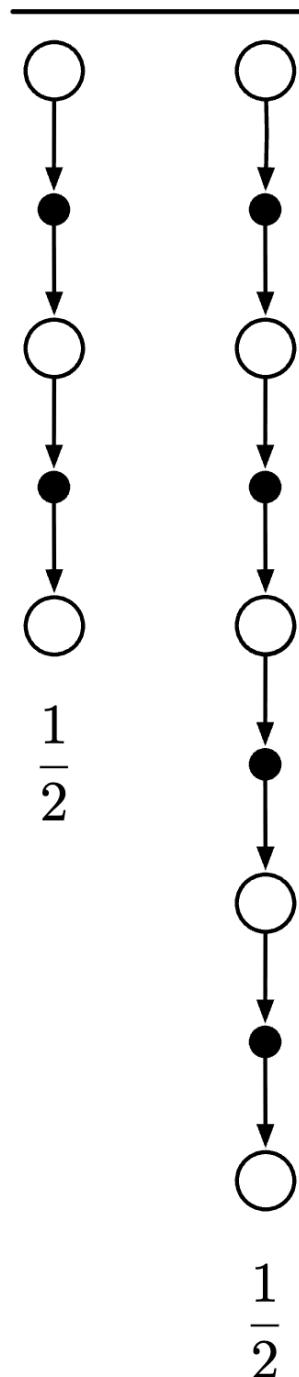


Averaging n-Step Returns

- We can average n-step returns over different n
- e.g. average the 2-step and 4-step returns
- Combines information from two different time-steps

$$\frac{1}{2}G_{t:t+2} + \frac{1}{2}G_{t:t+4}$$

- Can we efficiently combine information from all time-steps?



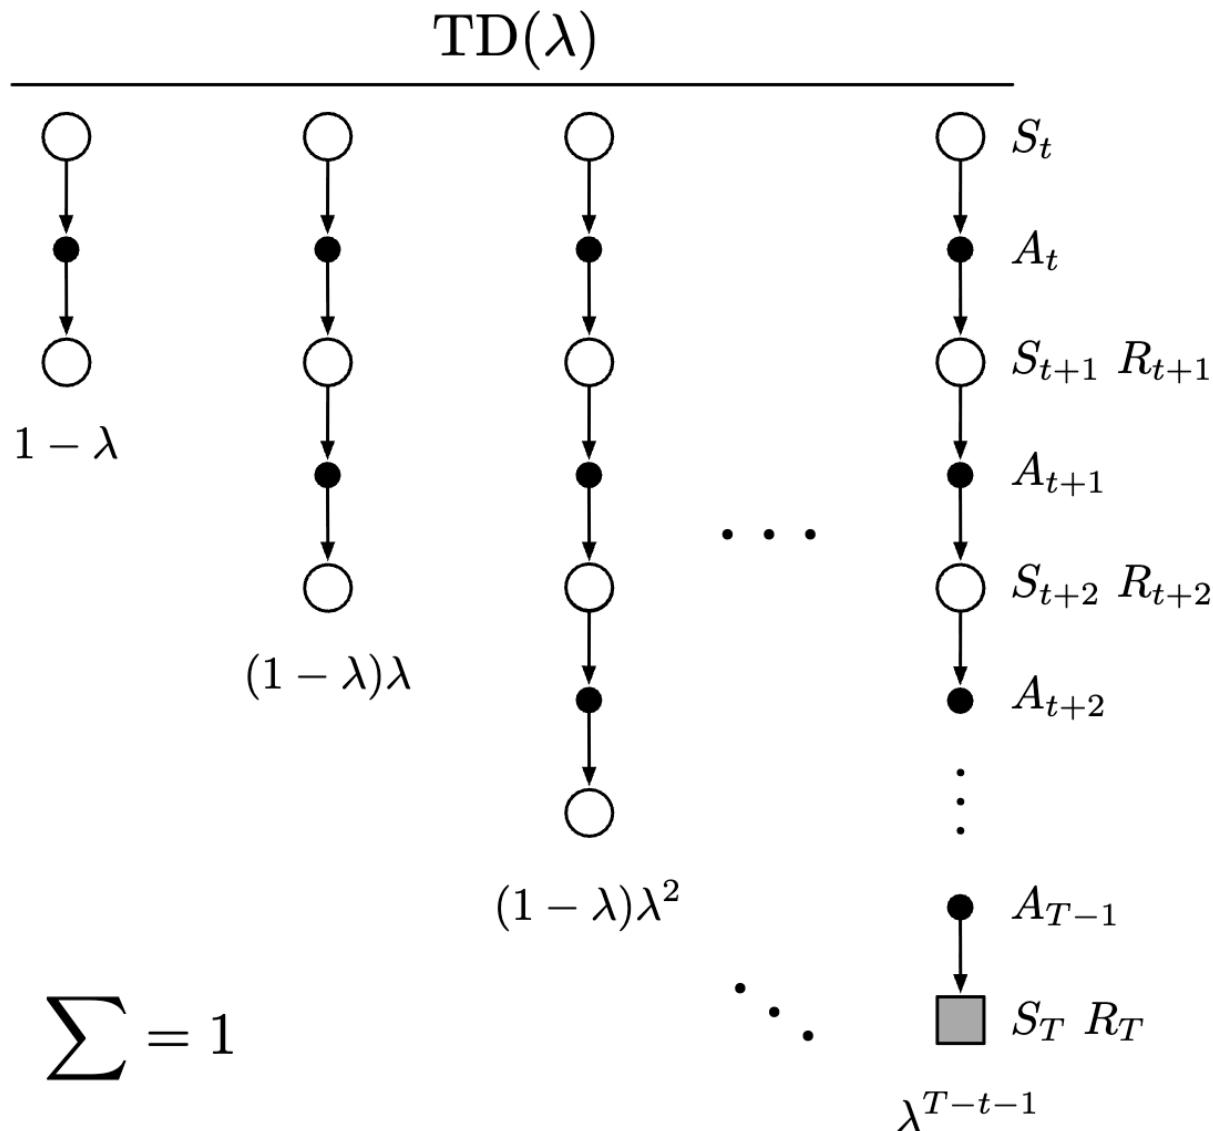
λ - return

- We consider the λ - return G_t^λ that combines all n-step returns
- It is a weighted sum of all the returns

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$$

- Instead of the ‘usual’ return we update

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$



λ - return

- We consider the λ - return G_t^λ that combines all n-step returns
- It is a weighted sum of all the returns

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$$

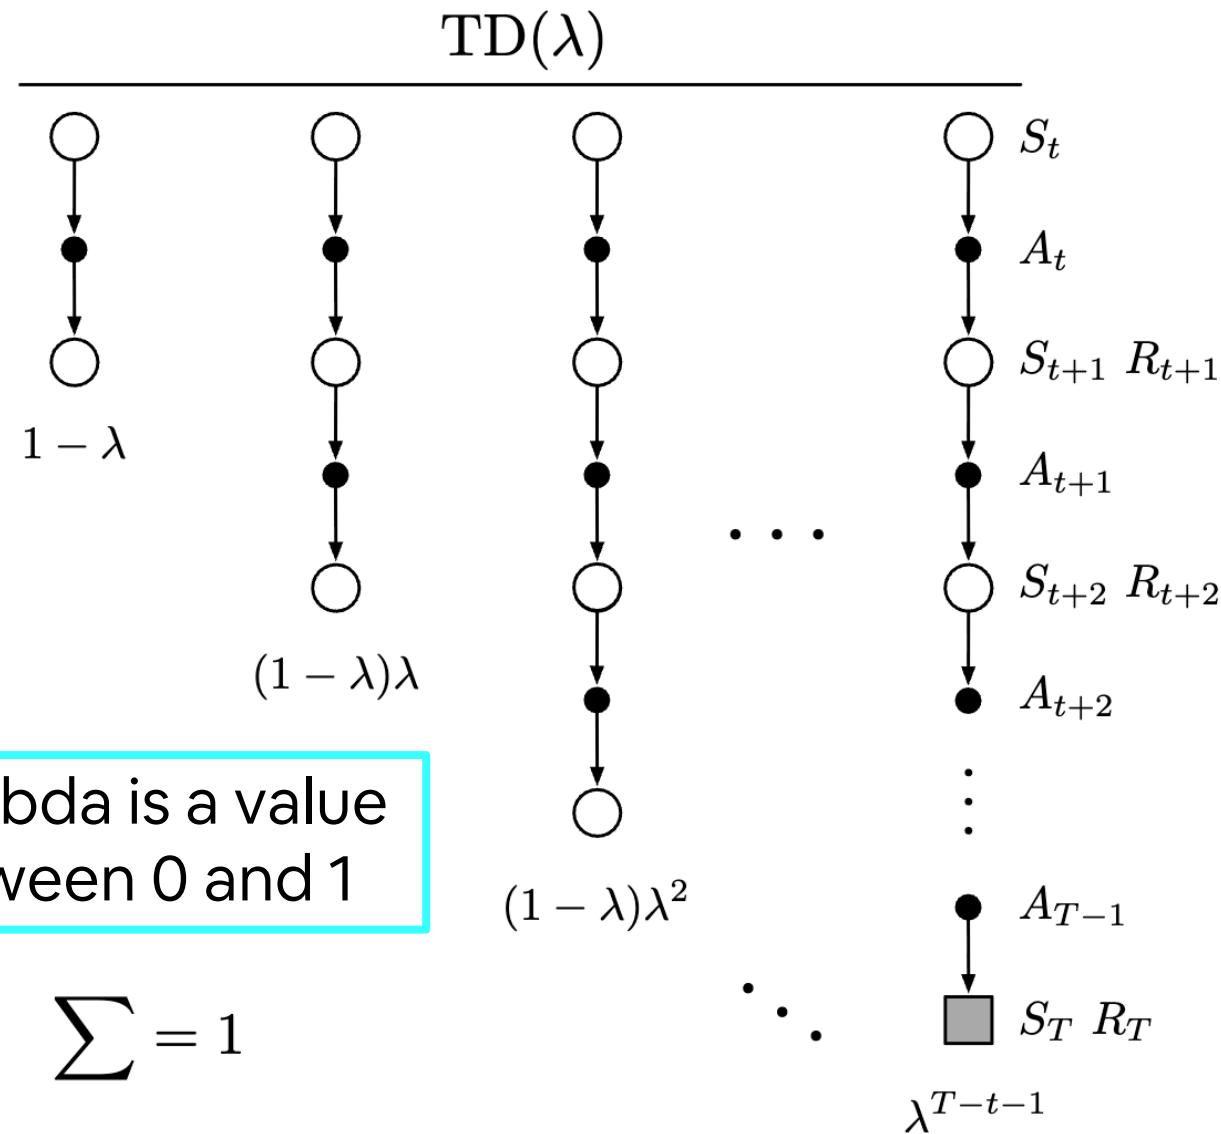
- Instead of the ‘usual’ return we update

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$

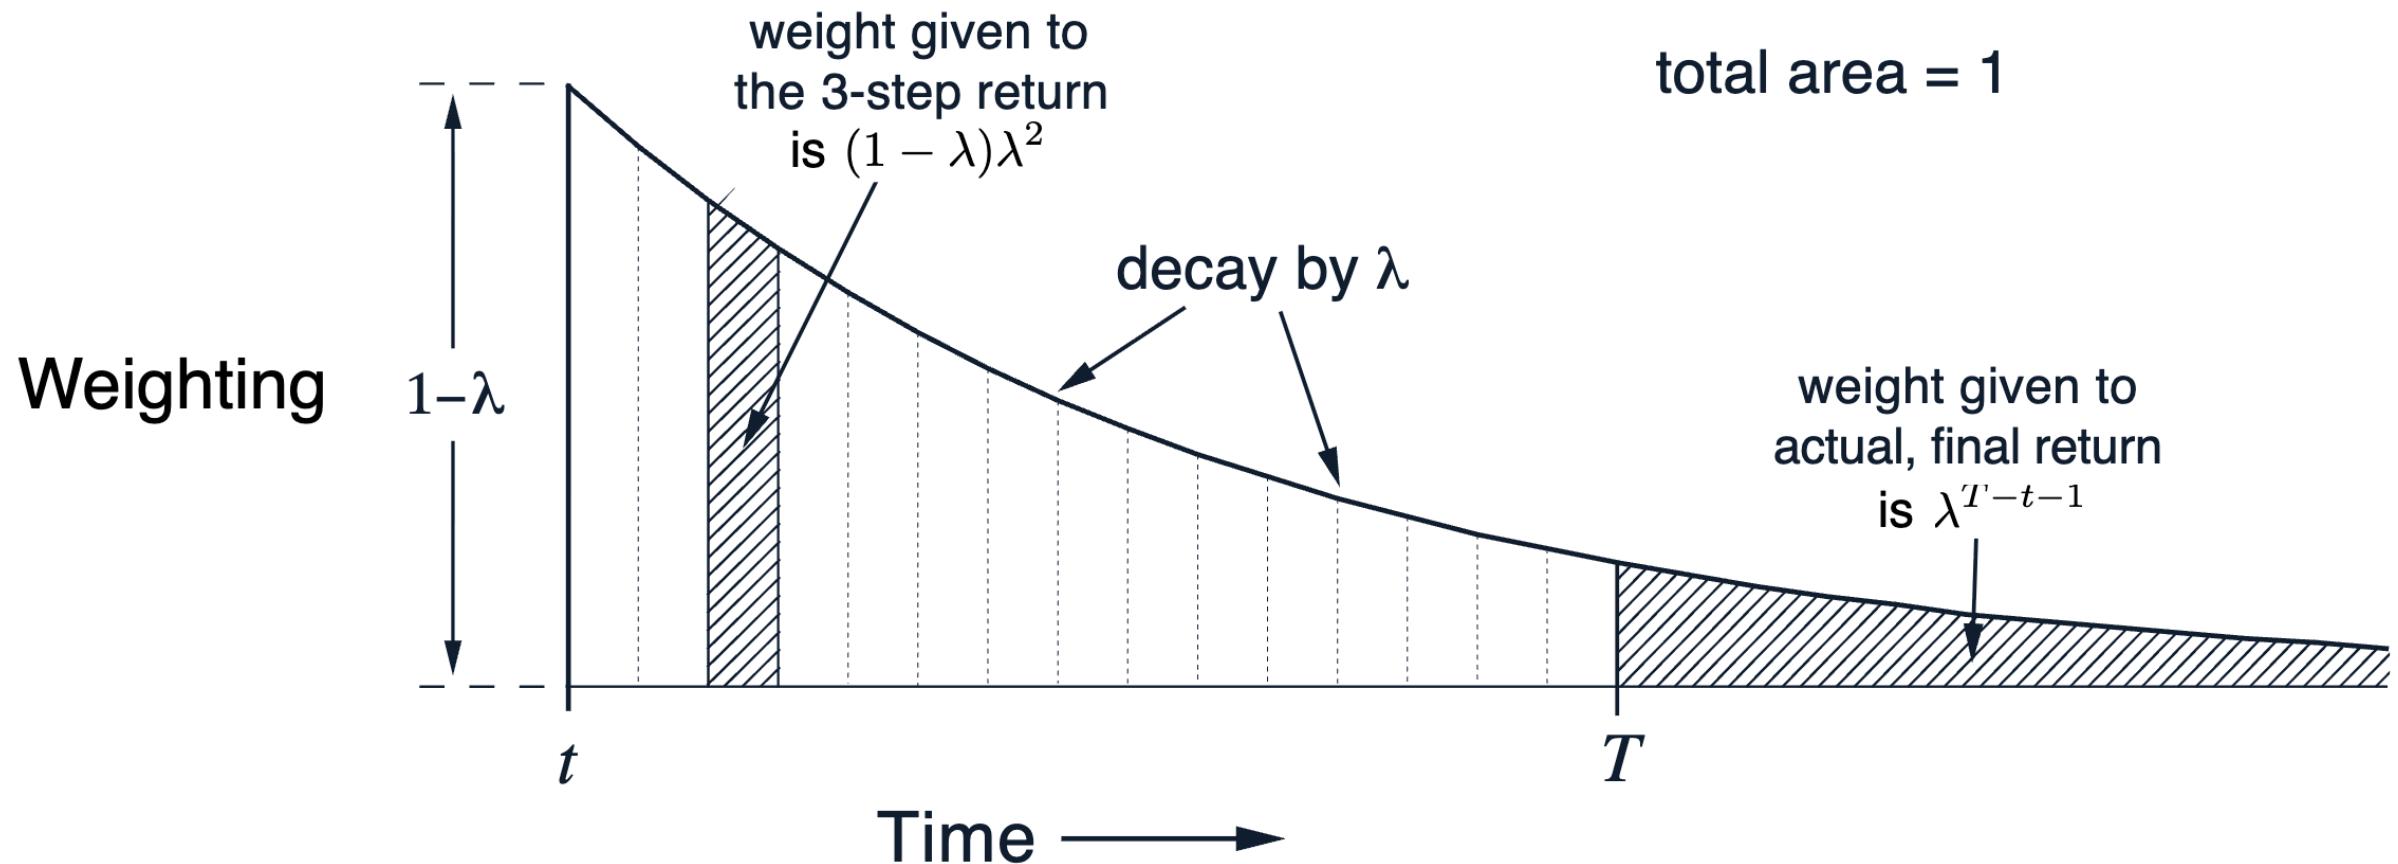
Lambda is a value between 0 and 1

$$\sum = 1$$

$$\sum_{n=1}^{\infty} \lambda^{n-1} = \sum_{n=0}^{\infty} \lambda^n = \frac{1}{1 - \lambda}$$



TD(λ) Weighting Function



Why TD(λ) Weighting Function?

Why we weight the steps?

The weights form a geometric decay.

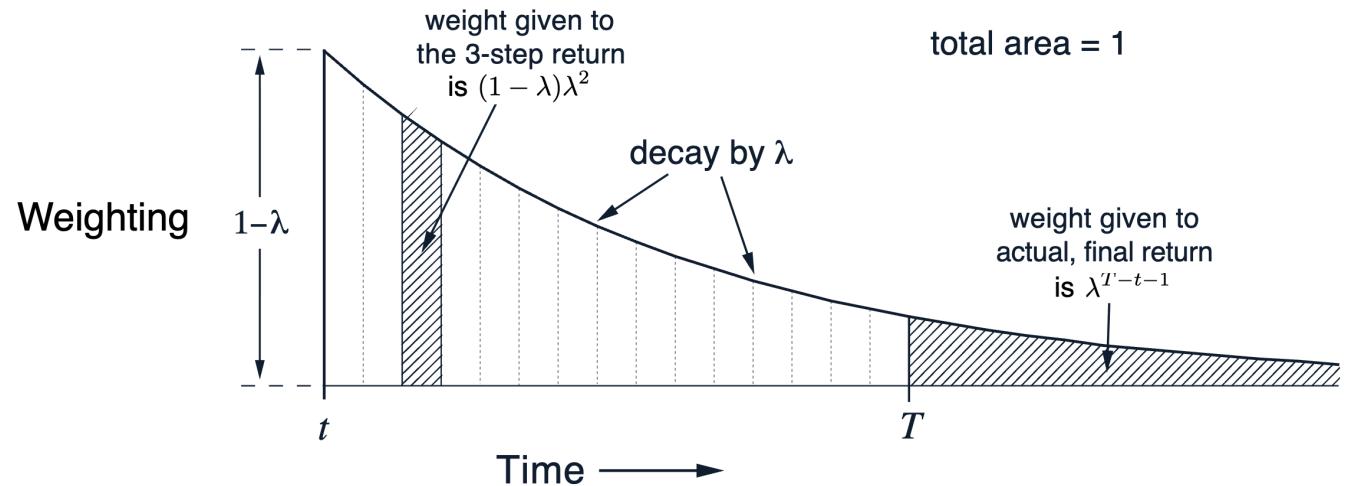
That means:

- 1-step return gets the highest weight
- 2-step, 3-step returns get smaller and smaller weights

This is **intentional** — it expresses how far into the future we trust our rollouts.

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$



Why TD(λ) Weighting Function?

Why not give all steps equal weight?

With that idea:

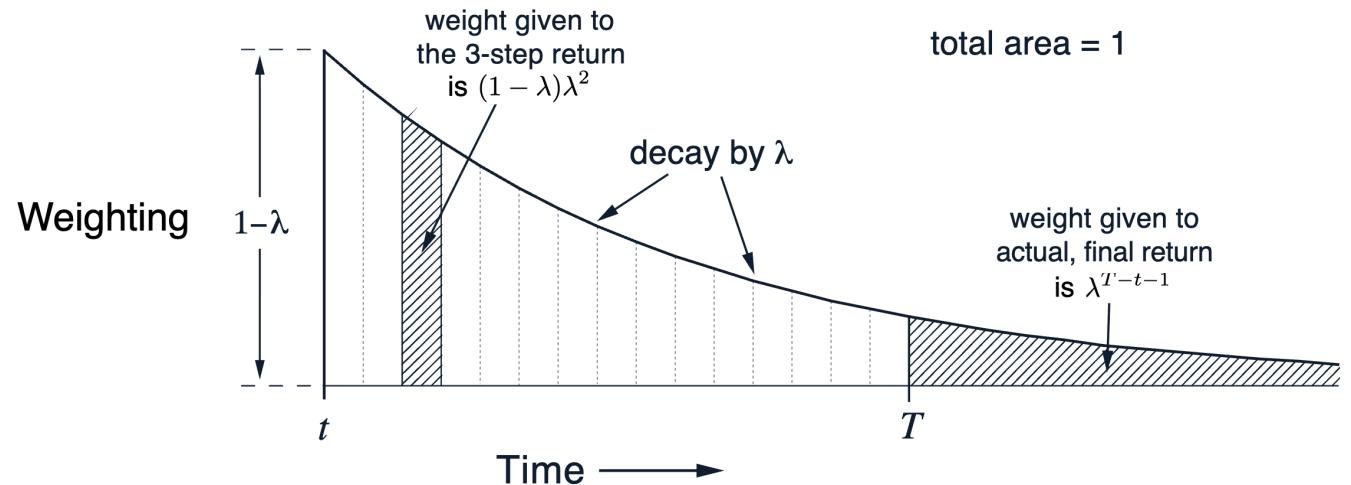
- You'd be mixing together very noisy long-term estimates with short-term, low-variance estimates.
- The result wouldn't correspond to any meaningful bias-variance trade-off.
- In practice, it would be unstable, especially when the environment is long or continuous.

The exponential decay in λ ensures that:

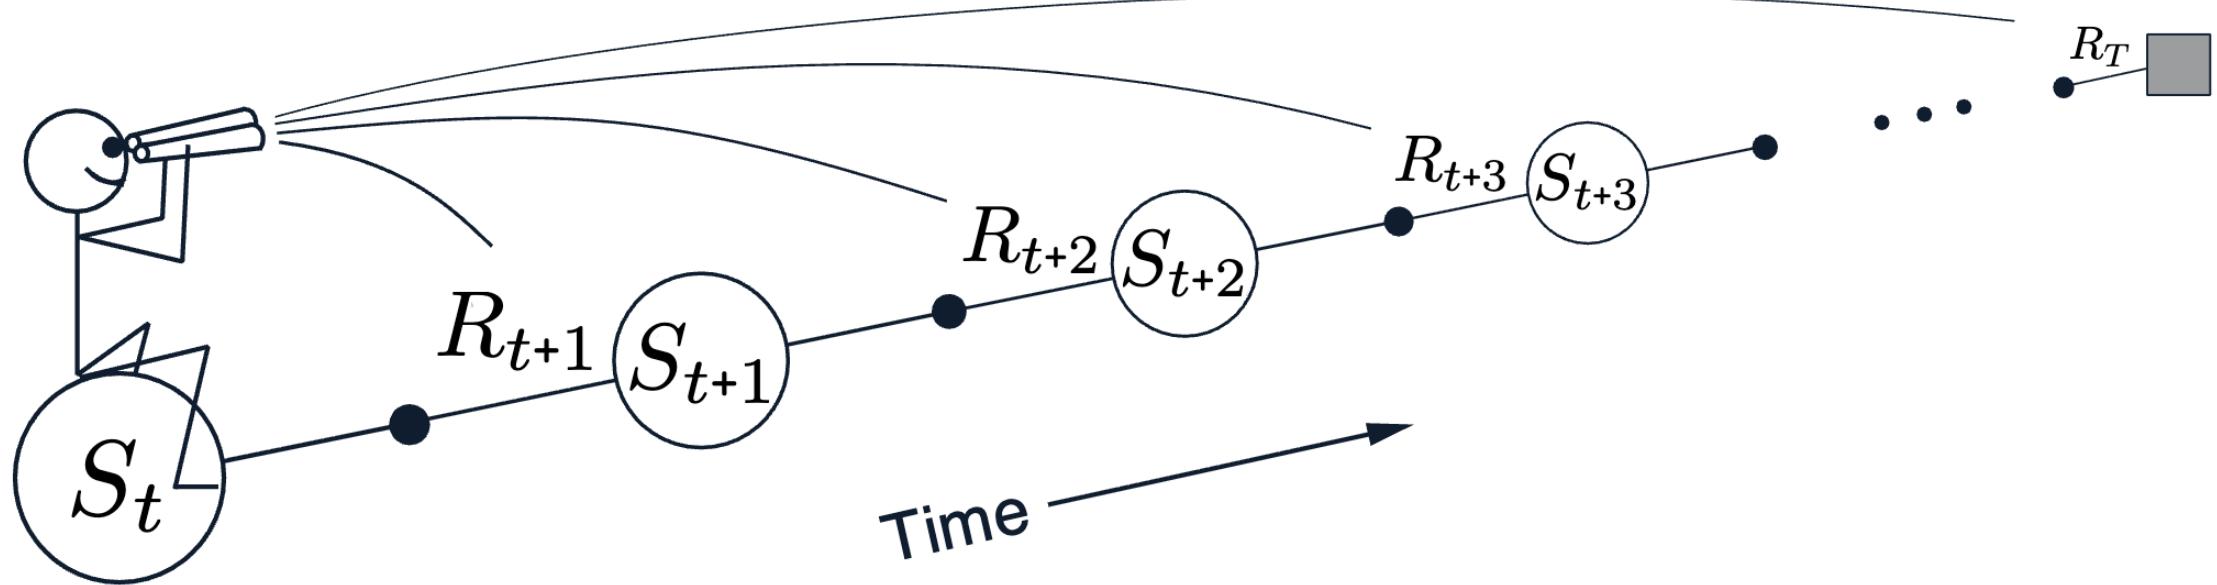
- recent TD errors matter more (because they're more reliable)
- but you still get some credit assignment from later steps.

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$



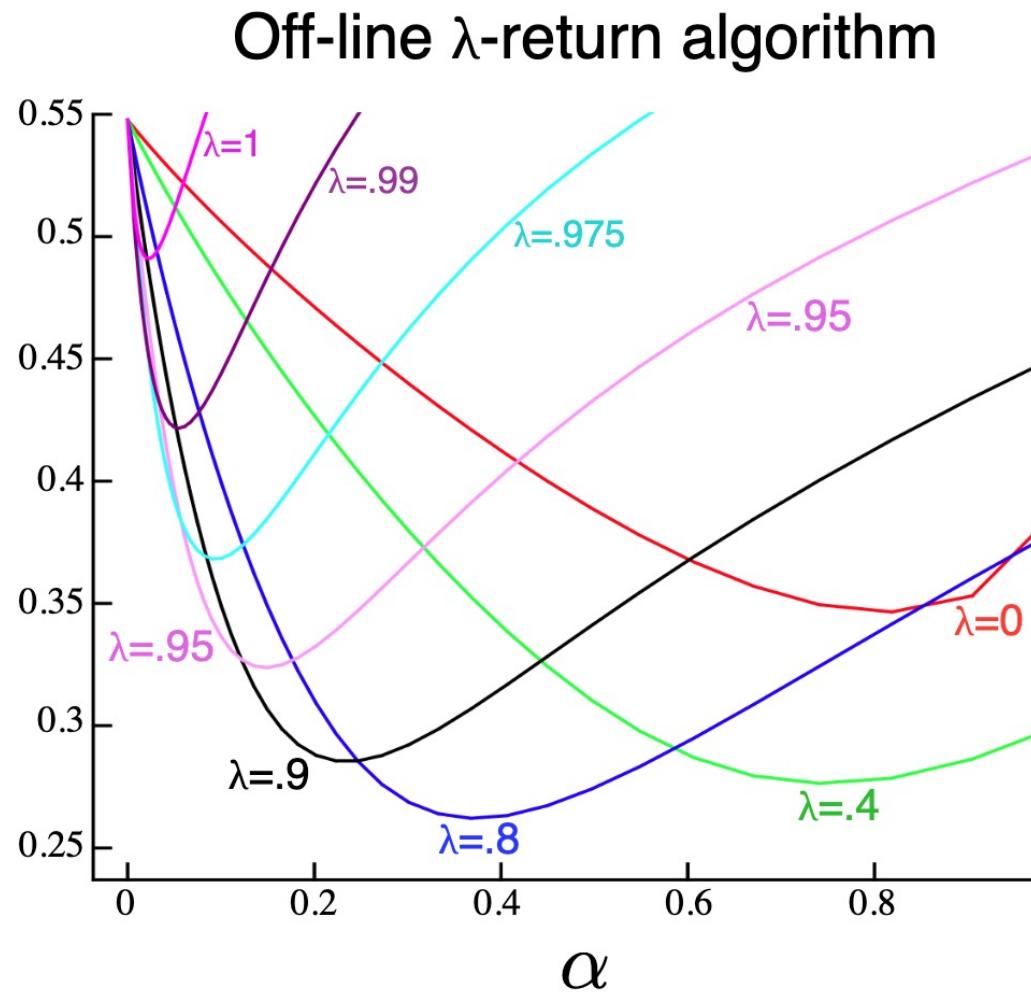
(Forward-view) TD(λ)



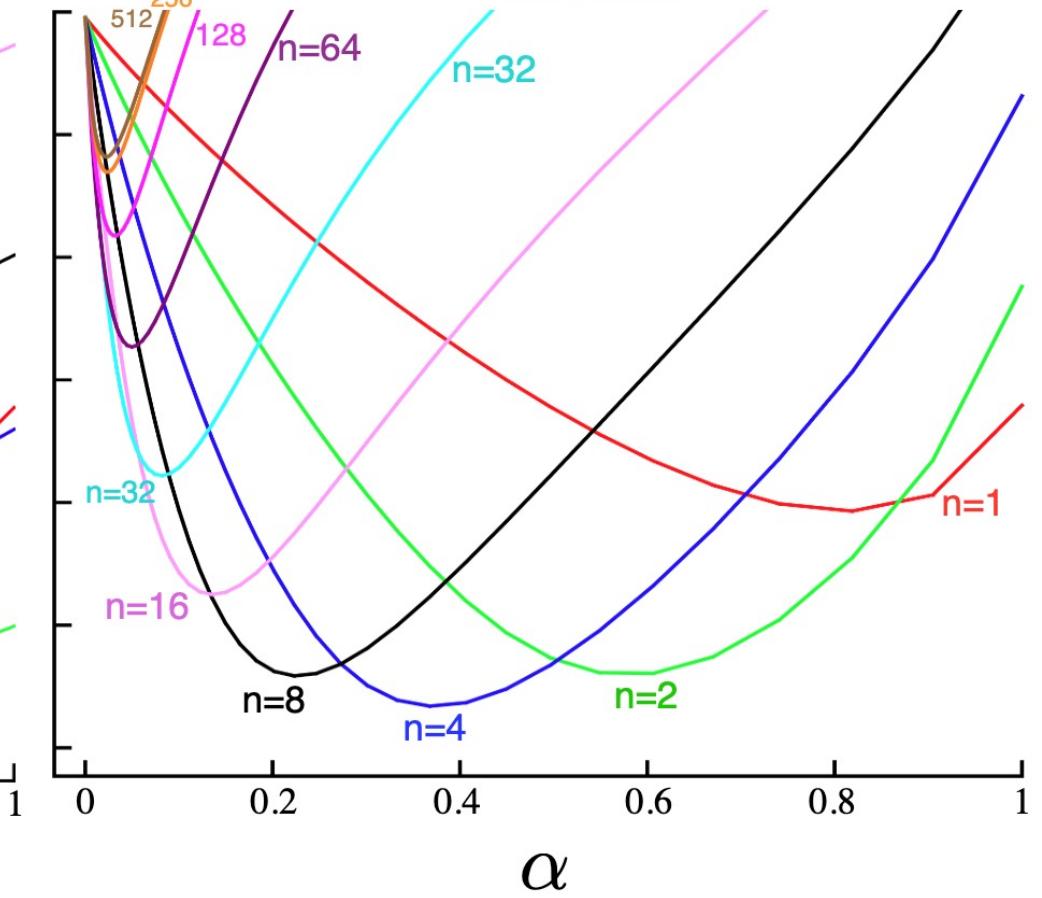
- Update value function towards the λ -return
- Forward-view looks into the future to compute λ - return G_t^λ
- Like Monte Carlo, it can only be computed from complete episodes

TD(λ) in action: ex. Prediction random walk

RMS error
at the end
of the episode
over the first
10 episodes



n-step TD methods
(from [Chapter 7](#))

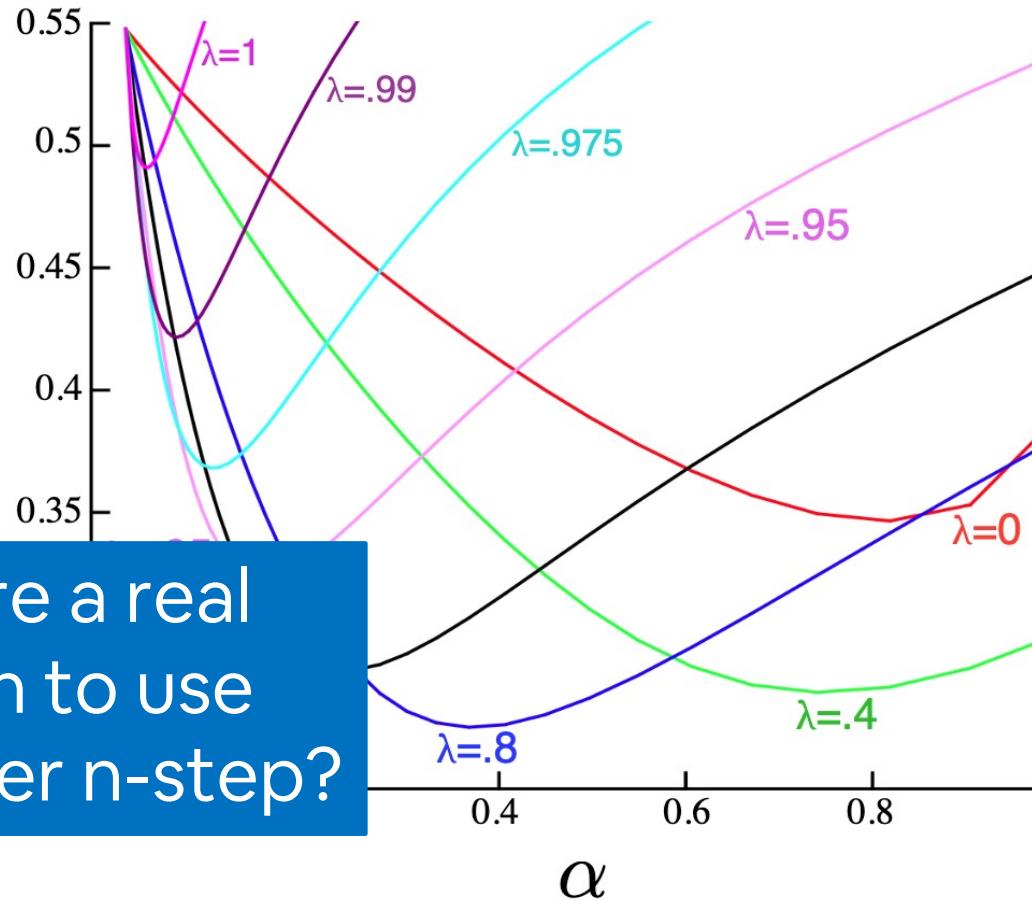


TD(λ) in action: ex. Prediction random walk

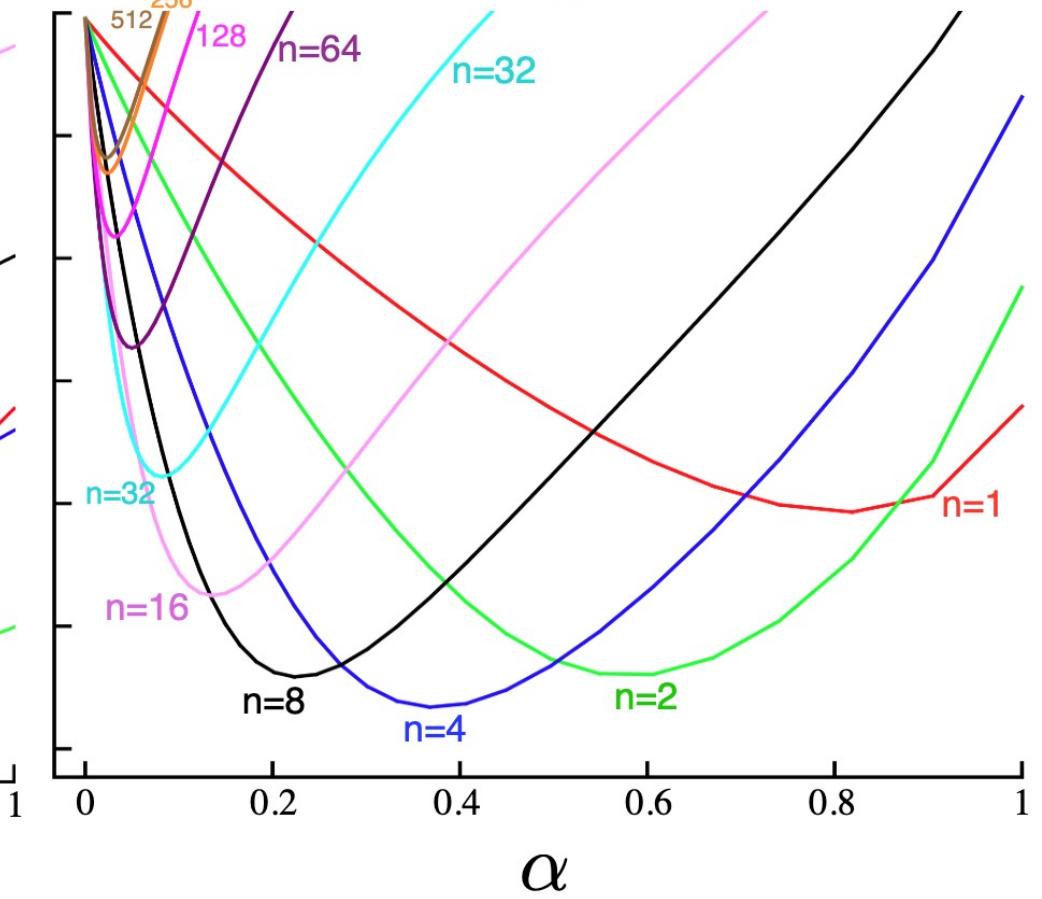
RMS error
at the end
of the episode
over the first
10 episodes

Is there a real
reason to use
TD(λ) over n-step?

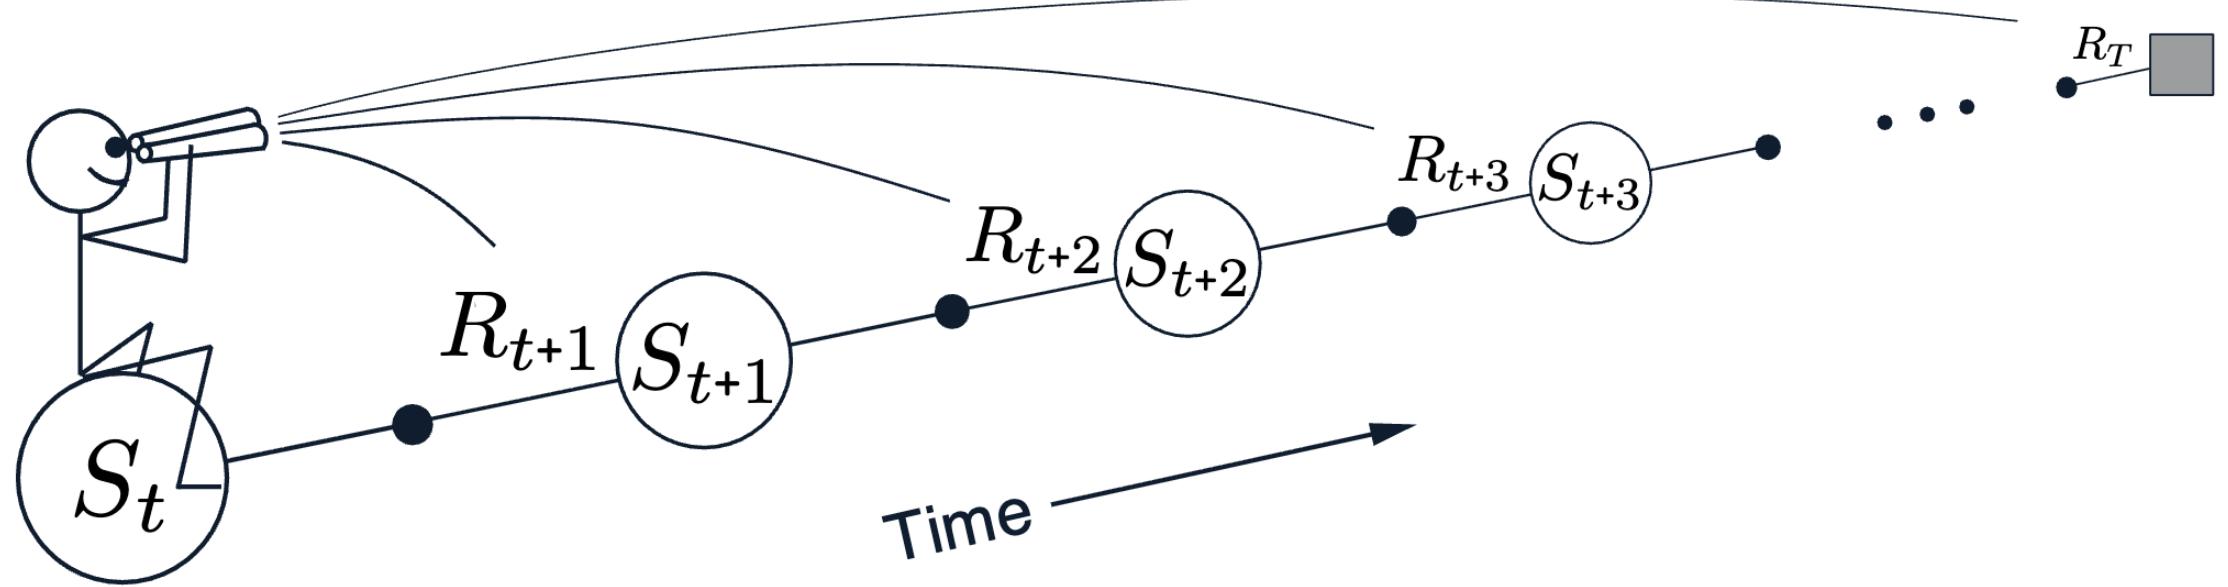
Off-line λ -return algorithm



n-step TD methods
(from [Chapter 7](#))

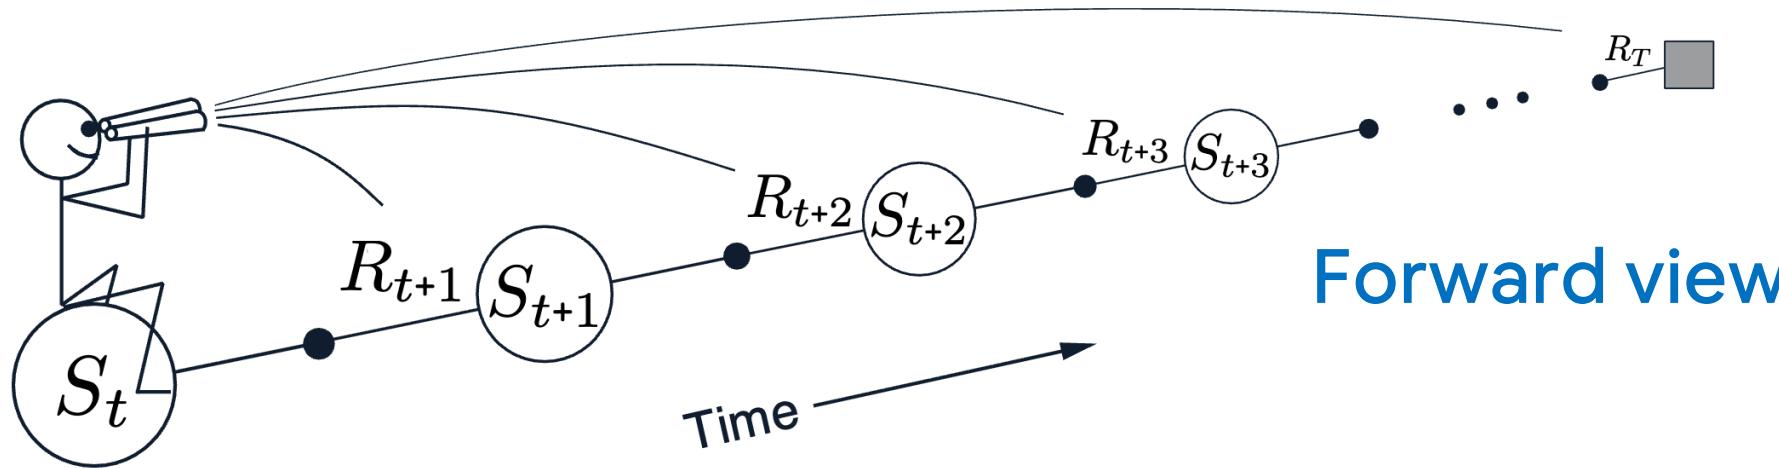


Forward-view TD(λ)

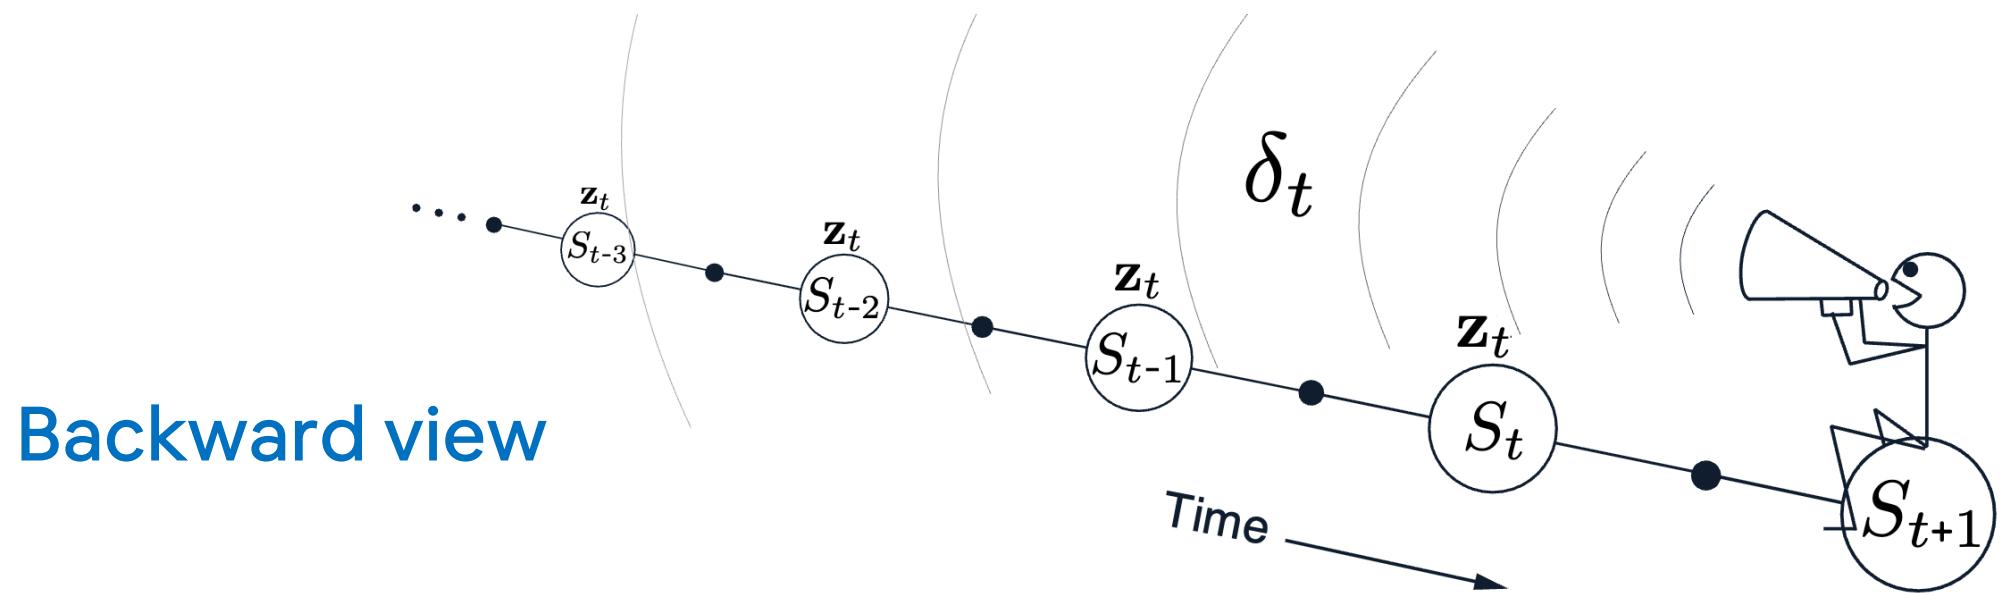


Forward views are somehow complex to implement because the update of each state depends on later events or rewards that are not available at current time

Let's change our perspective!

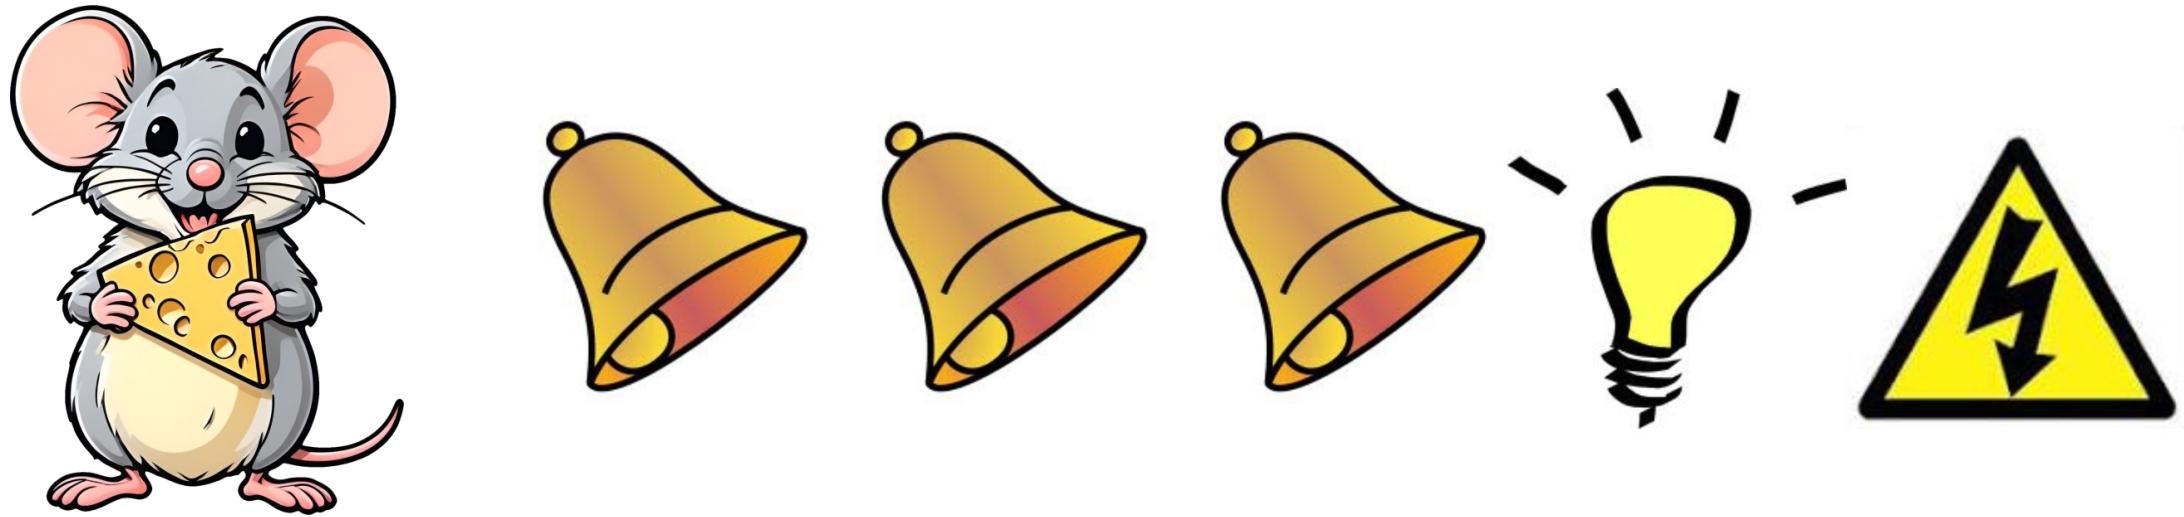


Forward view



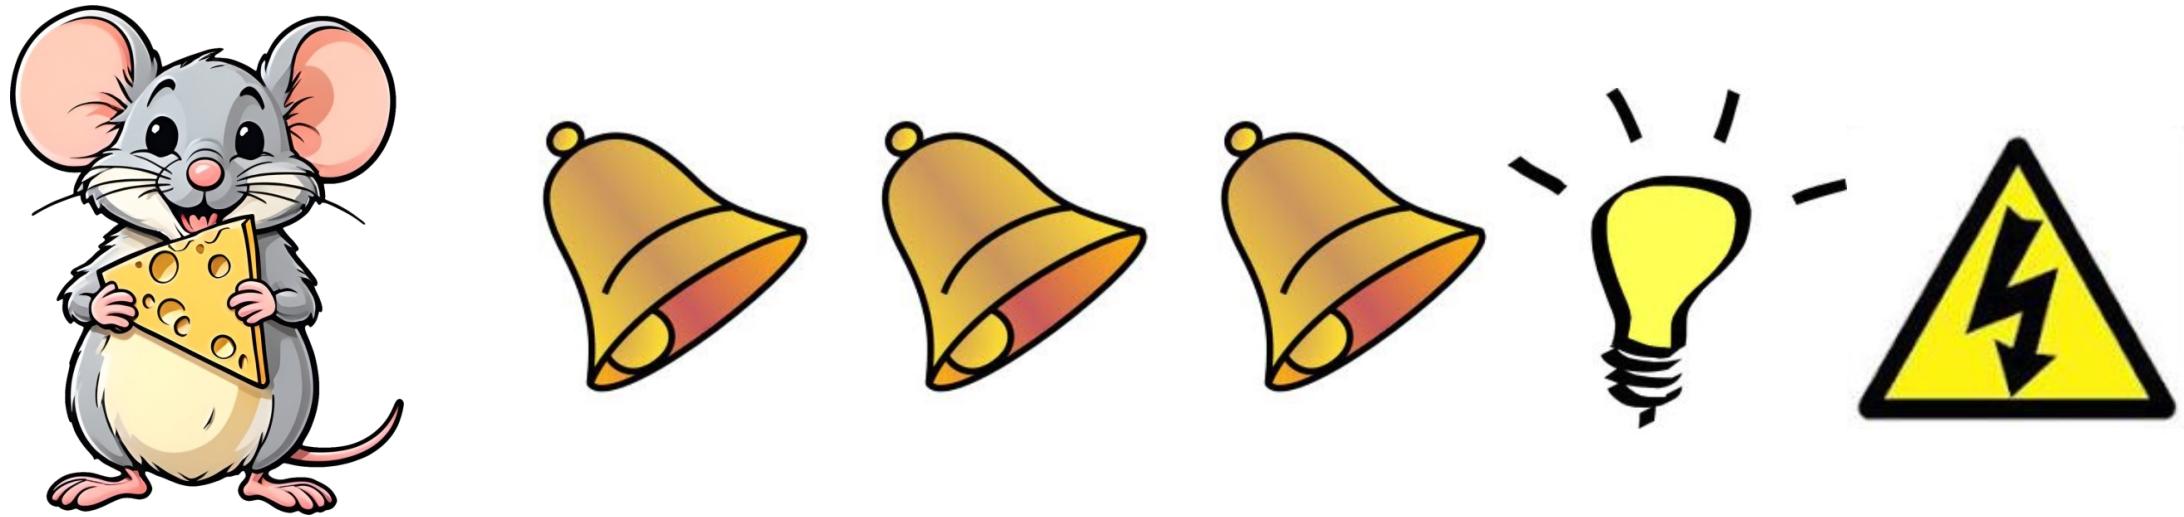
Backward view

Eligibility traces



- Credit assignment problem: did bell or light cause shock?

Eligibility traces

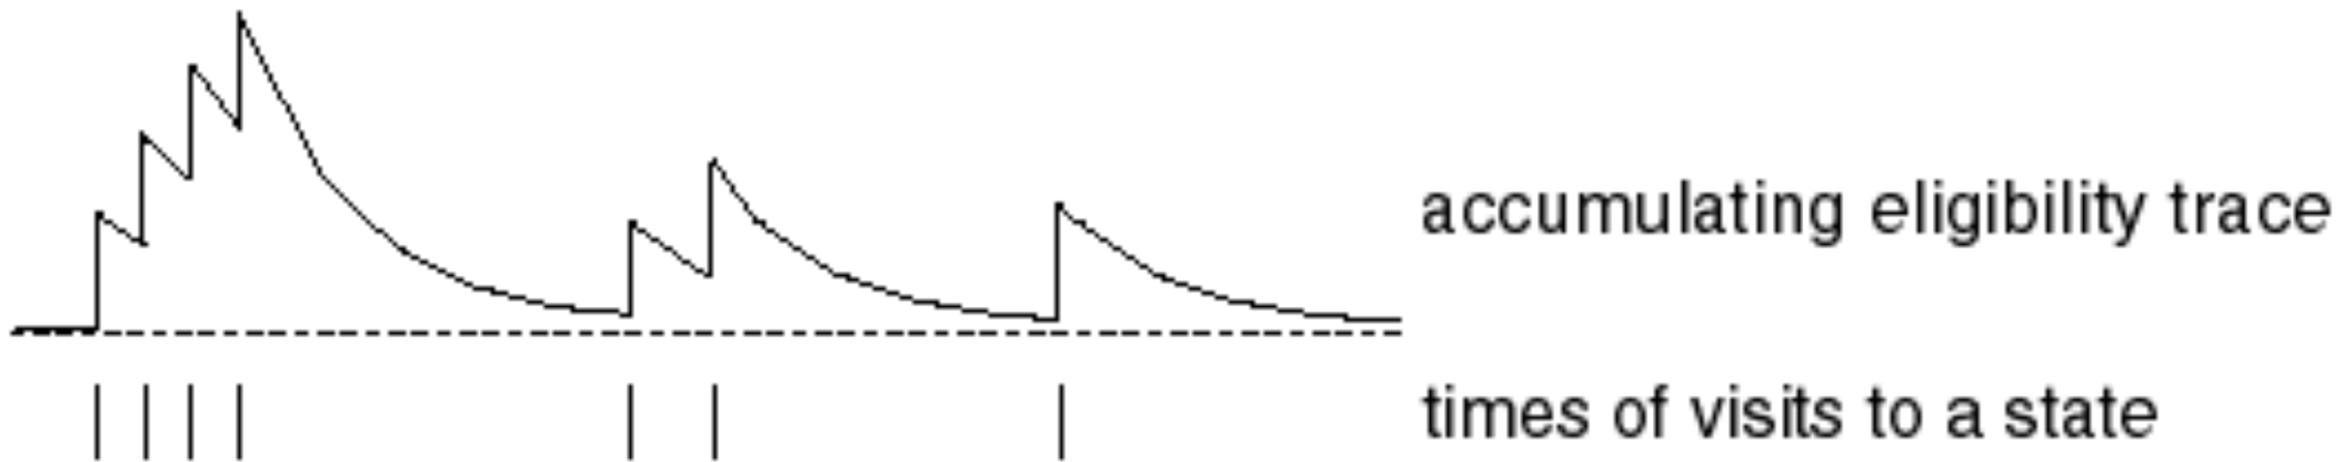


- Credit assignment problem: did bell or light cause shock?
- Frequency heuristic: assign credit to most frequent states
- Recency heuristic: assign credit to most recent states

Eligibility traces

$$E_0(s) = 0$$

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$



I have one eligibility trace for every state in the MDP and they evolve over the episode!

Eligibility traces (prediction)

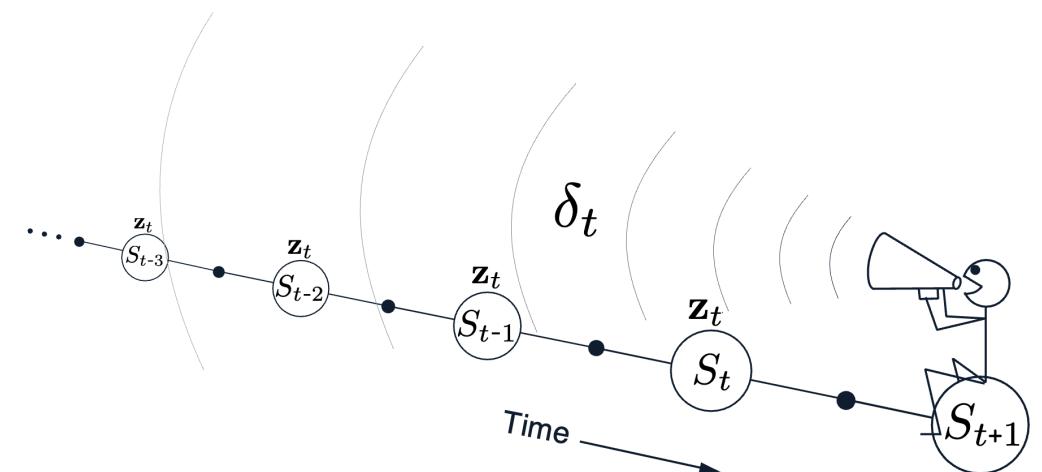
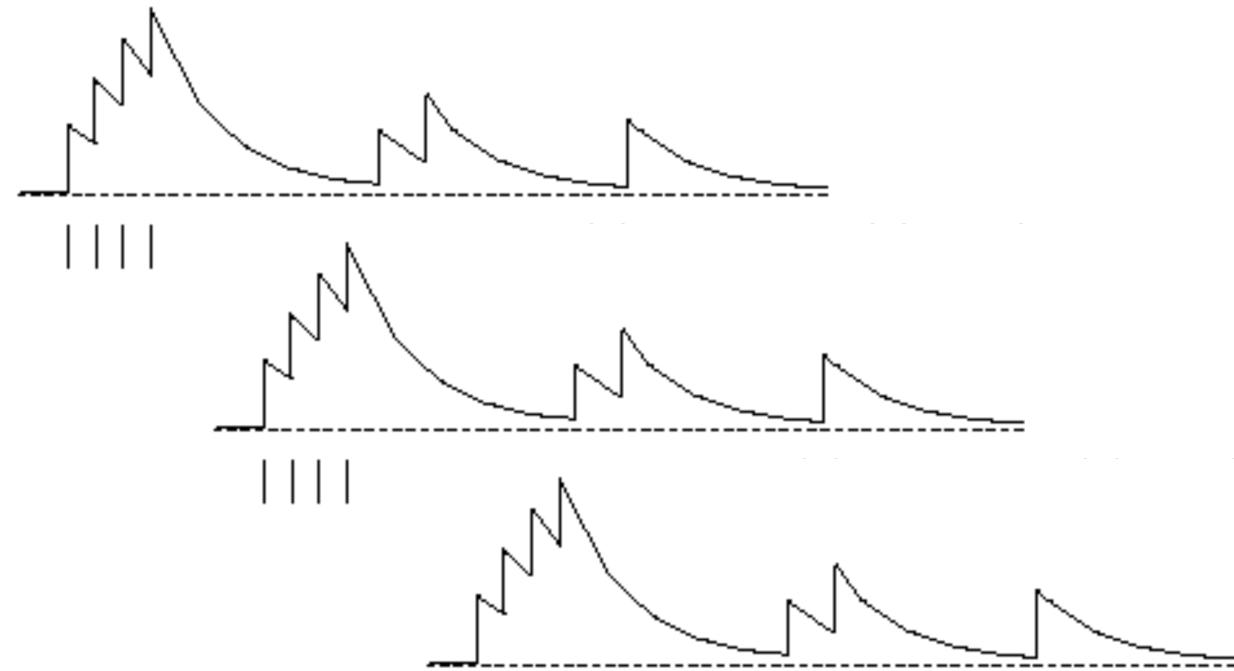
$$E_0(s) = 0$$

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$

We have an update **for all s (!)** at each step (like in TD-0) that depends on the **TD-error** and the **eligibility traces**

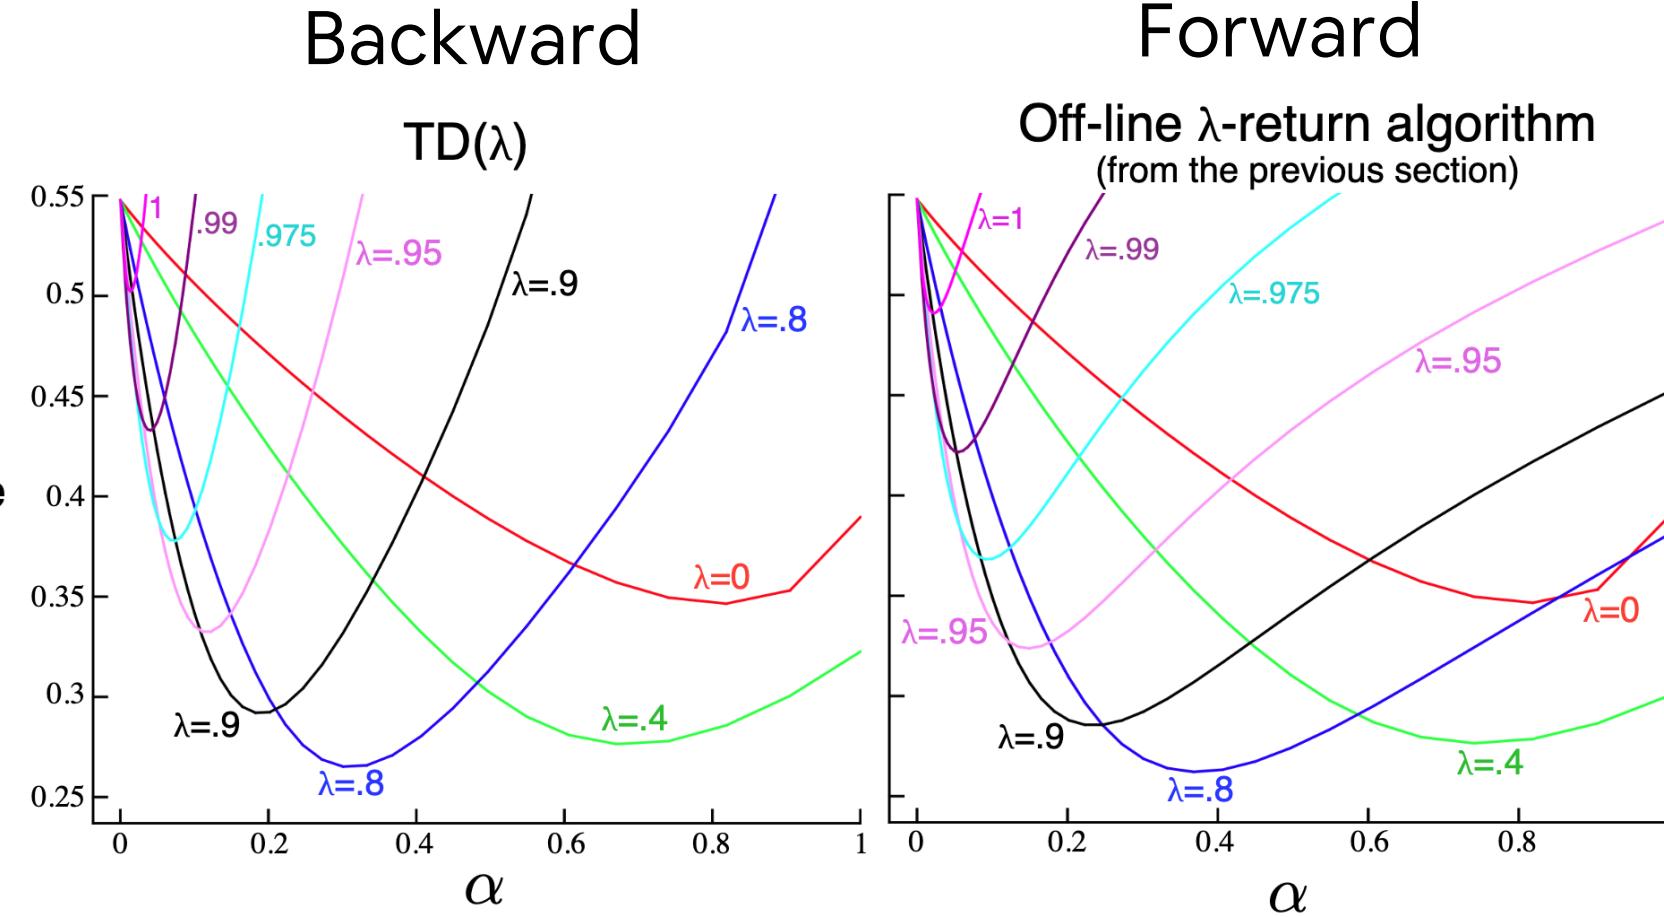
$$\delta_t = R_{t+1} + \gamma V_t(S_{t+1}) - V_t(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



Forward vs Backward view (prediction) - offline

RMS error at the end of the episode over the first 10 episodes



Off-line (after we have completed the episode) we have equivalence of the two approaches... but backward view works without waiting the end of the episode!

Credits

- Image of the course is taken from C. Mahoney 'Reinforcement Learning' <https://towardsdatascience.com/reinforcement-learning-fda8ff535bb6>



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Reinforcement Learning

2024/2025



Thank you! Questions?

Gian Antonio Susto
Ruggero Carli

