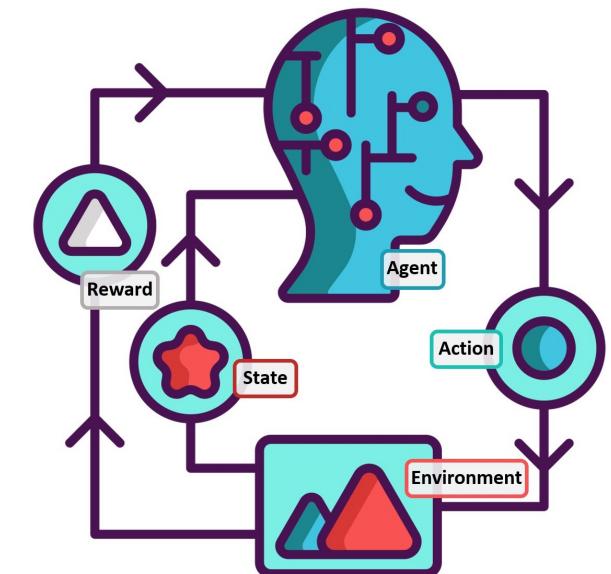


Lecture #14

Value Function Approximation

Gian Antonio Susto



Announcements before starting: 34 Importance Sampling Algorithm submissions...

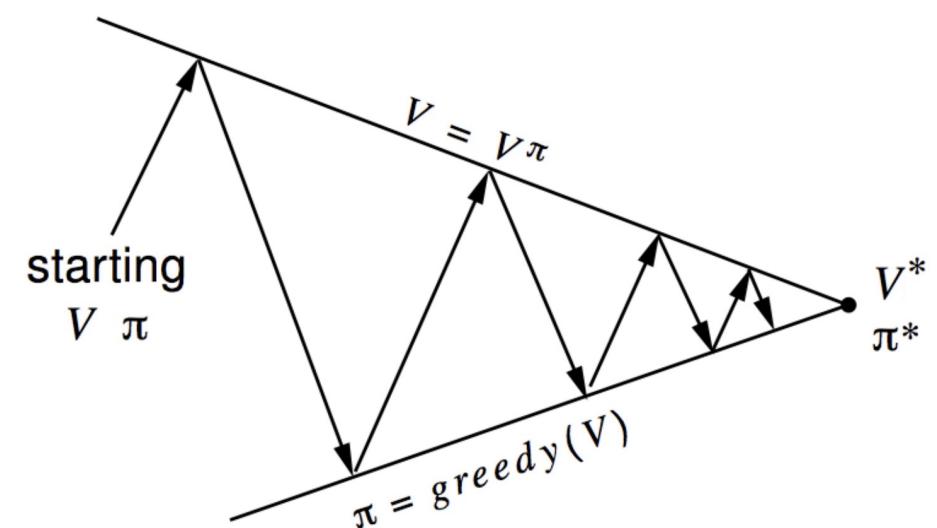
An example of good
solution

Algorithm 1 Monte Carlo Off-Policy Prediction for $V_\pi(s)$ (Weighted Importance Sampling - Average Implementation)

```
1: Input: target policy  $\pi$ , behaviour policy  $b$ 
2: Initialize  $V(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$ 
3: Initialize  $W_{\text{sum}}(s) \leftarrow 0$  for all  $s \in \mathcal{S}$                                 ▷ Cumulative Sum of weights
4: Initialize  $G_{\text{sum}}(s) \leftarrow 0$  for all  $s \in \mathcal{S}$                                 ▷ Cumulative sum of Returns
5: loop                                                                                   ▷ For each episode
6:   Generate an episode following  $b$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$ 
7:    $G \leftarrow 0$ 
8:    $W \leftarrow 1$                                                                ▷ Importance sampling weight
9:   for  $t = T - 1, T - 2, \dots, 0$  do
10:     $G \leftarrow \gamma G + R_{t+1}$                                               ▷ Policy Evaluation
11:     $W \leftarrow W \cdot \frac{\pi(A_t|S_t)}{b(A_t|S_t)}$                             ▷ Importance Sampling ratio
12:    if  $W > 0$  then                                                 ▷ Only update if weight is non-zero
13:       $G_{\text{sum}}(S_t) \leftarrow G_{\text{sum}}(S_t) + W \cdot G$ 
14:       $W_{\text{sum}}(S_t) \leftarrow W_{\text{sum}}(S_t) + W$ 
15:       $V(S_t) \leftarrow \frac{G_{\text{sum}}(S_t)}{W_{\text{sum}}(S_t)}$                       ▷ Weighted average
16:    end if
17:  end for
18: end loop
```

Recap: from Tabular RL...

- Up until now, all our approaches for **prediction** and **control** are based on the availability of good estimations of v and q
- So far we have represented value function by a lookup table
 - > Every state s has an entry $V(s)$
 - > Or every state-action pair s, a has an entry $Q(s, a)$
- We are therefore dealing with the so-called **tabular** RL: we have look-up tables that we reply on



... to Value Function Approximation

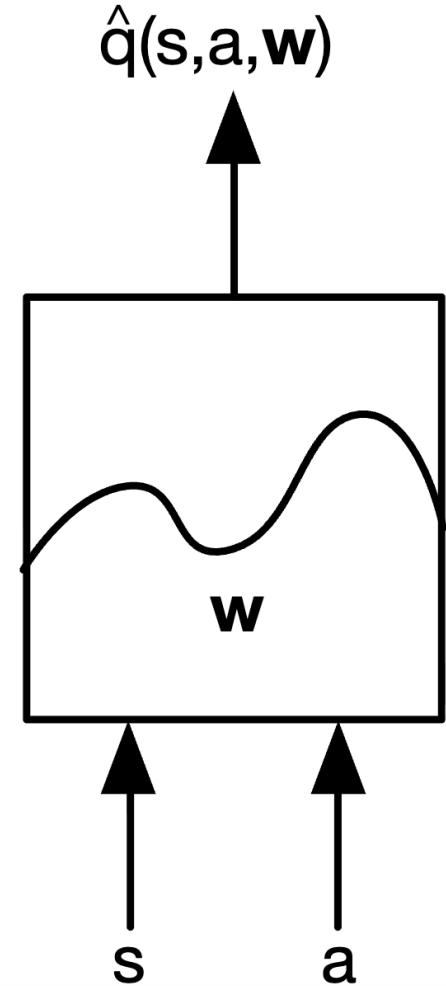
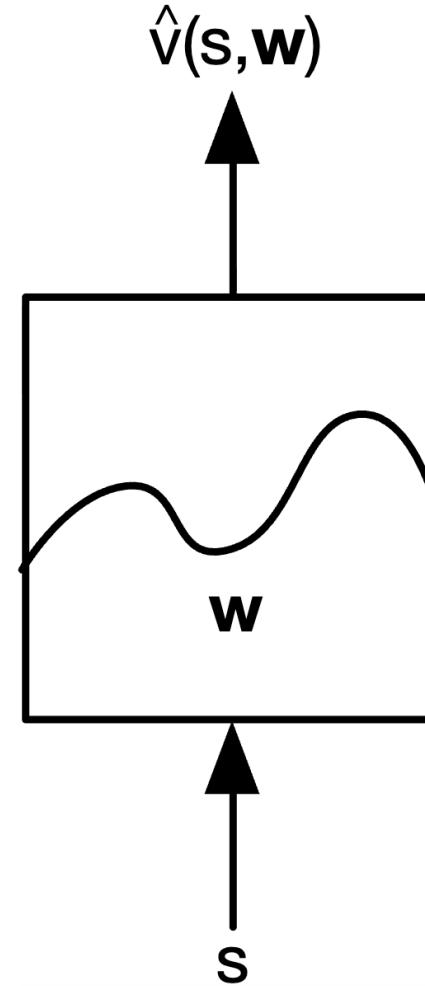
Solution for large MDPs:

- Estimate value function with function approximation

$$\hat{v}(s, \mathbf{w}) \approx v_{\pi}(s)$$

$$\hat{q}(s, a, \mathbf{w}) \approx q_{\pi}(s, a)$$

- Generalise from seen states to unseen states



Recap: Limits of Tabular RL

Unfortunately, in the naïve version, the tabular RL algorithms seen so far have some strong limitations: if the scale of problem become bigger, tables can become enormous

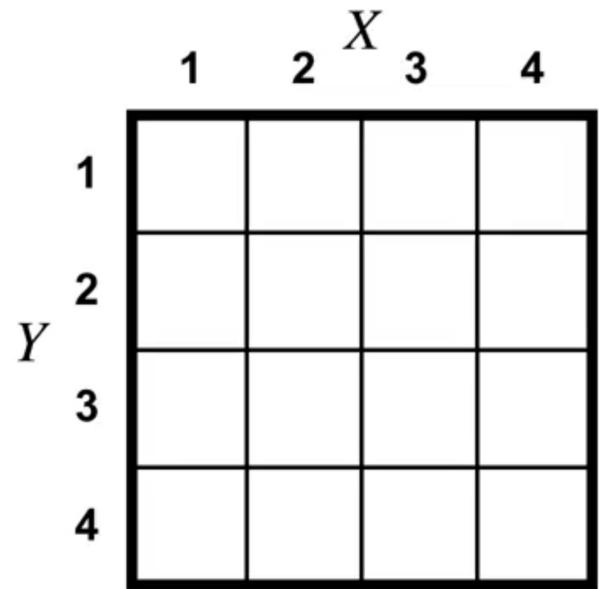
- > Backgammon: 10^{20} states
- > Go: 10^{170} states
- > Autonomous driving (continuous problem): continuous state space

Problems

- Memory: tables impossible to store
- Data/Time: to complete all entries on the tables with enough data to have good estimations, we need huge amount of data
- No generalization: a single update only impact an element of the table

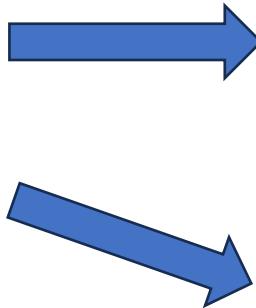
Generalization with VFA

$$\hat{v}(s, \mathbf{w}) \doteq w_1 X + w_2 Y$$



One change of parameters (1 step, 1 episode), **affect potentially the estimations for all states!**

$$w_1 = 1$$
$$w_2 = 1$$



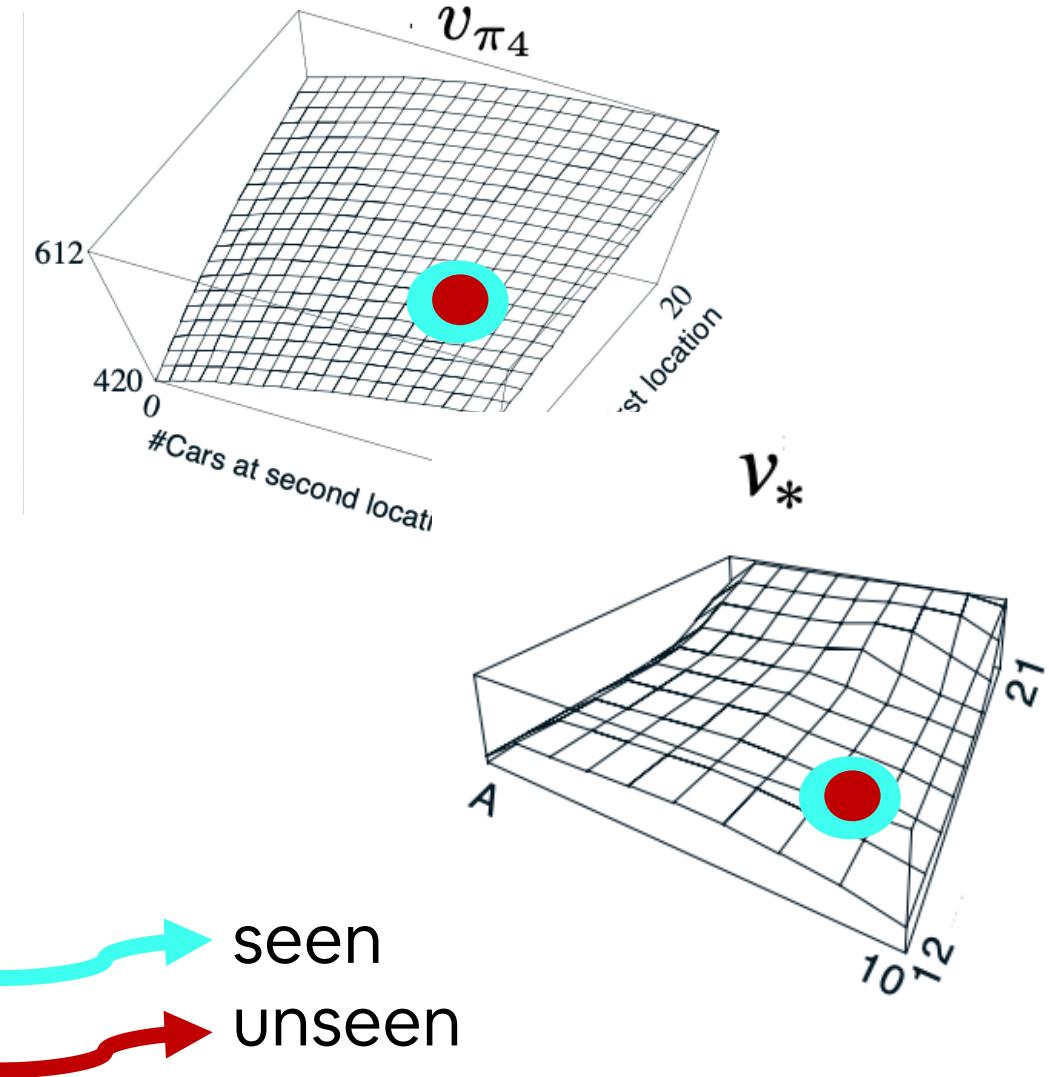
	1	2	X	3	4
1	2	3	4	5	
2	3	4	5	6	
3	4	5	6	7	
4	5	6	7	8	

$$w_1 = 4$$
$$w_2 = 1$$

	1	2	X	3	4
1	5	9	13	17	
2	6	10	14	18	
3	7	11	15	19	
4	8	12	16	20	

Generalization with VFA

Moreover, this is typically what we do in supervised learning! We don't have data for all possible situations; instead, we aim to find a function capable of generalizing—i.e., one that can describe the underlying phenomena for **unseen data**.



Recap: From true Q and V to approximation - example

Small grid-world: 16 states

Can we find an alternative, more ‘compact’, description for v?

For example, we can introduce **descriptors** (**features**) of the state $(x, y) = ([0, \dots, 3], [0, \dots, 3])$ we can see V is

$$V(x, y) = -(x + y)$$

A function approximator can, for example, have the form:

$$\hat{V}(s; \theta) = \theta_0 + \theta_1 x + \theta_2 y$$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

v_7

Recap: From true Q and V to approximation - example

Small grid-world: 16 states

Can we find an alternative, more ‘compact’, description for v?

For example, we can introduce **descriptors (features)** of the state $(x, y) = ([0, \dots, 3], [0, \dots, 3])$ we can see V is

$$V(x, y) = -(x + y)$$

A function approximator can, for example, have the form:

$$\hat{V}(s; \theta) = \theta_0 + \theta_1 x + \theta_2 y$$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

v_7

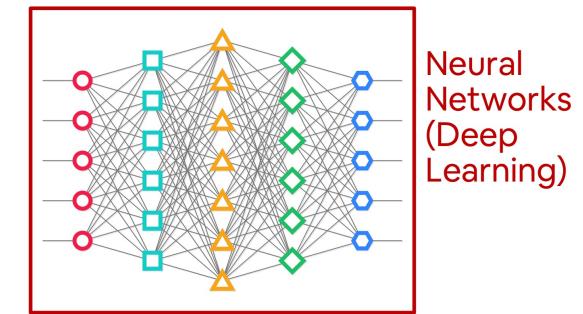
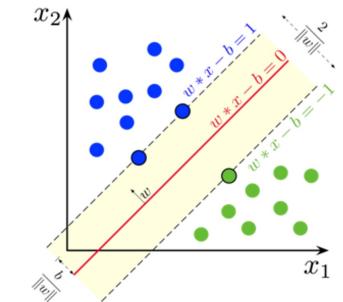
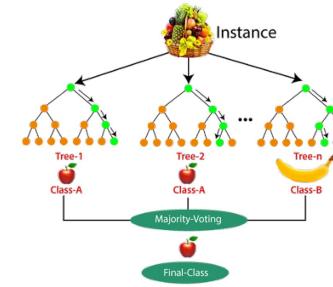
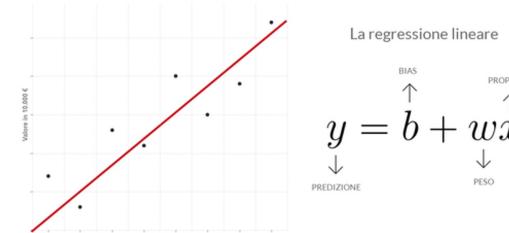
Depending on the problem we can come out with different type of features/compact representations of the problem

Recap: Which function approximator?

In principle all supervised learning approaches are fine!

We can use:

- Linear regression
- Neural networks / Deep Learning
- SVM
- Decision trees
- Ensemble approaches (Random Forest, XGBoost, ...)



However, we have a strong preference for **differentiable approaches**: some approach that I can learn along the way (ie. while I'm collecting experience interacting with the real world)

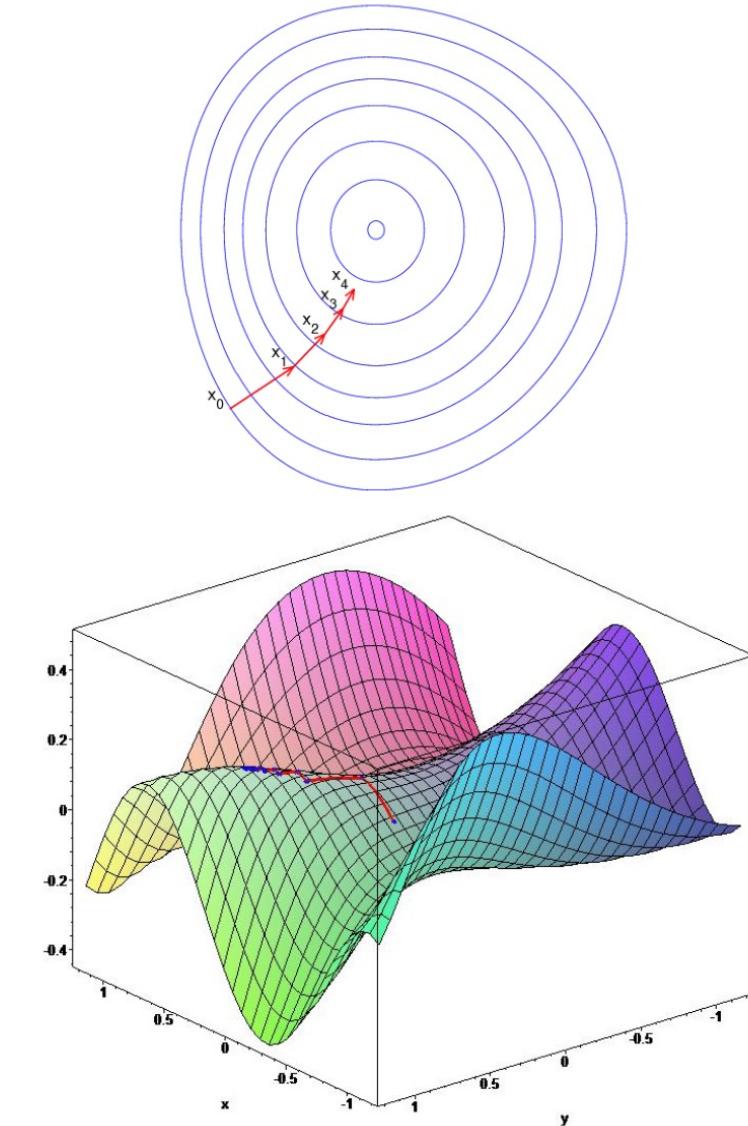
Differentiable Supervised Learning approaches

With:

- Linear approaches
- Neural networks / Deep Learning

we can use **stochastics gradient descent** and update (for example with new data)

The procedure is **incremental**: we can see this as 'learning along the way'



Differentiable Supervised Learning approaches

With:

- Linear approaches
- Neural networks / Deep Learning

we can use **stochastics gradient descent** and update (for example with new data)

The procedure is **incremental**: we can see this as 'learning along the way'

Linear approaches:

- Easy intuitions and guarantees
- Require the definition of features

Neural networks / Deep Learning approaches:

- Self-learning of features
- More complicated function approximators
- No guarantees and more complexity

Next week, lecture 15-16 will be dedicated to basics of NN/DL!

We need a Prediction Objective!

- In the tabular case we could aim for perfect matching of the true value function v_π . But when using function approximation, this is no longer feasible: changes in one state's estimate affect other states & therefore **not possible to get the values of all states exactly correct!**
- Moreover, **making one state's estimate more accurate might make others' less accurate!**
- We therefore need a formal objective that defines what “good” approximation means globally — over all states.
- Define the approximate value function $\hat{v}_\pi(s, \mathbf{w})$, parameterized by weight vector \mathbf{w}
- Define **a state-distribution $\mu(s) \geq 0$, with $\sum \mu(s) = 1$** (sum over all the states), which captures **how much we care about each state**.
- Then the objective is the so-called Mean-Squared Value Error (VE):

$$\overline{VE}(\mathbf{w}) \doteq \sum_{s \in \mathcal{S}} \mu(s) [v_\pi(s) - \hat{v}(s, \mathbf{w})]^2$$

We need a Prediction Objective!

- In on-policy continuous tasks $\mu(s)$ is the stationary distribution under π
- For simplicity we will consider this case onwards

$$\overline{\text{VE}}(\mathbf{w}) \doteq \sum_{s \in \mathcal{S}} \mu(s) [v_\pi(s) - \hat{v}(s, \mathbf{w})]^2$$



$$J(\mathbf{w}) = \mathbb{E}_\pi [(v_\pi(s) - \hat{v}(s, \mathbf{w}))^2]$$

Recap on Stochastic Gradient Descent (SGD)

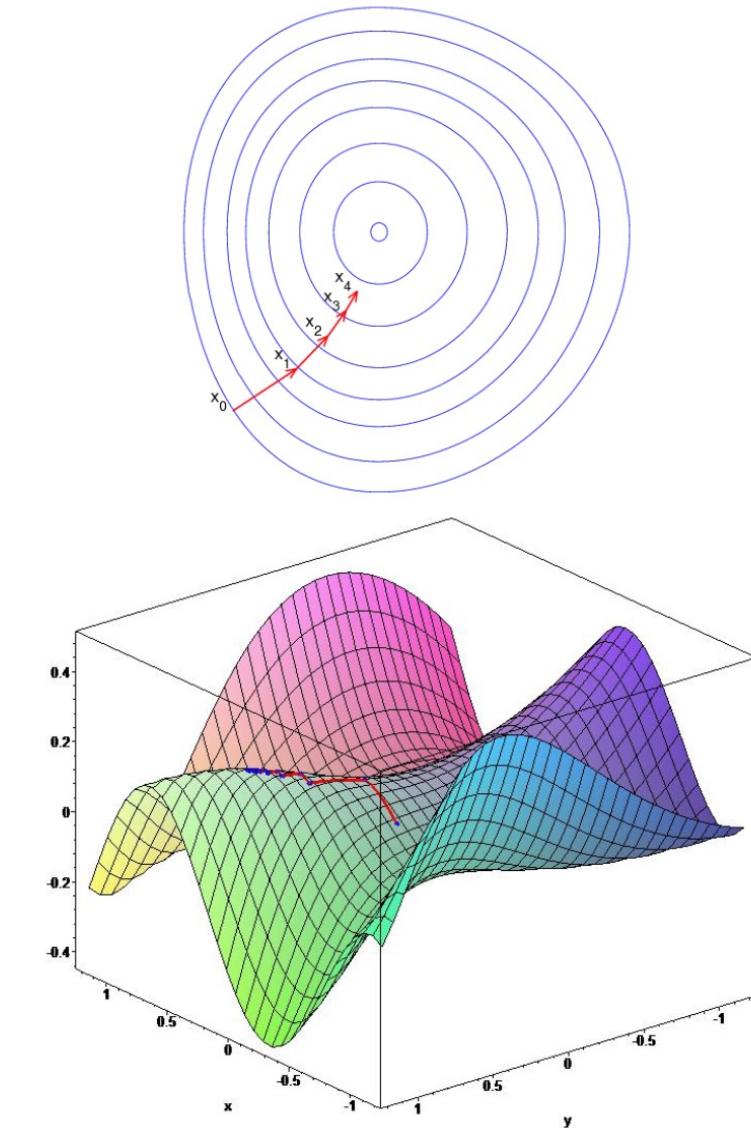
- Let $J(\mathbf{w})$ be a differentiable function of parameter vector \mathbf{w}
- Define the *gradient* of $J(\mathbf{w})$ to be

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \begin{pmatrix} \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}_1} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}_n} \end{pmatrix}$$

- To find a local minimum of $J(\mathbf{w})$
- Adjust \mathbf{w} in direction of -ve gradient

$$\Delta \mathbf{w} = -\frac{1}{2}\alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$

where α is a step-size parameter



Stochastic Gradient Descent (SGD)

- Goal: find parameter vector \mathbf{w} minimising mean-squared error between approximate value fn $\hat{v}(s, \mathbf{w})$ and true value fn $v_\pi(s)$

$$J(\mathbf{w}) = \mathbb{E}_\pi [(v_\pi(S) - \hat{v}(S, \mathbf{w}))^2]$$

- Gradient descent finds a local minimum

$$\begin{aligned}\Delta \mathbf{w} &= -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w}) \\ &= \alpha \mathbb{E}_\pi [(v_\pi(S) - \hat{v}(S, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w})]\end{aligned}$$

- Stochastic gradient descent *samples* the gradient

$$\Delta \mathbf{w} = \alpha (v_\pi(S) - \hat{v}(S, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w})$$

Only computable if the oracle is available!

A positive step-size parameter, following this condition to obtain convergence (more later)

$$\sum_{t=1}^{\infty} \alpha_t = \infty \quad \sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

- Expected update is equal to full gradient update $\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \nabla \mathbf{w}_t$

$$= \mathbf{w}_t + \alpha [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

In the linear case

- Represent value function by a linear combination of features

$$\hat{v}(S, \mathbf{w}) = \mathbf{x}(S)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S) w_j$$

- Objective function is quadratic in parameters \mathbf{w}

$$J(\mathbf{w}) = \mathbb{E}_\pi \left[(v_\pi(S) - \mathbf{x}(S)^\top \mathbf{w})^2 \right]$$

- Stochastic gradient descent converges on *global* optimum
- Update rule is particularly simple

$$\nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w}) = \mathbf{x}(S)$$

$$\Delta \mathbf{w} = \alpha (v_\pi(S) - \hat{v}(S, \mathbf{w})) \mathbf{x}(S)$$

Update = *step-size* \times *prediction error* \times *feature value*

In the linear case

- Represent value function by a linear combination of features

$$\hat{v}(S, \mathbf{w}) = \mathbf{x}(S)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S) w_j$$

- Objective function is quadratic in parameters \mathbf{w}

$$J(\mathbf{w}) = \mathbb{E}_{\pi} \left[(v_{\pi}(S) - \mathbf{x}(S)^\top \mathbf{w})^2 \right]$$

- Stochastic gradient descent converges on *global* optimum
- Update rule is particularly simple

$$\nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w}) = \mathbf{x}(S)$$

$$\Delta \mathbf{w} = \alpha (v_{\pi}(S) - \hat{v}(S, \mathbf{w})) \mathbf{x}(S)$$

Update = step-size \times prediction error \times feature value

Tabular RL is a special case of linear VFA

$$\mathbf{x}^{table}(S) = \begin{pmatrix} \mathbf{1}(S = s_1) \\ \vdots \\ \mathbf{1}(S = s_n) \end{pmatrix}$$

$$\hat{v}(S, \mathbf{w}) = \begin{pmatrix} \mathbf{1}(S = s_1) \\ \vdots \\ \mathbf{1}(S = s_n) \end{pmatrix} \cdot \begin{pmatrix} \mathbf{w}_1 \\ \vdots \\ \mathbf{w}_n \end{pmatrix}$$

Recap: Supervised Learning

- Supervised learning approaches requires tagged examples! We will need a dataset of examples

$$(s_1, v(s_1)), (s_2, v(s_2)), \dots, (s_D, v(s_D))$$

$$(s_1, a_1, q(s_1, a_1)), (s_2, a_2, q(s_2, a_2)), \dots, (s_D, a_D, q(s_D, a_D))$$



Core idea of Value Function Approximation

To create a supervised setting, instead of the true v and q , we use a **target**

- Monte Carlo target: the return

$$(s_1, G_t^1), (s_2, G_t^2), \dots, (s_D, G_t^D)$$

- TD(0) target: the TD target

$$(s_1, R_{t+1}^1 + \gamma \hat{v}(S_{t+1}^1, \mathbf{w})), (s_2, R_{t+1}^2 + \gamma \hat{v}(S_{t+1}^2, \mathbf{w})), \dots, (s_D, R_{t+1}^D + \gamma \hat{v}(S_{t+1}^D, \mathbf{w}))$$

- TD(λ) target: the λ -return

$$(s_1, G_1^\lambda), (s_2, G_2^\lambda), \dots, (s_D, G_D^\lambda)$$

Value Function Approximation: Monte Carlo

- Return G_t is an unbiased, noisy sample of true value $v_\pi(S_t)$
- Can therefore apply supervised learning to “training data”:

$$\langle S_1, G_1 \rangle, \langle S_2, G_2 \rangle, \dots, \langle S_T, G_T \rangle$$

- For example, using *linear Monte-Carlo policy evaluation*

$$\begin{aligned}\Delta \mathbf{w} &= \alpha(\textcolor{red}{G_t} - \hat{v}(S_t, \mathbf{w}))\nabla_{\mathbf{w}}\hat{v}(S_t, \mathbf{w}) \\ &= \alpha(G_t - \hat{v}(S_t, \mathbf{w}))\mathbf{x}(S_t)\end{aligned}$$

- Monte-Carlo evaluation converges to a local optimum
- Even when using non-linear value function approximation

Prediction - Value Function Approximation: Monte Carlo

Gradient Monte Carlo Algorithm for Estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameter: step size $\alpha > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop forever (for each episode):

 Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T$ using π

 Loop for each step of episode, $t = 0, 1, \dots, T - 1$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

Being the MC target (the return), an unbiased estimator, the gradient-descent version of MC prediction converges to a:

- [linear case] globally optimal approximation of v_π
- [non-linear case] locally optimal approximation of v_π

Prediction - Value Function Approximation: Monte Carlo

Gradient Monte Carlo Algorithm for Estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameter: step size $\alpha > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop forever (for each episode):

Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T$ using π

Loop for each step of episode, $t = 0, 1, \dots, T - 1$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

We don't have
the same
guarantees with
TD(0)* or
TD(lambda)

Being the MC target (the return), an unbiased estimator, the gradient-descent version of MC prediction converges to a:

- [linear case] globally optimal approximation of v_π
- [non-linear case] locally optimal approximation of v_π

* But we are
'close'

Value Function Approximation: TD(0)

- The TD-target $R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})$ is a *biased* sample of true value $v_\pi(S_t)$
- Can still apply supervised learning to “training data”:

$$\langle S_1, R_2 + \gamma \hat{v}(S_2, \mathbf{w}) \rangle, \langle S_2, R_3 + \gamma \hat{v}(S_3, \mathbf{w}) \rangle, \dots, \langle S_{T-1}, R_T \rangle$$

- For example, using *linear TD(0)*

$$\begin{aligned}\Delta \mathbf{w} &= \alpha(R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w}) \\ &= \alpha \delta \mathbf{x}(S)\end{aligned}$$

- Linear TD(0) converges (close) to global optimum

Prediction - Value Function Approximation: TD(0)

Semi-gradient TD(0) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameter: step size $\alpha > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose $A \sim \pi(\cdot | S)$

 Take action A , observe R, S'

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$$

$S \leftarrow S'$

 until S is terminal

Why ‘semi-gradient’?

An important observation:

$$\Delta \mathbf{w} = \alpha(v_\pi(S) - \hat{v}(S, \mathbf{w}))\nabla_{\mathbf{w}}\hat{v}(S, \mathbf{w})$$

Target does not depend on \mathbf{w}

$$\Delta \mathbf{w} = \alpha(R + \gamma\hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w}))\nabla_{\mathbf{w}}\hat{v}(S, \mathbf{w})$$

It depends on \mathbf{w} (a typical property of all bootstrapping methods)

They are not true gradient descent method: *they take into account the effect of changing the weight vector w_t on the estimate but ignore its effect on the target*

Value Function Approximation: TD(lambda)

- The λ -return G_t^λ is also a biased sample of true value $v_\pi(s)$
- Can again apply supervised learning to “training data”:

$$\langle S_1, G_1^\lambda \rangle, \langle S_2, G_2^\lambda \rangle, \dots, \langle S_{T-1}, G_{T-1}^\lambda \rangle$$

- Forward view linear TD(λ)

$$\begin{aligned}\Delta \mathbf{w} &= \alpha(G_t^\lambda - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w}) \\ &= \alpha(G_t^\lambda - \hat{v}(S_t, \mathbf{w})) \mathbf{x}(S_t)\end{aligned}$$

- Backward view linear TD(λ)

$$\begin{aligned}\delta_t &= R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}) \\ E_t &= \gamma \lambda E_{t-1} + \mathbf{x}(S_t) \\ \Delta \mathbf{w} &= \alpha \delta_t E_t\end{aligned}$$

Value Function Approximation: TD(lambda)

- The λ -return G_t^λ is also a biased sample of true value $v_\pi(s)$
- Can again apply supervised learning to “training data”:

$$\langle S_1, G_1^\lambda \rangle, \langle S_2, G_2^\lambda \rangle, \dots, \langle S_{T-1}, G_{T-1}^\lambda \rangle$$

Each ‘couple’
has info for
many states

- Forward view linear TD(λ)

$$\begin{aligned}\Delta \mathbf{w} &= \alpha(G_t^\lambda - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w}) \\ &= \alpha(G_t^\lambda - \hat{v}(S_t, \mathbf{w})) \mathbf{x}(S_t)\end{aligned}$$

- Backward view linear TD(λ)

$$\begin{aligned}\delta_t &= R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}) \\ E_t &= \gamma \lambda E_{t-1} + \mathbf{x}(S_t) \\ \Delta \mathbf{w} &= \alpha \delta_t E_t\end{aligned}$$

Prediction - Value Function Approxim ation: n- step TD

n -step semi-gradient TD for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameters: step size $\alpha > 0$, a positive integer n

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

All store and access operations (S_t and R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take an action according to $\pi(\cdot | S_t)$

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

 If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n \hat{v}(S_{\tau+n}, \mathbf{w})$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G - \hat{v}(S_\tau, \mathbf{w})] \nabla \hat{v}(S_\tau, \mathbf{w})$$

 Until $\tau = T - 1$

$$(G_{\tau:\tau+n})$$

Prediction - Value Function Approximation: convergence

Method	Linear FA	Nonlinear FA
TD(0)	✓ converges	✗ no guarantee, can diverge
TD(λ)	✓ converges for all λ	✗ no guarantee
n-step TD	✓ converges	✗ no guarantee
MC	✓ globally	✓ <i>locally</i> (unbiased target avoids bootstrapping issues)

What about features? How to build them?

- Linear learning methods are of great interest:
 - convergence properties
 - very efficient in terms of both data and computation

...but the overall goodness depends on how the states are represented in terms of the **features!**
- Adding prior domain knowledge to RL systems is important for feature design (it is usually important, also in other phases of an RL and ML project)

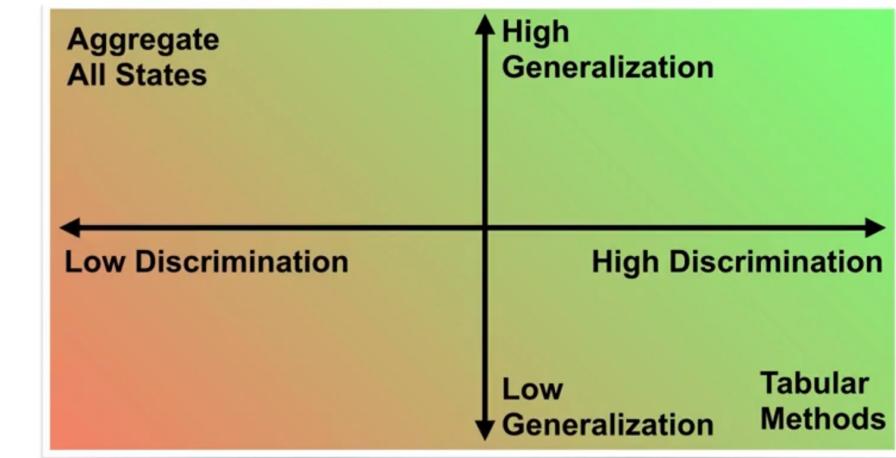
Ie. HVAC system -> room occupancy, hour of the day, month, current temperature...

Ie. Mobile robot -> locations, remaining battery, recent sensor readings...

What about features? How to build them?

In the search of ‘good’ features, we need a good trade-off between:

- **Generalization**, ie. how features allow the value function approximator to treat similar states as having similar values.
- **Discrimination**, ie. how features allow the approximator to distinguish states that have different values.



Generalization is important because:

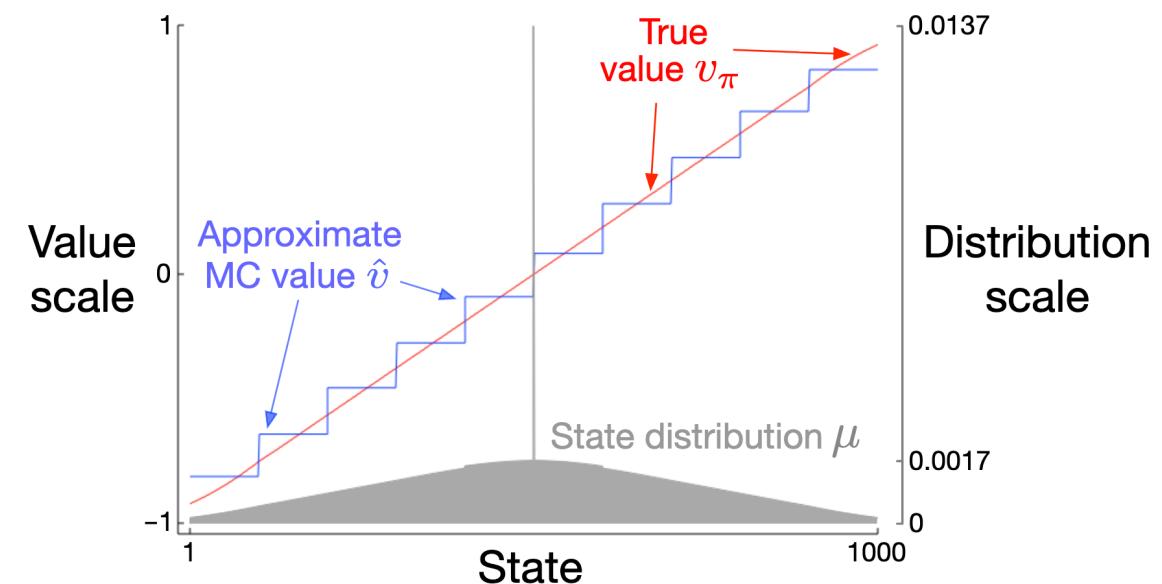
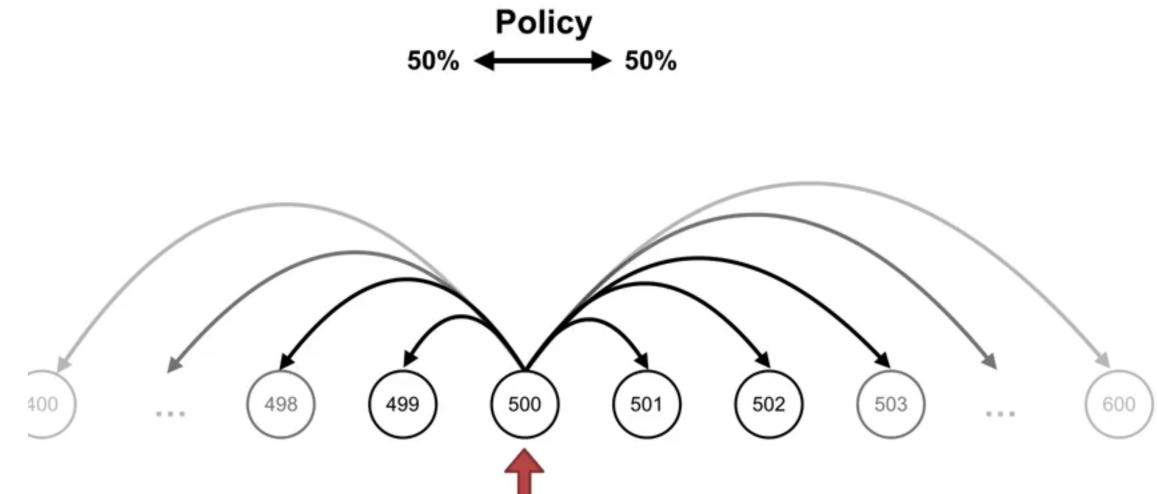
- In large or continuous state spaces, we cannot learn a separate value for every state.
- Instead, the agent must generalize information learned from one state to nearby or similar states.

Discrimination is important because:

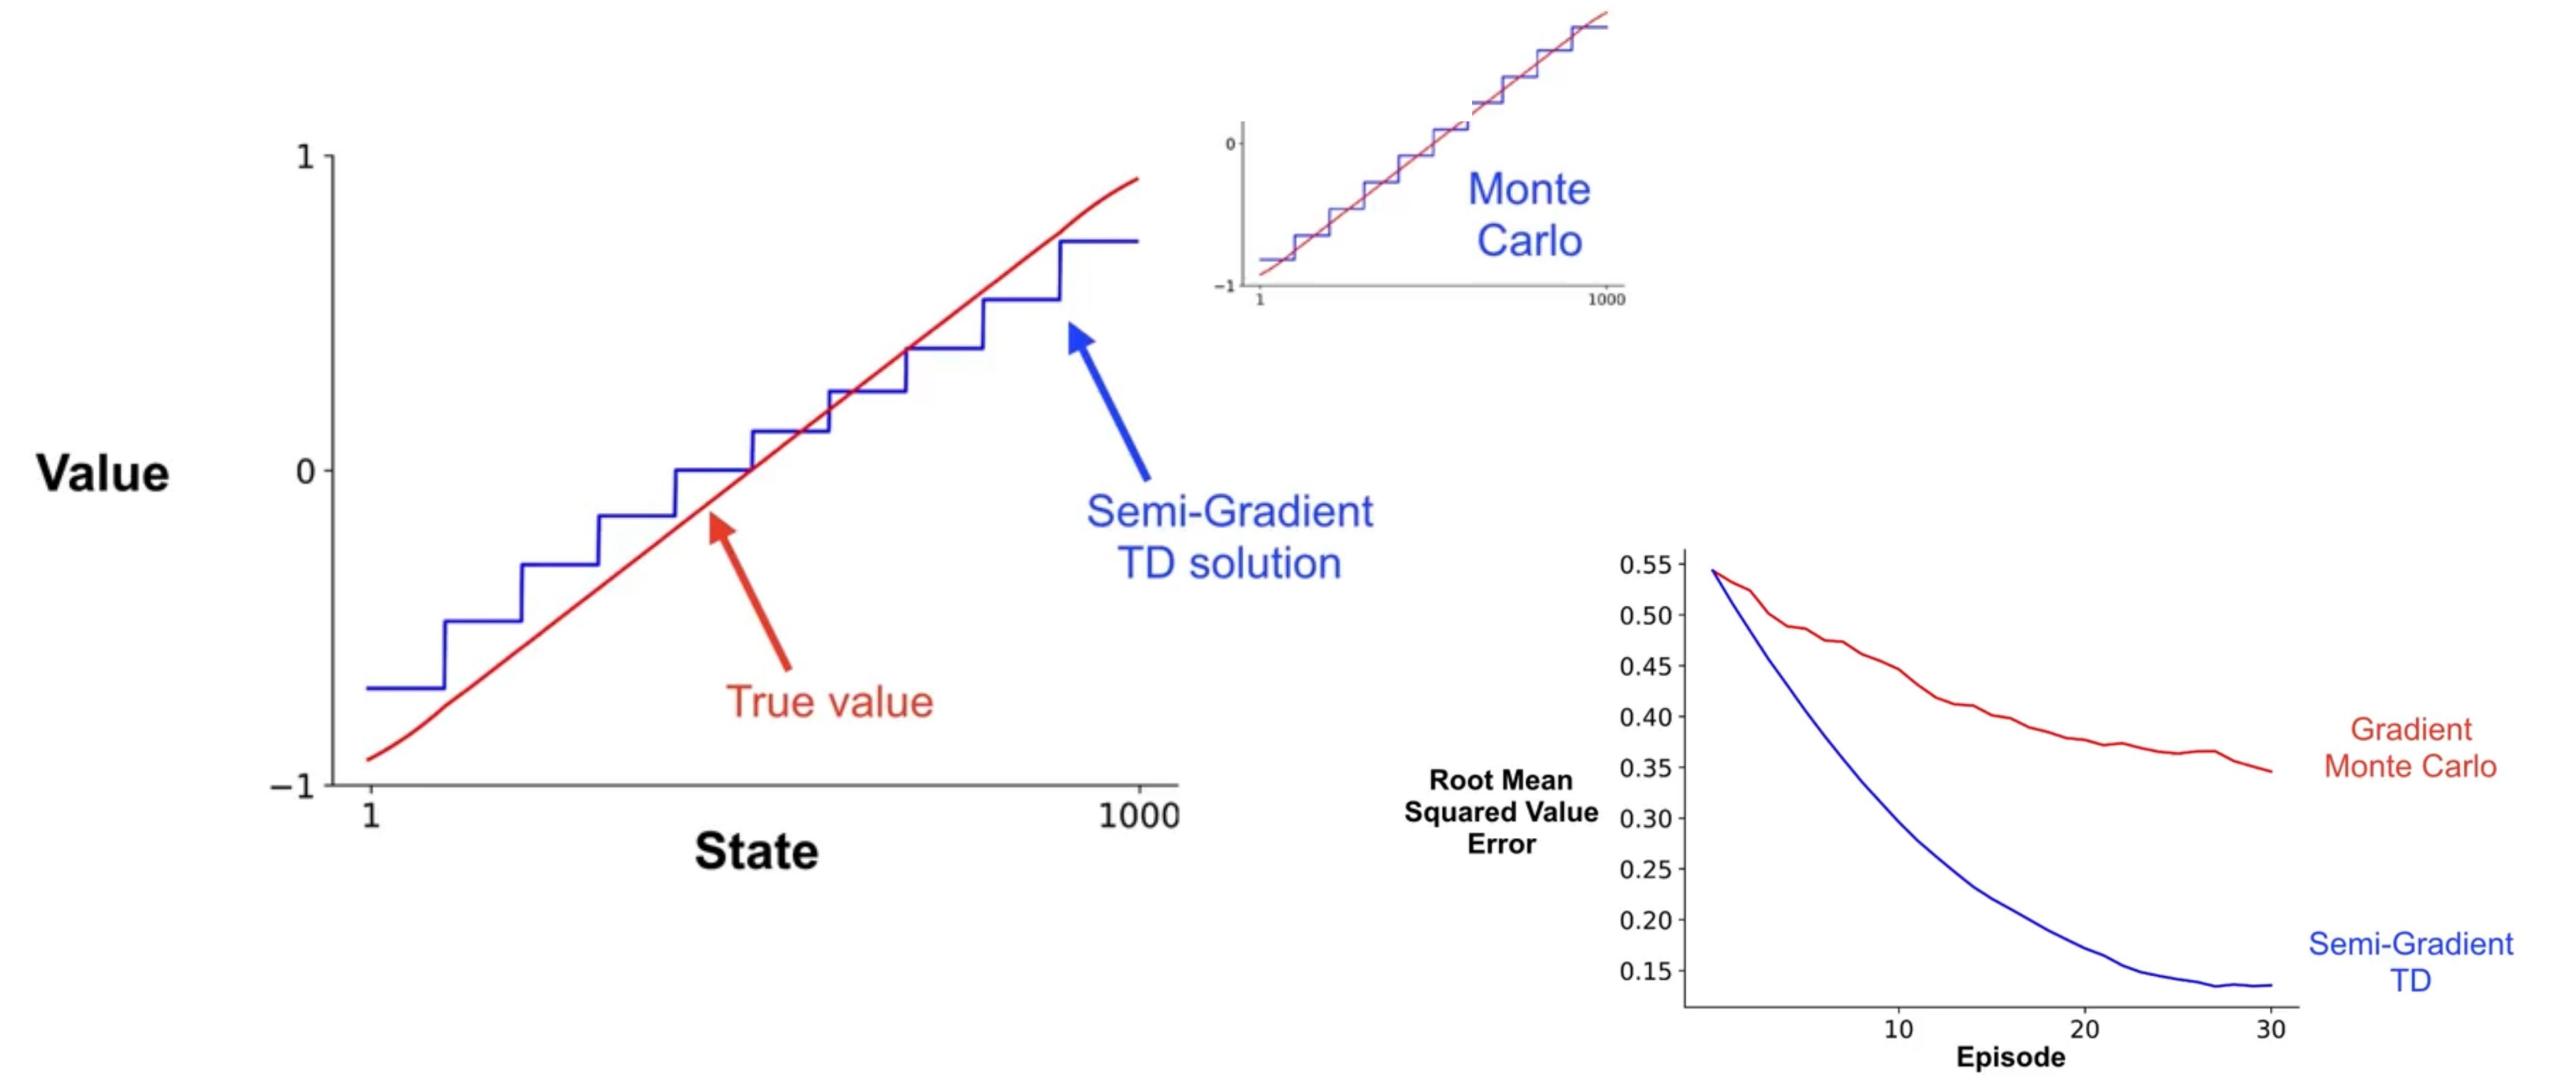
- If two states must be valued differently for optimal behavior, the function approximator must have the capacity to separate them.
- Without discrimination, value estimates will blur together, causing suboptimal policy learning.

Feature Design #01: state aggregation

- A simple form of generalizing function approximation in which states are grouped together, with one estimated value (one component of the weight vector w) for each group
- Large scale random walk example in the book

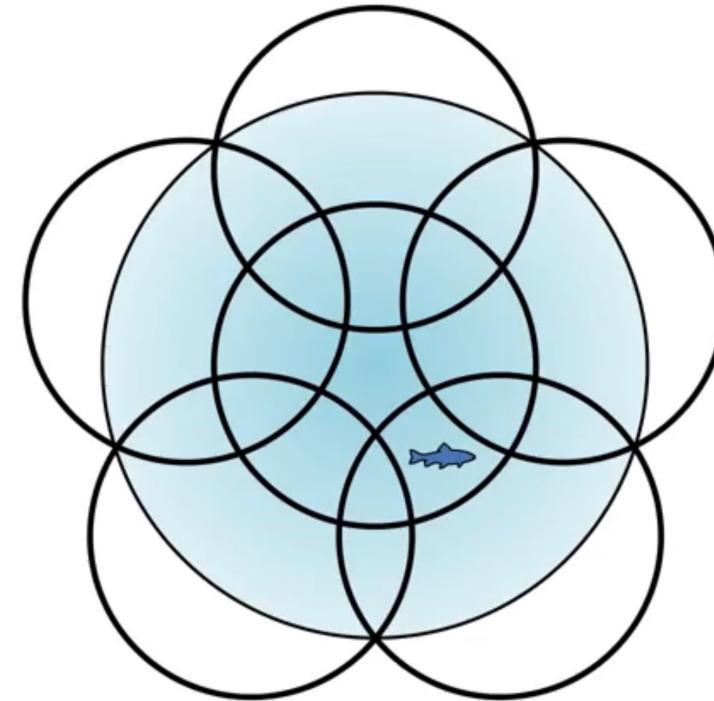
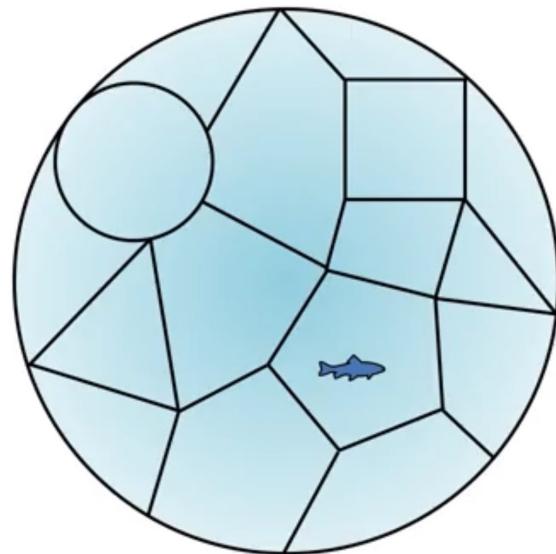


Feature Design #01: state aggregation, large random walk



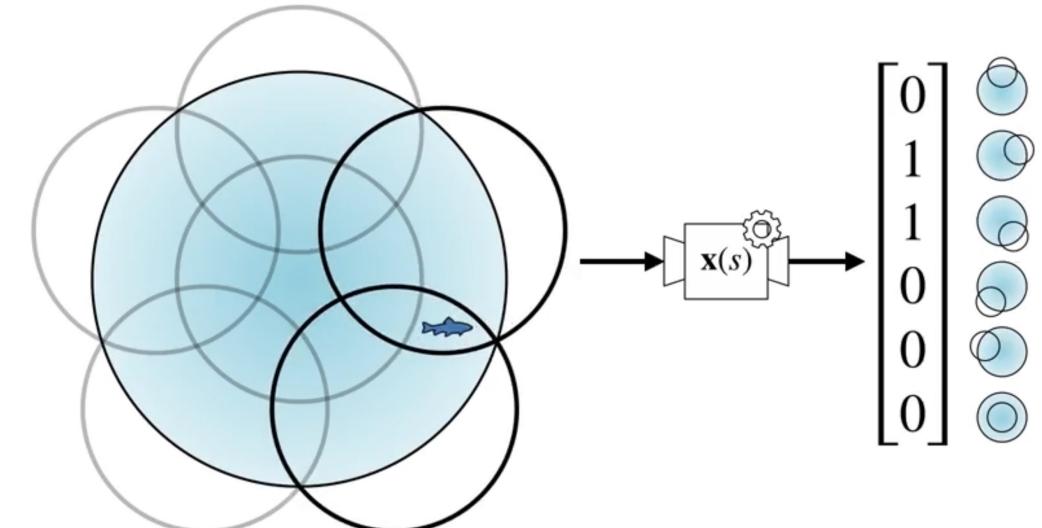
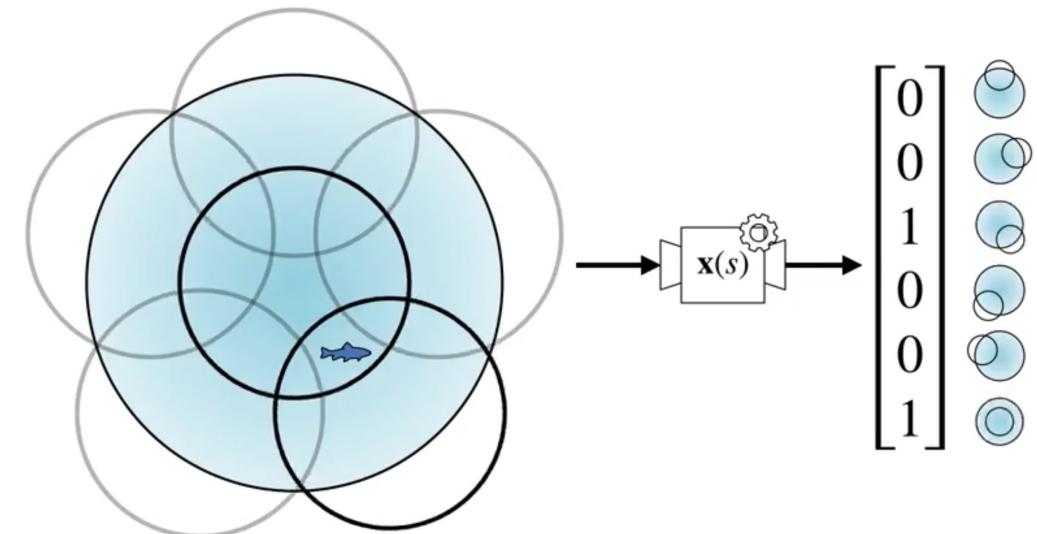
Feature Design #02: coarse coding

- Coarse coding is a feature-representation method where each feature responds to a broad region of the state space, and many features **overlap**.



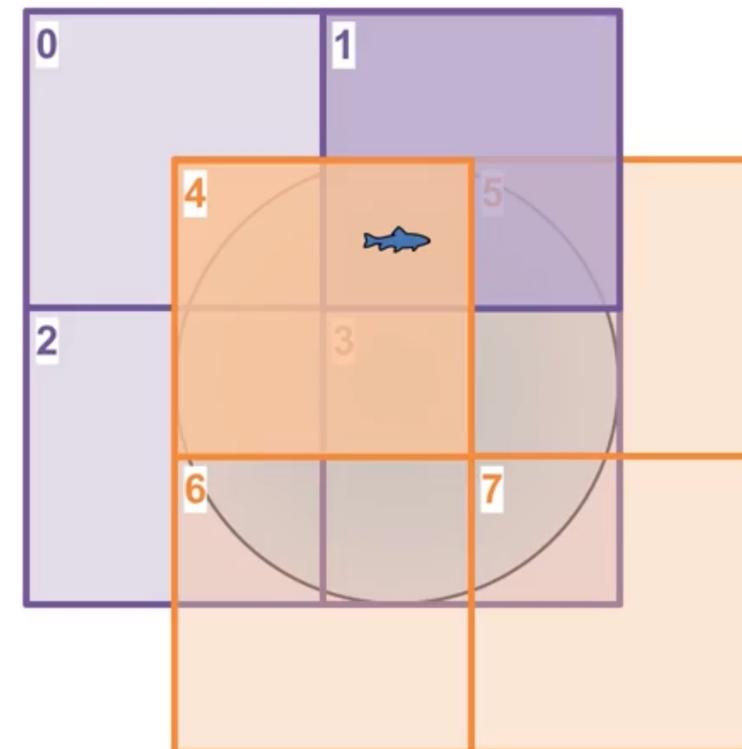
Feature Design #02: coarse coding

- With coarse coding we have a set of features that each respond to a coarse (wide) region of the state space; any given state activates several of these **overlapping features**, and the unique pattern of activations represents the state
- Coarse coding uses fewer, broader features, making the **representation compact**, but because states activate different overlapping subsets of these features, **it can still capture fine-grained distinctions**.
- Aggregated states behave like one-hot features, while coarse coding is more efficient because each state activates multiple overlapping features instead of just one.



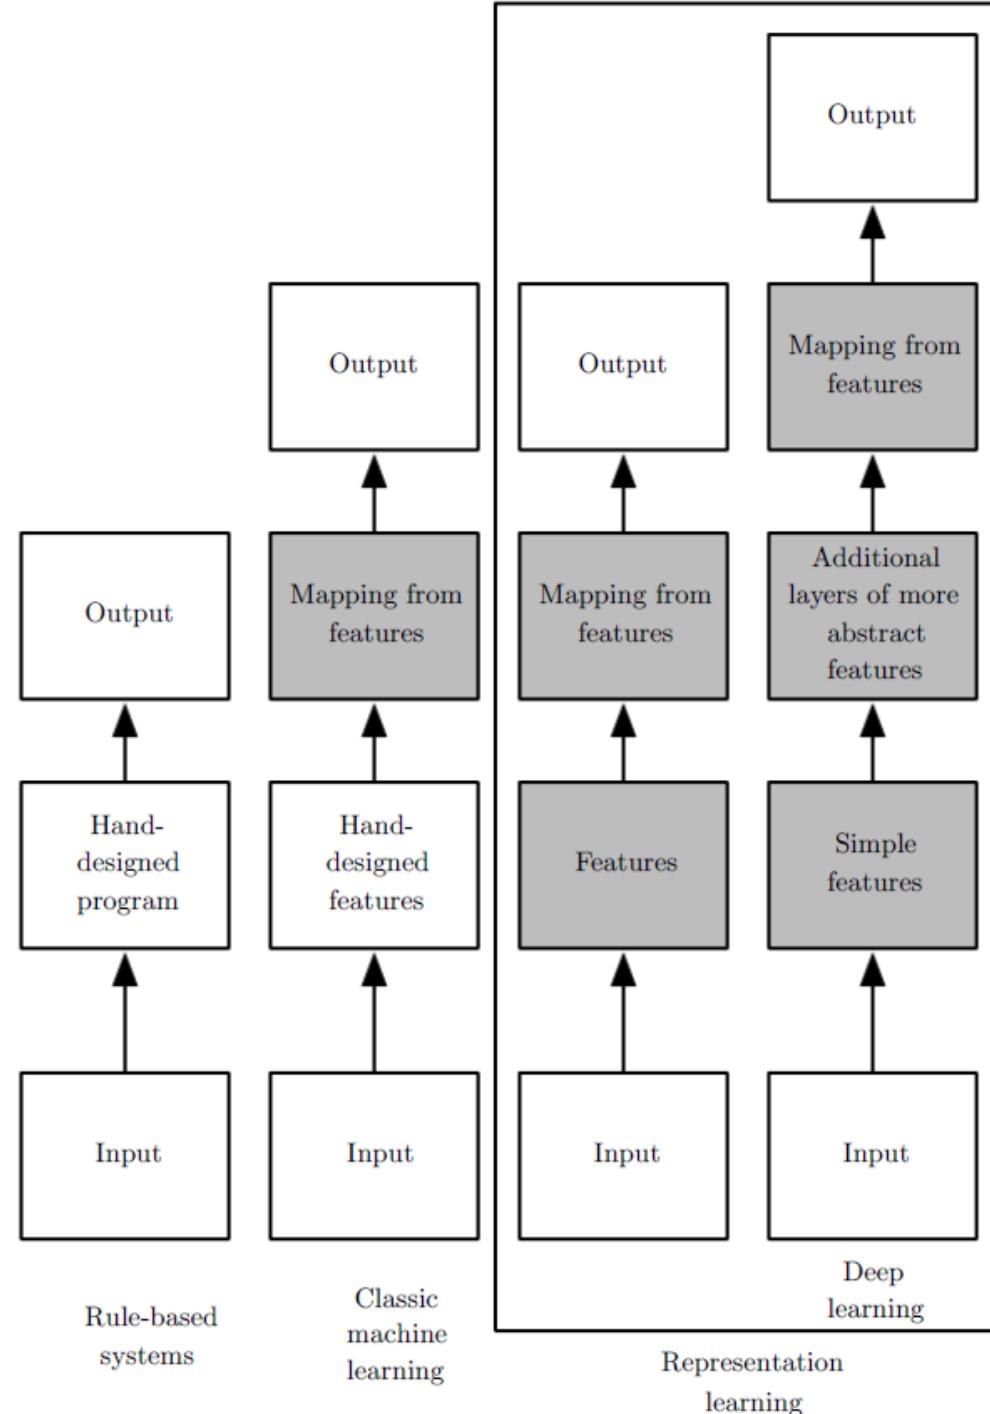
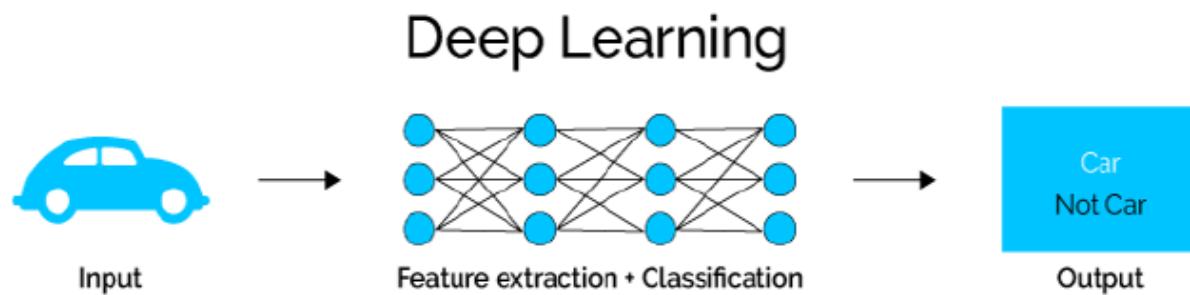
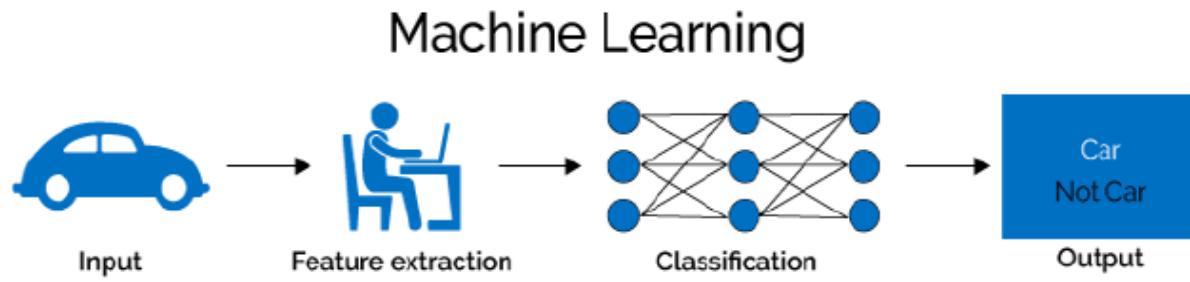
Feature Design #03: tile coding

- A special, structured case of coarse coding is tile coding
- Tile coding: **rectangular regions** arranged in multiple **overlapping grids (tilings)**
- Very efficient, sparse and easy to use in RL
- Overlap is binary (inside tile = 1, outside = 0)
- Fixed costs at each iteration (we are always updating the same amount of weight)

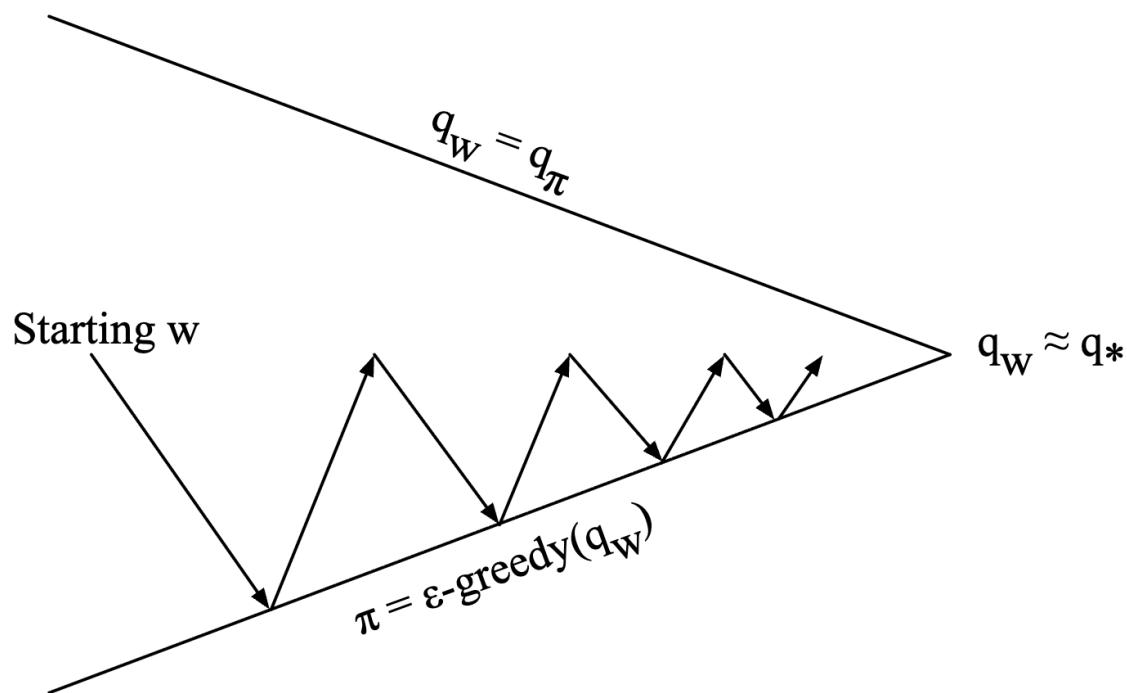


$$\mathbf{w} = \begin{bmatrix} 0.0 \\ 1.5 \\ 1.1 \\ 0.2 \\ 0.5 \\ -0.3 \\ 1.3 \\ -0.7 \end{bmatrix} \quad \mathbf{x}(s) = \begin{bmatrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \\ 7 \end{bmatrix}$$

With Deep Learning, it is easier!



Control - Even if we are dealing with a new ‘framework’, we leverage on the algorithms we have already seen!



Policy evaluation **Approximate** policy evaluation, $\hat{q}(\cdot, \cdot, \mathbf{w}) \approx q_\pi$
Policy improvement ϵ -greedy policy improvement

Control - VFA

- Approximate the action-value function

$$\hat{q}(S, A, \mathbf{w}) \approx q_\pi(S, A)$$

- Minimise mean-squared error between approximate action-value fn $\hat{q}(S, A, \mathbf{w})$ and true action-value fn $q_\pi(S, A)$

$$J(\mathbf{w}) = \mathbb{E}_\pi [(q_\pi(S, A) - \hat{q}(S, A, \mathbf{w}))^2]$$

- Use stochastic gradient descent to find a local minimum

$$-\frac{1}{2} \nabla_{\mathbf{w}} J(\mathbf{w}) = (q_\pi(S, A) - \hat{q}(S, A, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w})$$

$$\Delta \mathbf{w} = \alpha (q_\pi(S, A) - \hat{q}(S, A, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w})$$

Control - VFA

- Represent state *and* action by a *feature vector*

$$\mathbf{x}(S, A) = \begin{pmatrix} \mathbf{x}_1(S, A) \\ \vdots \\ \mathbf{x}_n(S, A) \end{pmatrix}$$

- Represent action-value fn by linear combination of features

$$\hat{q}(S, A, \mathbf{w}) = \mathbf{x}(S, A)^\top \mathbf{w} = \sum_{j=1}^n \mathbf{x}_j(S, A) \mathbf{w}_j$$

- Stochastic gradient descent update

$$\nabla_{\mathbf{w}} \hat{q}(S, A, \mathbf{w}) = \mathbf{x}(S, A)$$

$$\Delta \mathbf{w} = \alpha(q_\pi(S, A) - \hat{q}(S, A, \mathbf{w})) \mathbf{x}(S, A)$$

Control - VFA

- Like prediction, we must substitute a *target* for $q_\pi(S, A)$

- For MC, the target is the return G_t

$$\Delta \mathbf{w} = \alpha(G_t - \hat{q}(S_t, A_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S_t, A_t, \mathbf{w})$$

- For TD(0), the target is the TD target $R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$

$$\Delta \mathbf{w} = \alpha(R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}) - \hat{q}(S_t, A_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S_t, A_t, \mathbf{w})$$

- For forward-view TD(λ), target is the action-value λ -return

$$\Delta \mathbf{w} = \alpha(q_t^\lambda - \hat{q}(S_t, A_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{q}(S_t, A_t, \mathbf{w})$$

- For backward-view TD(λ), equivalent update is

$$\delta_t = R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}) - \hat{q}(S_t, A_t, \mathbf{w})$$

$$E_t = \gamma \lambda E_{t-1} + \nabla_{\mathbf{w}} \hat{q}(S_t, A_t, \mathbf{w})$$

$$\Delta \mathbf{w} = \alpha \delta_t E_t$$

Episodic Semi-gradient Sarsa for Estimating $\hat{q} \approx q_*$

Input: a differentiable action-value function parameterization $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameters: step size $\alpha > 0$, small $\varepsilon > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'

If S' is terminal:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

Go to next episode

Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

$$S \leftarrow S'$$

$$A \leftarrow A'$$

Episodic Semi-gradient Sarsa for Estimating $\hat{q} \approx q_*$

Input: a differentiable action-value function parameterization $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameters: step size $\alpha > 0$, small $\varepsilon > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'

If S' is terminal:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

Go to next episode

Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

$$S \leftarrow S'$$

$$A \leftarrow A'$$

We can also have optimistic value initialization strategies to foster exploration at the early stages of the episode

Episodic Semi-gradient Sarsa for Estimating $\hat{q} \approx q_*$

Input: a differentiable action-value function parameterization $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameters: step size $\alpha > 0$, small $\varepsilon > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'

If S' is terminal:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

Go to next episode

Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

$S \leftarrow S'$

$A \leftarrow A'$

The update of the parameters of q is in fact the improvement step!

Episodic Semi-gradient Sarsa for Estimating $\hat{q} \approx q_*$

Input: a differentiable action-value function parameterization $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameters: step size $\alpha > 0$, small $\varepsilon > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'

If S' is terminal:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

Go to next episode

Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$$

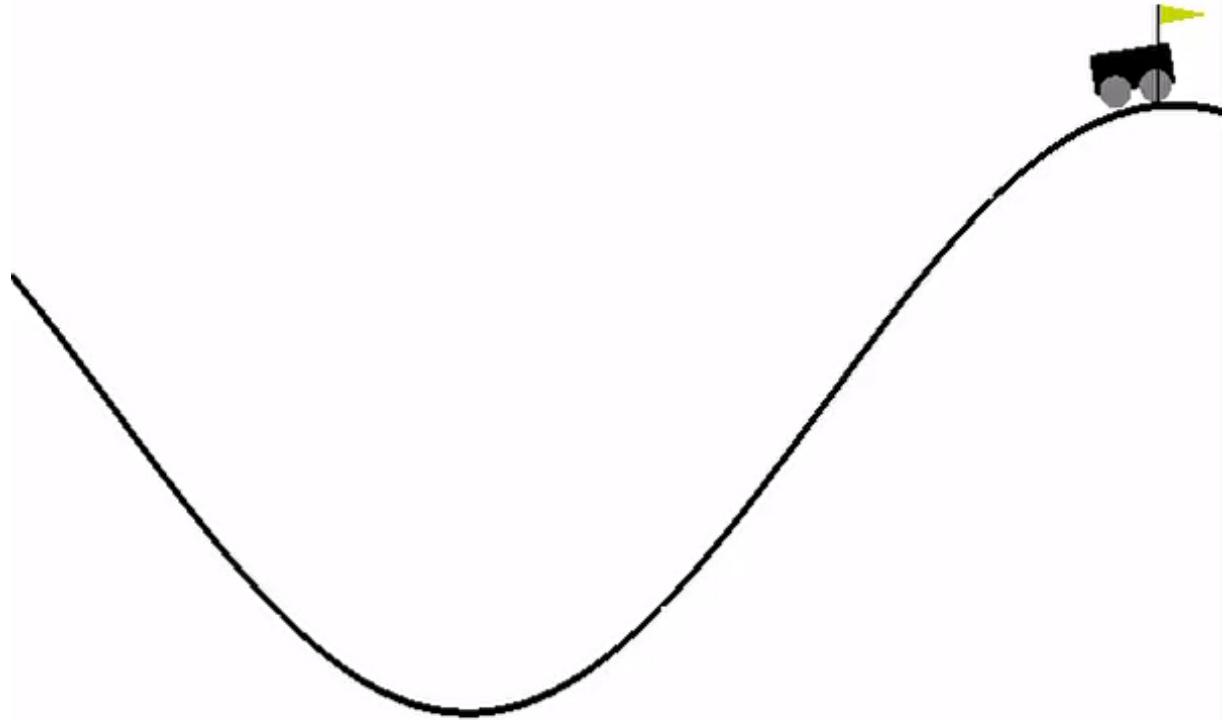
$$S \leftarrow S'$$

$$A \leftarrow A'$$

On policy

Example: Mountain Car Problem

Example: Mountain
Car problem
(Example 10.1 of the
book)



D. Unzueta <https://towardsdatascience.com/reinforcement-learning-applied-to-the-mountain-car-problem-1c4fb16729ba>

Example: Mountain Car Problem

Example: Mountain
Car problem
(Example 10.1 of the
book)

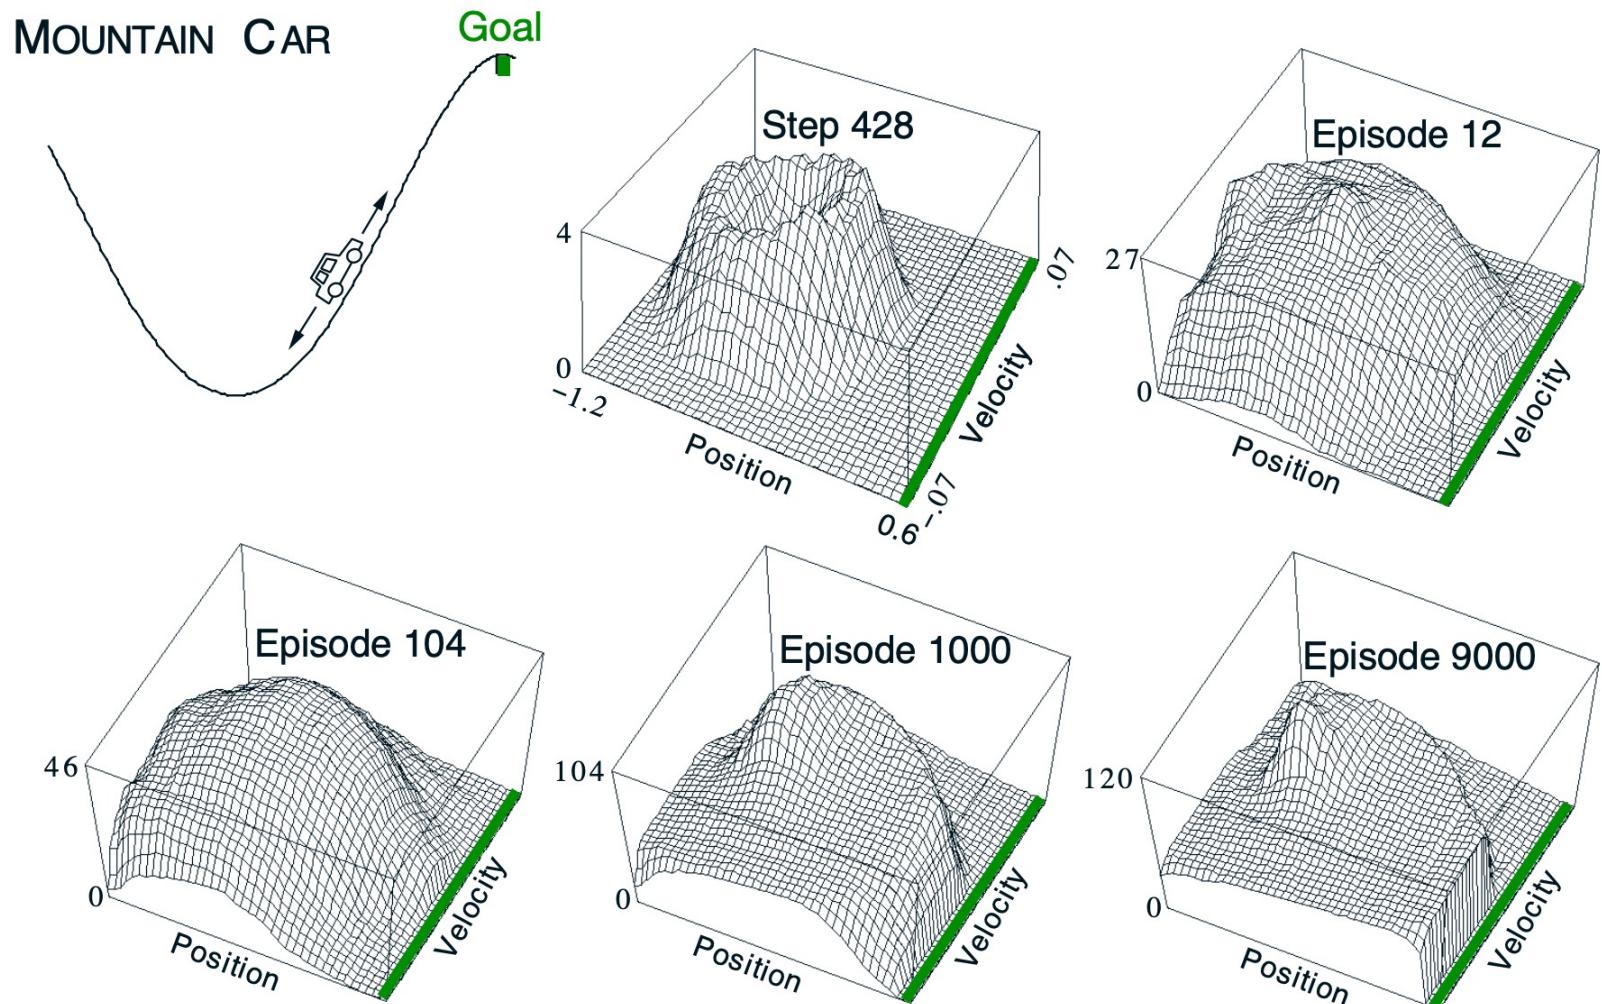


Figure 10.1: The Mountain Car task (upper left panel) and the cost-to-go function ($-\max_a \hat{q}(s, a, \mathbf{w})$) learned during one run.

Credits

- David Silver ‘Lecture on RL’, UCL 2015
- Image of the course is taken from C. Mahoney ‘Reinforcement Learning’ <https://towardsdatascience.com/reinforcement-learning-fda8ff535bb6>

**Thank you!
Questions?**

**Lecture #14: Value Function
Approximation**

Gian Antonio Susto

